


Performance, Scalability Messaging

Ashish Sharma

Final Project

- 
1. Team Name
Team Members (1 – 3)
Feb 29
 2. Project Proposal
March 14
 3. Project Update
Apr 6
 4. Final Presentation
Apr 25-27



Quiz: Recap

Q1. Specs of the fastest machine

Q2. Hierarchy vs. Locality

- 3120000 cores; 54.9 PFLops
- Memory hierarchy refers to the hierarchy in speed and storage that is observed in accessible memory. Ranges from the fastest and smallest (L1 cache), to the slowest and largest (disk)
- Memory Locality refers to the proximity in space (contiguous memory blocks) or time (access the same block across multiple cycles)

Q3. Spatial vs. Temporal Locality

Q4. R_{PEAK} of quad core, 2.25Ghz, 6ops/
cycle

- Reduce Latency
→ Temporal Locality
- Increase Bandwidth
→ Spatial Locality
- $2.25 * 10^9 * 4 * 6 = 54\text{GFlops}$

Q5. 32KB L1 cache with a 1ns latency
4GBs via a bus with 100ns latency
bus is 1 word wide
Multiply 32x32

$$\text{FLOPS} = \text{Computation} / (T_{\text{comp}} + T_{\text{comp}})$$

64K Ops

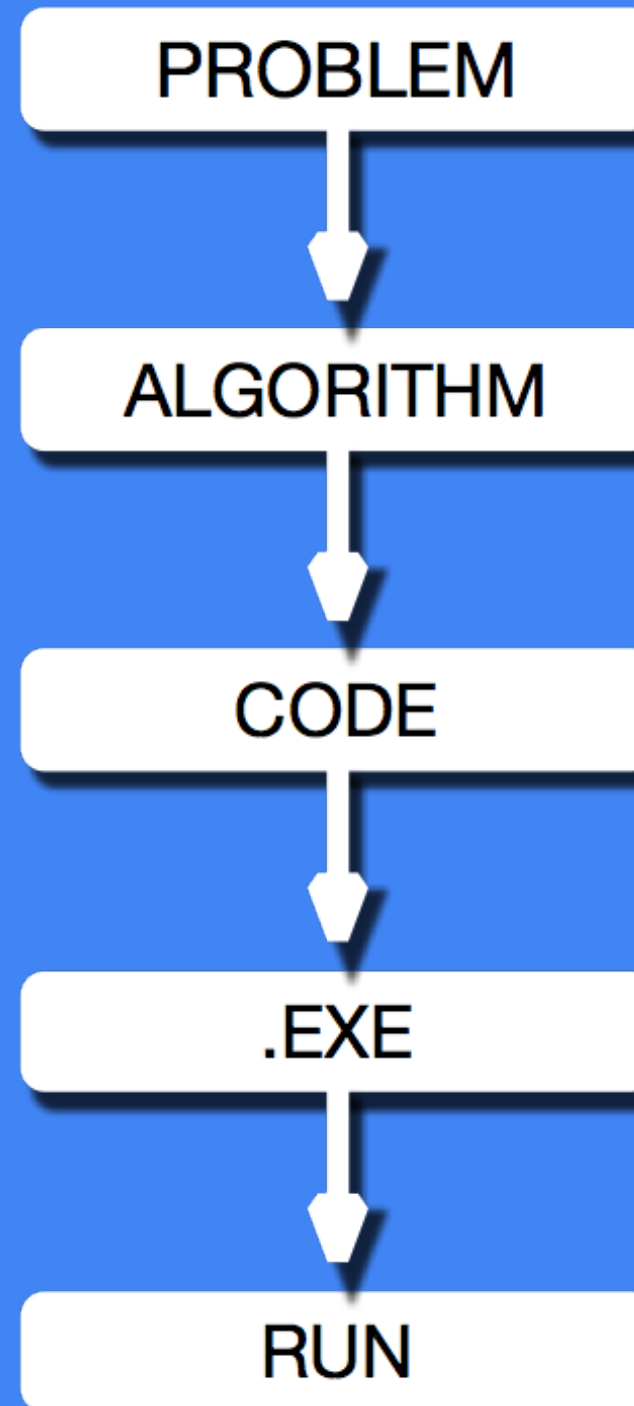
$$T_{\text{net}} \rightarrow 100\text{ns} * 1024 = \sim 100\mu\text{s}$$
$$A, B \rightarrow 100\mu\text{s each}$$

$$T_{\text{comp}} \rightarrow 64\text{Kops} @ 6 \text{ ops/cycle}$$
$$\sim 11\text{K cycles}$$
$$\rightarrow 11\text{K} * 2.25 * 10^9 = \sim 25\mu\text{s}$$

$$\text{FLOPS: } 64\text{K} / (100 + 100 + 25)\mu\text{s} = \sim 280\text{MFlops}$$

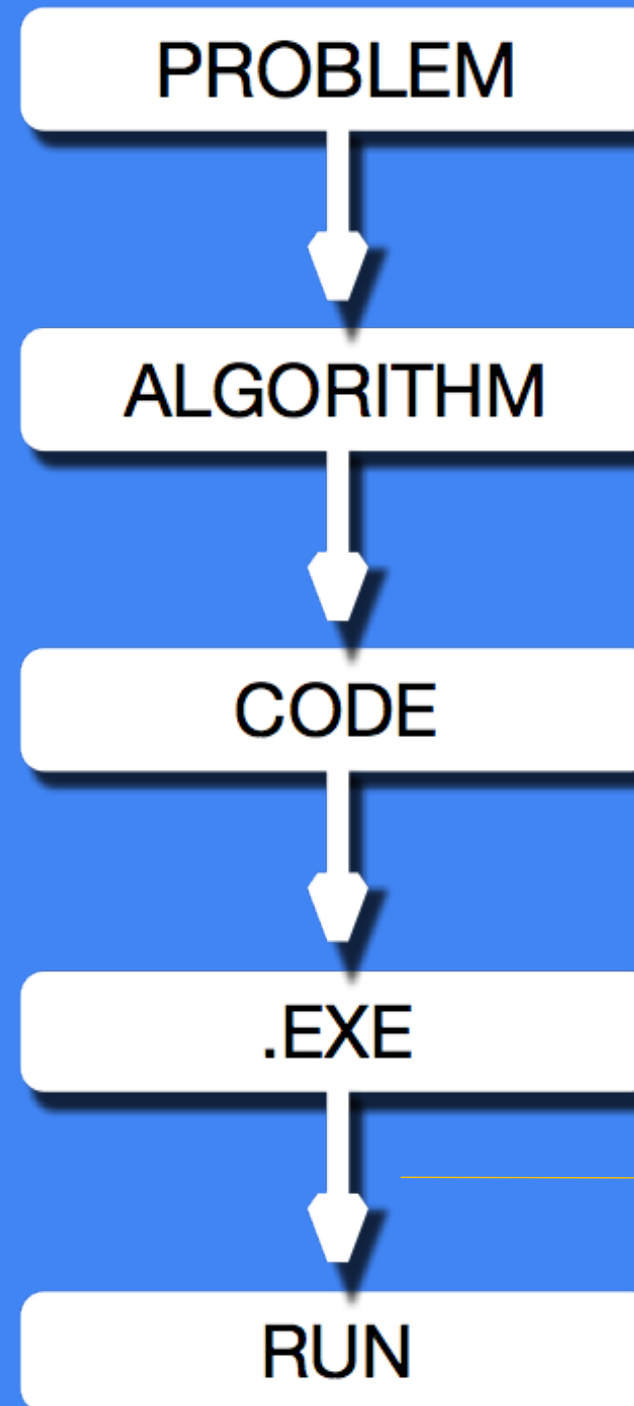
Goal?

Learn about Code
Performance and Scalability



Goal?

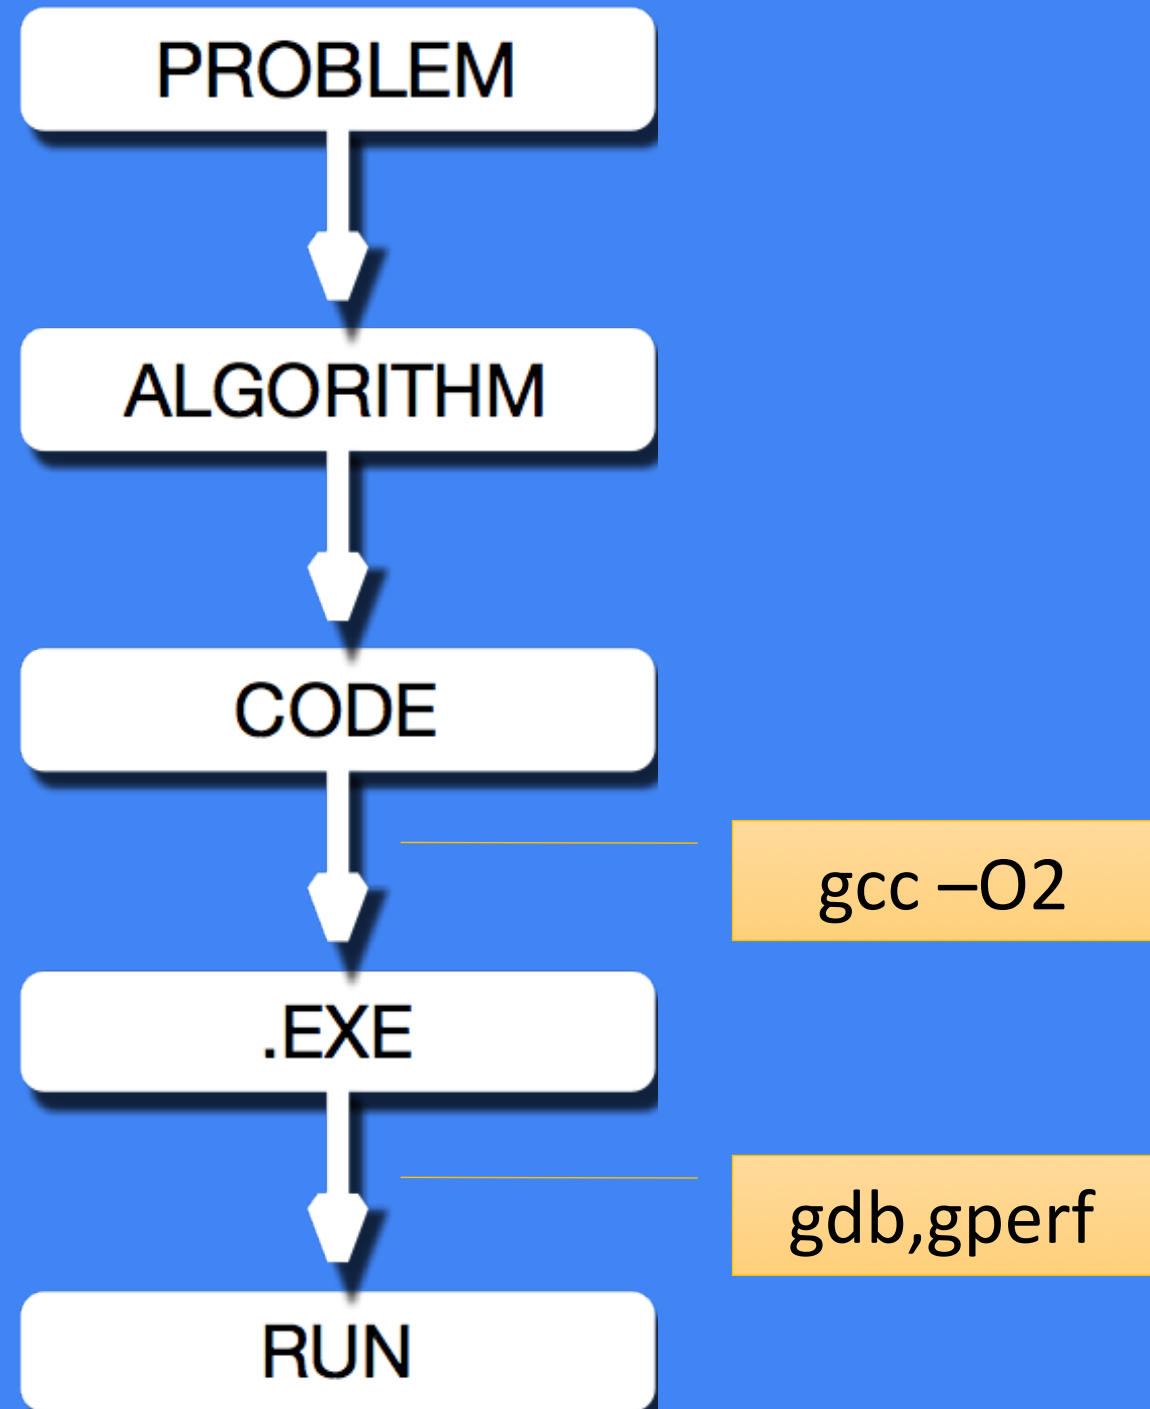
Learn about Code
Performance and Scalability



gdb,gperf

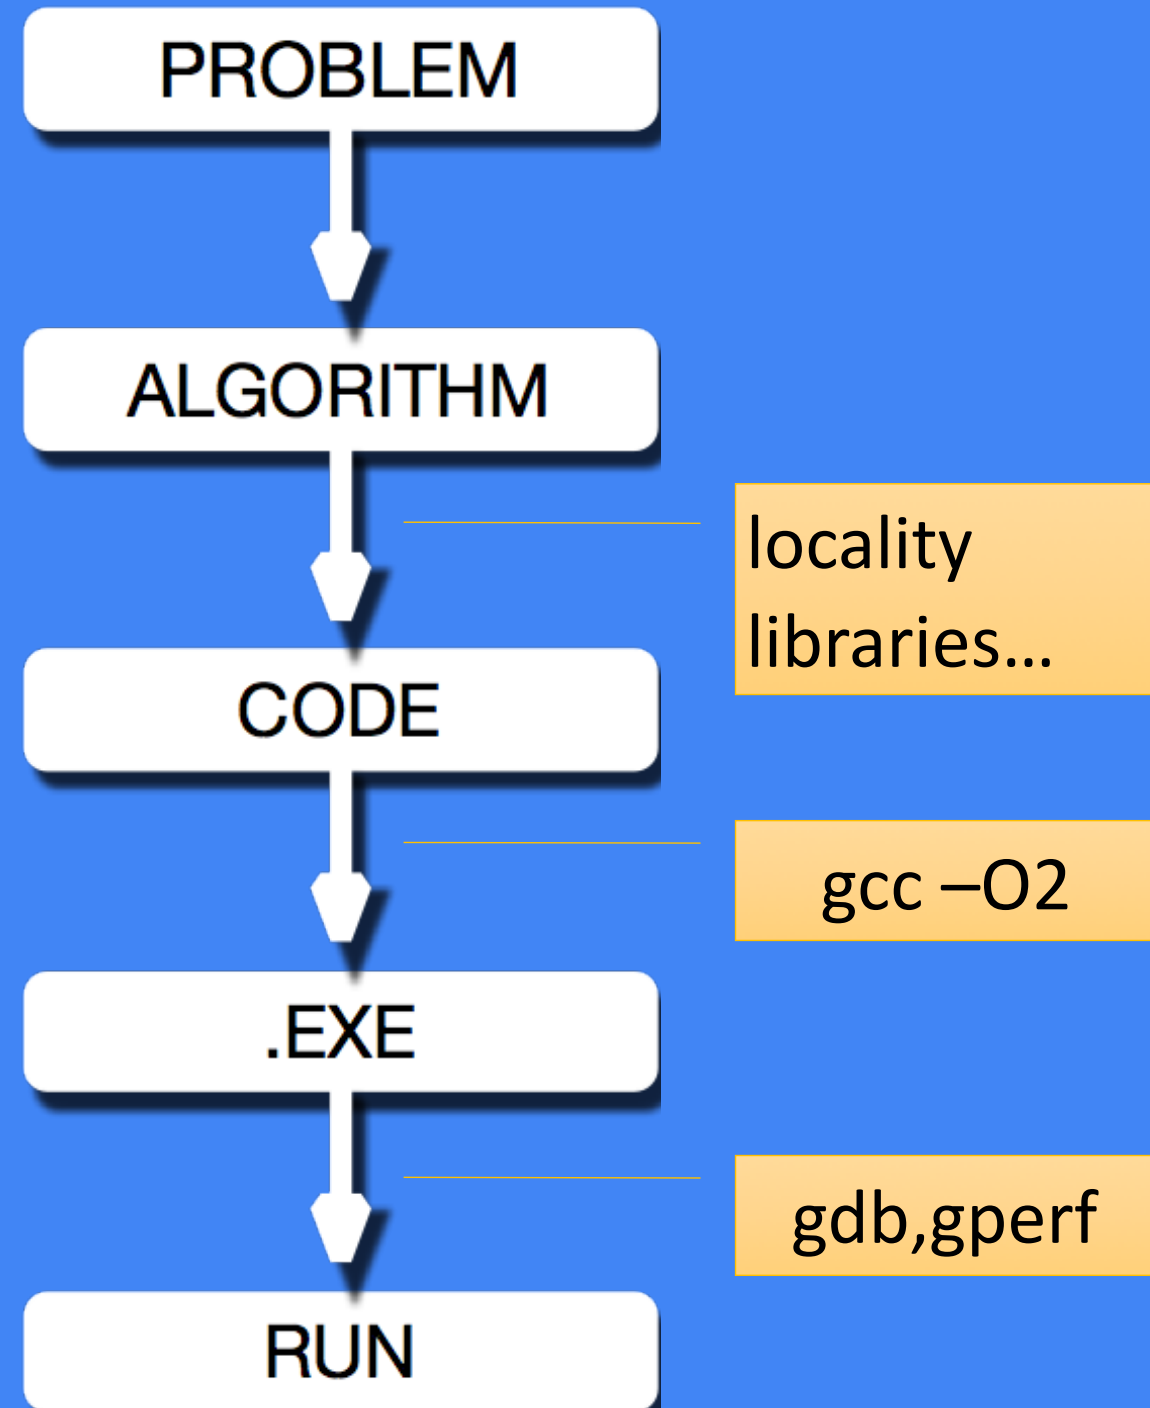
Goal?

Learn about Code
Performance and Scalability



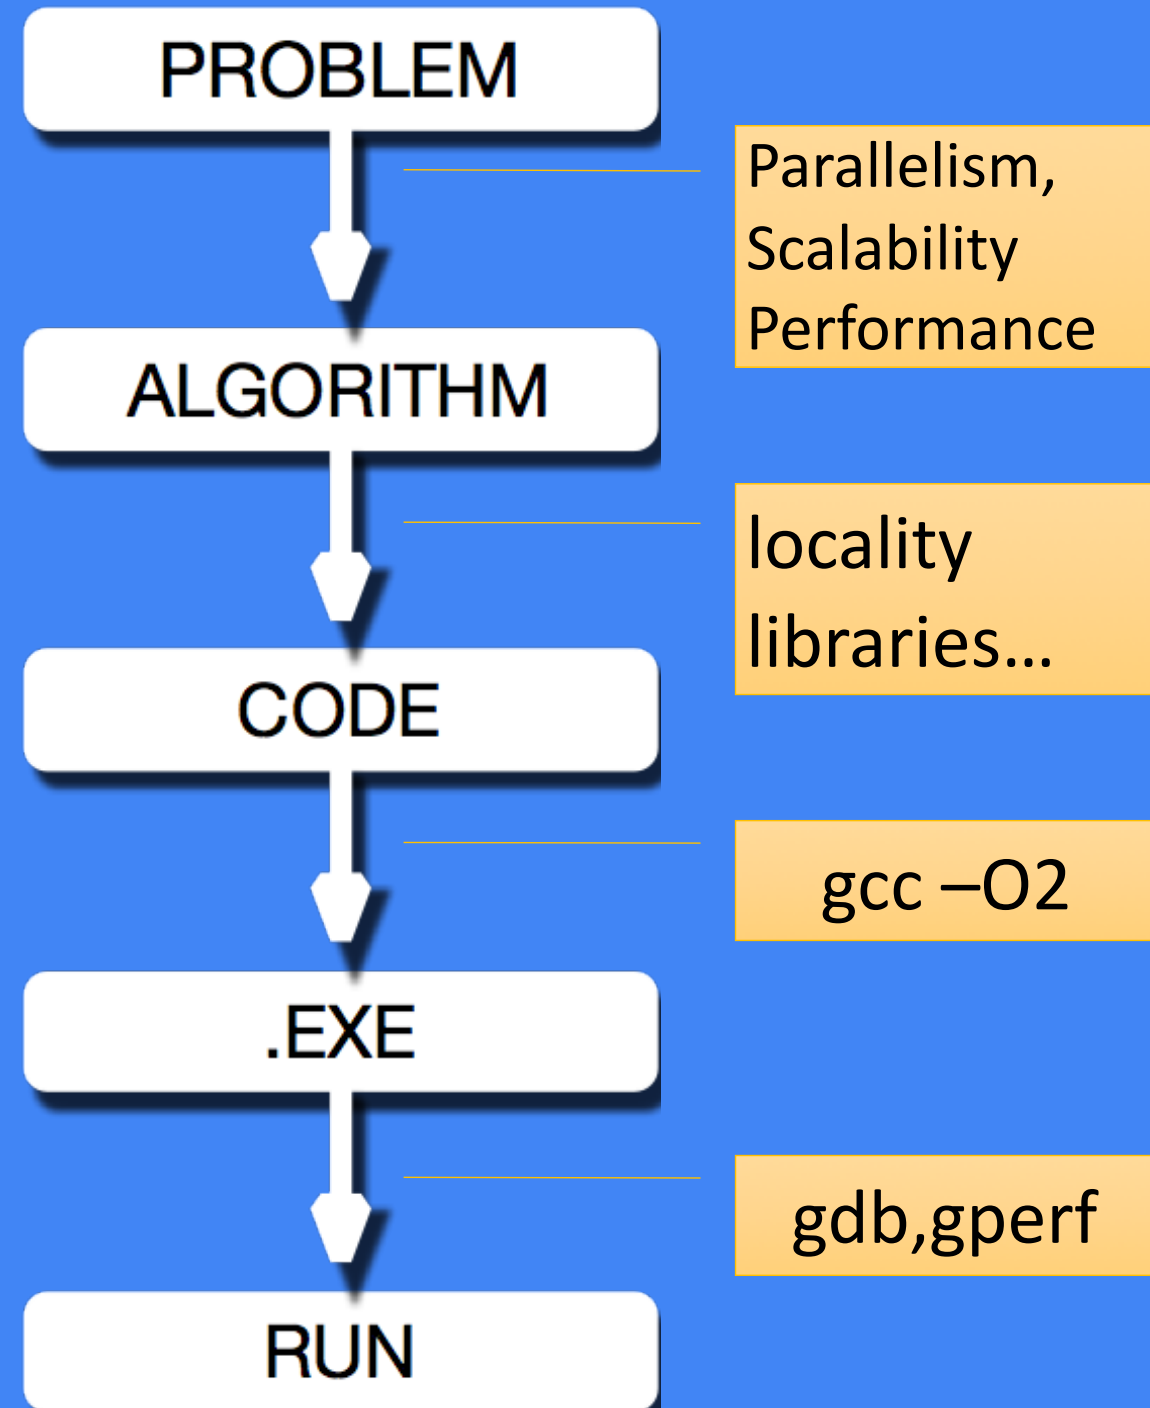
Goal?

Learn about Code
Performance and Scalability



Goal?

Learn about Code
Performance and Scalability



Step 1: Choosing an algorithm for PROBLEM

8

Algorithm → Biggest impact on efficiency and scalability

Step 1: Choosing an algorithm for PROBLEM

8

Algorithm → Biggest impact on efficiency and scalability

- It may not be intuitively obvious which algorithm is best.

Step 1: Choosing an algorithm for PROBLEM

8

Algorithm → Biggest impact on efficiency and scalability

- It may not be intuitively obvious which algorithm is best.
 - A method that is highly efficient on one or a few processors may not be optimal for a large-scale parallel system.

Step 1: Choosing an algorithm for PROBLEM

8

Algorithm → Biggest impact on efficiency and scalability

- It may not be intuitively obvious which algorithm is best.
 - A method that is highly efficient on one or a few processors may not be optimal for a large-scale parallel system.
 - REASON: Code = $\Sigma(\text{sub-algorithms, sub-tasks})$

Step 1: Choosing an algorithm for PROBLEM

8

Algorithm → Biggest impact on efficiency and scalability

- It may not be intuitively obvious which algorithm is best.
 - A method that is highly efficient on one or a few processors may not be optimal for a large-scale parallel system.
 - REASON: Code = $\Sigma(\text{sub-algorithms, sub-tasks})$
 - And these sub-algorithms/approaches may not have the best scalability for your problem, at the scales you want to work in.

Step 1: Choosing an algorithm for PROBLEM

8

Algorithm → Biggest impact on efficiency and scalability

- It may not be intuitively obvious which algorithm is best.
 - A method that is highly efficient on one or a few processors may not be optimal for a large-scale parallel system.
 - REASON: Code = $\Sigma(\text{sub-algorithms, sub-tasks})$
 - And these sub-algorithms/approaches may not have the best scalability for your problem, at the scales you want to work in.

SOLUTION: Start Early. Model Performance

So what is a **scalable** and **efficient** algorithm

So what is a **scalable** and **efficient** algorithm

9

1. Depends on the characteristics of your platform

So what is a **scalable** and **efficient** algorithm

1. Depends on the characteristics of your platform
2. Can decompose work into (sort-of) independent sub-tasks

So what is a **scalable** and **efficient** algorithm

1. Depends on the characteristics of your platform
2. Can decompose work into (sort-of) independent sub-tasks
3. Can compute >> Communicate. Why??

Memory Hierarchy

So what is a **scalable** and **efficient** algorithm

1. Depends on the characteristics of your platform
2. Can decompose work into (sort-of) independent sub-tasks
3. Can compute >> Communicate. Why??
Memory Hierarchy
4. Can minimize communication between sub-tasks

So what is a **scalable** and **efficient** algorithm

1. Depends on the characteristics of your platform
2. Can decompose work into (sort-of) independent sub-tasks
3. Can compute >> Communicate. Why??
Memory Hierarchy
4. Can minimize communication between sub-tasks
5. Minimal global communication. Why??

So what is a **scalable** and **efficient** algorithm

1. Depends on the characteristics of your platform
2. Can decompose work into (sort-of) independent sub-tasks
3. Can compute >> Communicate. Why??
Memory Hierarchy
4. Can minimize communication between sub-tasks
5. Minimal global communication. Why??
6. Initiate as many tasks as possible

SCALABILITY

SCALABILITY

10

- Common metric for measuring scalability is SPEEDUP

SCALABILITY

10

- Common metric for measuring scalability is SPEEDUP
- How much faster is your serial code vs. your parallel code

SCALABILITY

10

- Common metric for measuring scalability is SPEEDUP
- How much faster is your serial code vs. your parallel code
 - $T_s \rightarrow$ Time to complete serial version of the problem

SCALABILITY

10

- Common metric for measuring scalability is SPEEDUP
- How much faster is your serial code vs. your parallel code
 - $T_s \rightarrow$ Time to complete serial version of the problem
 - $T_p \rightarrow$ Time to complete parallel version of the problem

SCALABILITY

10

- Common metric for measuring scalability is SPEEDUP
- How much faster is your serial code vs. your parallel code
 - $T_s \rightarrow$ Time to complete serial version of the problem
 - $T_p \rightarrow$ Time to complete parallel version of the problem
- $\text{Speedup} = T_s / T_p$

SCALABILITY

10

- Common metric for measuring scalability is SPEEDUP
- How much faster is your serial code vs. your parallel code
 - $T_s \rightarrow$ Time to complete serial version of the problem
 - $T_p \rightarrow$ Time to complete parallel version of the problem
- $\text{Speedup} = T_s / T_p$

SCALABILITY

10

- Common metric for measuring scalability is SPEEDUP
- How much faster is your serial code vs. your parallel code
 - $T_s \rightarrow$ Time to complete serial version of the problem
 - $T_p \rightarrow$ Time to complete parallel version of the problem
 - $\text{Speedup} = T_s / T_p$

Say 10 processors, finish in $1/10^{\text{th}}$ the time. $\text{Speedup} = 10 \rightarrow$ LINEAR

SCALABILITY

10

- Common metric for measuring scalability is SPEEDUP
- How much faster is your serial code vs. your parallel code
 - $T_s \rightarrow$ Time to complete serial version of the problem
 - $T_p \rightarrow$ Time to complete parallel version of the problem
- $\text{Speedup} = T_s / T_p$

Say 10 processors, finish in $1/10^{\text{th}}$ the time. $\text{Speedup} = 10 \rightarrow$ LINEAR
SUPERLINEAR $\rightarrow T_p > 1/10$

Why Parallelize?

Why Parallelize?

11

- One reason → DO IT FASTER

Why Parallelize?

11

- One reason → DO IT FASTER
- Remember, Scalability is a measure of Speedup

Why Parallelize?

11

- One reason → DO IT FASTER
- Remember, Scalability is a measure of Speedup
- One way to define scalability, using speedup is as follows:

Why Parallelize?

11

- One reason → DO IT FASTER
- Remember, Scalability is a measure of Speedup
- One way to define scalability, using speedup is as follows:
Fix the **problem size** and Increase **number of processors**

Why Parallelize?

11

- One reason → DO IT FASTER
- Remember, Scalability is a measure of Speedup
- One way to define scalability, using speedup is as follows:
Fix the **problem size** and Increase **number of processors**

STRONG SCALING

Why Parallelize?

11

- One reason → DO IT FASTER
- Remember, Scalability is a measure of Speedup
- One way to define scalability, using speedup is as follows:
Fix the **problem size** and Increase **number of processors**

STRONG SCALING

Problem size is 123456789.

Why Parallelize?

11

- One reason → DO IT FASTER
- Remember, Scalability is a measure of Speedup
- One way to define scalability, using speedup is as follows:
Fix the **problem size** and Increase **number of processors**

STRONG SCALING

Problem size is 123456789.

Takes 100s(1p) →

Why Parallelize?

11

- One reason → DO IT FASTER
- Remember, Scalability is a measure of Speedup
- One way to define scalability, using speedup is as follows:
Fix the **problem size** and Increase **number of processors**

STRONG SCALING

Problem size is 123456789.

Takes 100s(1p) →

10s(10p) →

Why Parallelize?

11

- One reason → DO IT FASTER
- Remember, Scalability is a measure of Speedup
- One way to define scalability, using speedup is as follows:
Fix the **problem size** and Increase **number of processors**

STRONG SCALING

Problem size is 123456789.

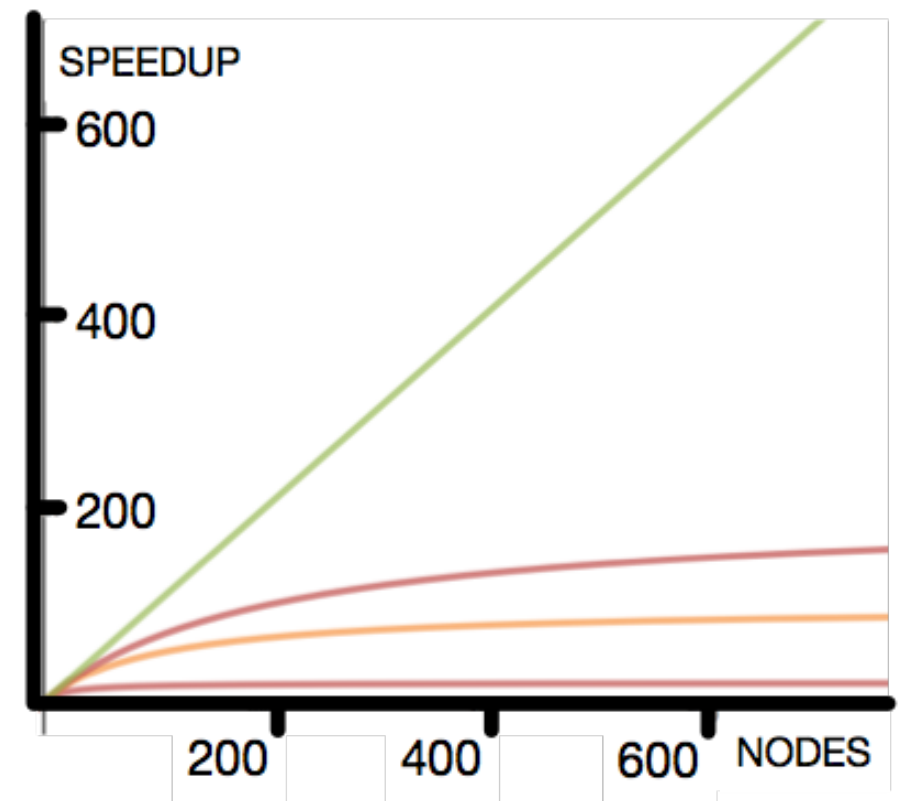
Takes 100s(1p) →

10s(10p) →

1s(100p)

Amdahl's Law

12



<https://goo.gl/cpOaoc>

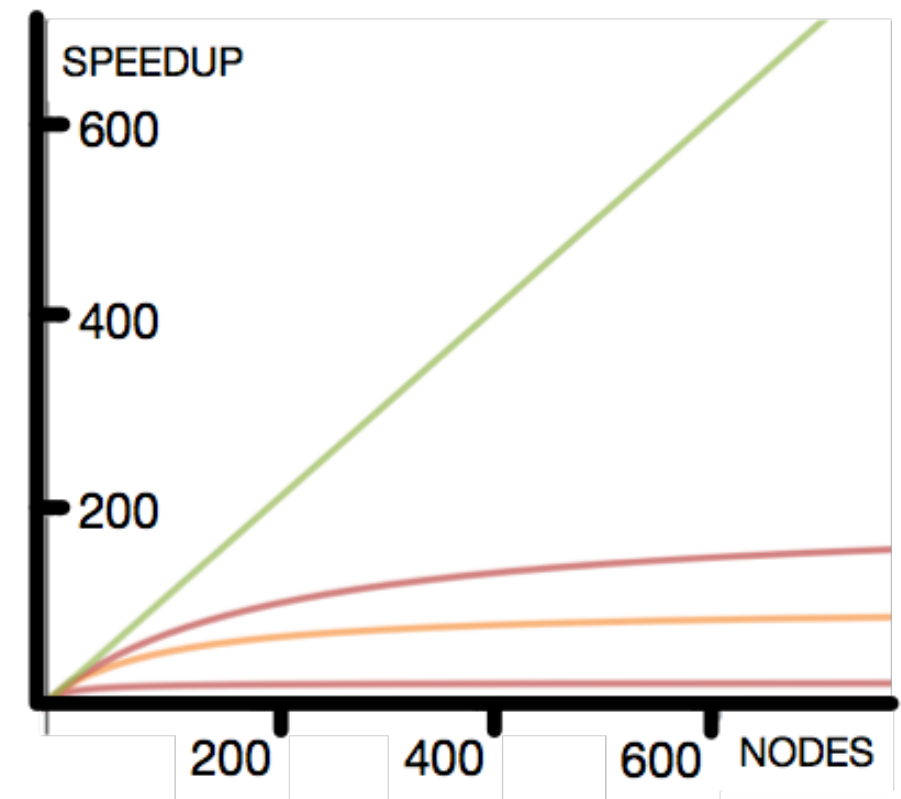
Amdahl's Law

12

$$S = (T_S / T_P)$$

$$S = \frac{T_S}{(1-P) * T_S + P * T_S / N}$$

$$S = \frac{1}{(1 - P) + P/N}$$



<https://goo.gl/cpOaoc>

Amdahl's Law

12

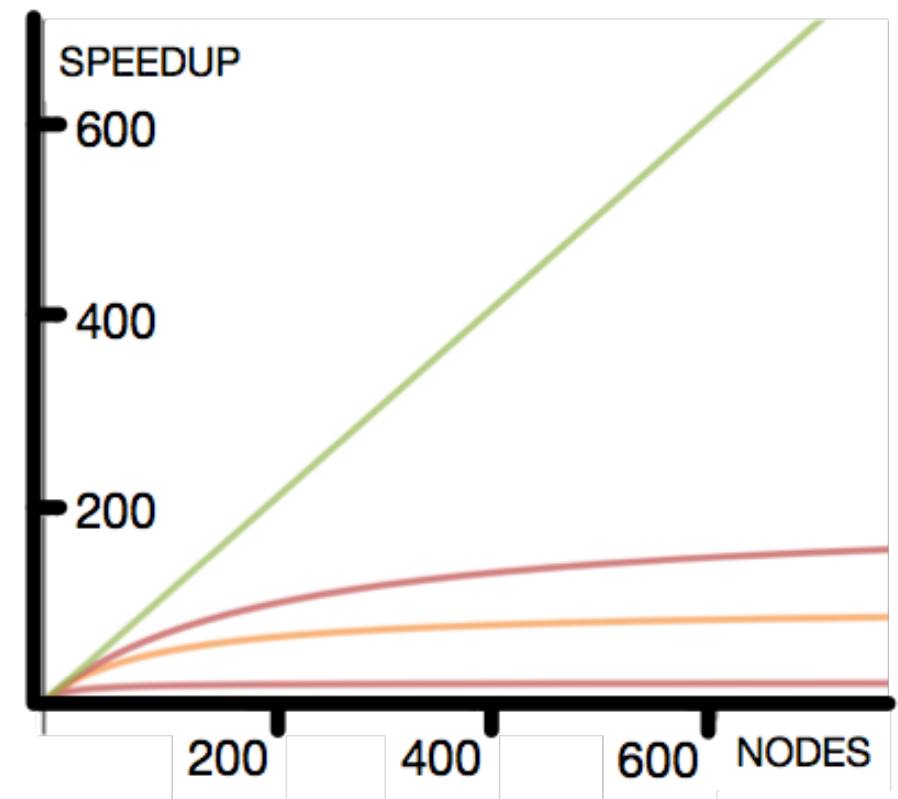
$$S = (T_S / T_P)$$

$$S = \frac{T_S}{(1-P) * T_S + P * T_S / N}$$

$$S = \frac{1}{(1 - P) + P/N}$$

Adding more cores does not always work.

Unless your code is 100% efficient at using multiple cores, you will receive less and less of a benefit by adding more cores



<https://goo.gl/cpOaoc>

Faster is hard. So let's do LARGER

13

- Scalability does not have to mean faster. It can mean → process larger
- **Aka Gustafson's Law** → *"an observation that Amdahl's assumptions don't match the way people use parallel processors. People scale their problems to match the power available, in contrast to Amdahl's assumption that the problem is always the same no matter how capable the computer"*
- Called **Weak Scaling**
$$S = (1 - P) + P * N$$

Suggested Readings

“Estimating CPU Performance using Amdahl's Law”

<https://goo.gl/Dg0Mc9>

“Amdahl's Law, Gustafson's Trend, and the Performance Limits of Parallel Applications”

<https://goo.gl/mBcvR9>

<https://goo.gl/4P28Xj>
(5mins)

What does it mean to communicate?

16

- Not a shouting match (~~All to All at the Same Time~~)
- More like some nodes to some other nodes
 - Node \rightarrow Processing Core
 - Wire \rightarrow Communication Channel
- Cost of sending a single message = $T_{\text{setup}} + T_{\text{transmission}}$

$$\rightarrow T_s + T_w \times \text{Size of data}$$

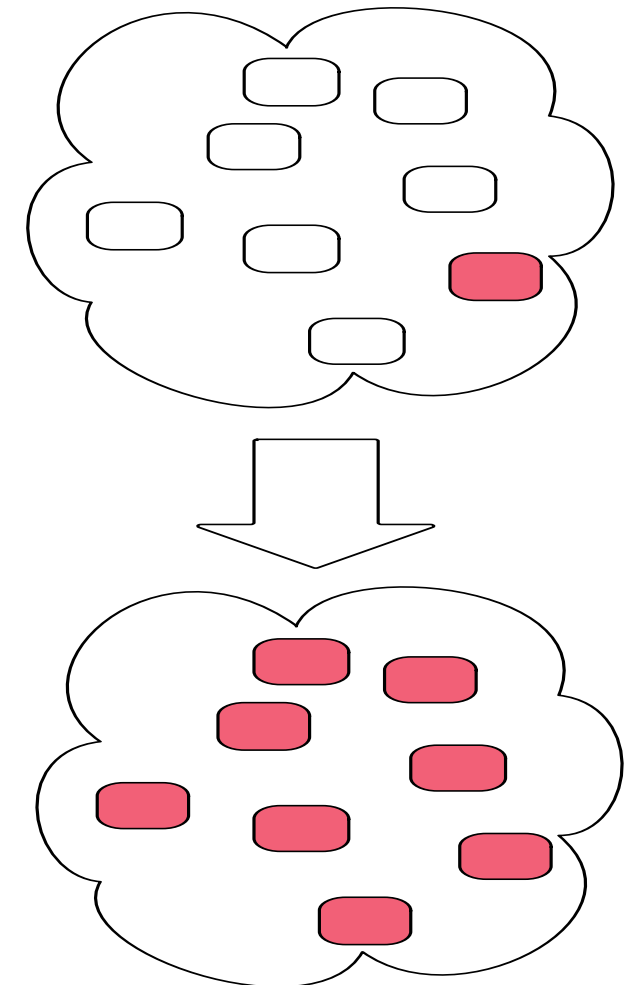
$$\rightarrow T_s + T_w \times m$$

$$\text{Total Cost} = (T_s + T_w \times m) \times \text{number of messages}$$

Strategies to communicate

17

Objective: One node will send to All other nodes
Commonly called **One-to-All Broadcast**

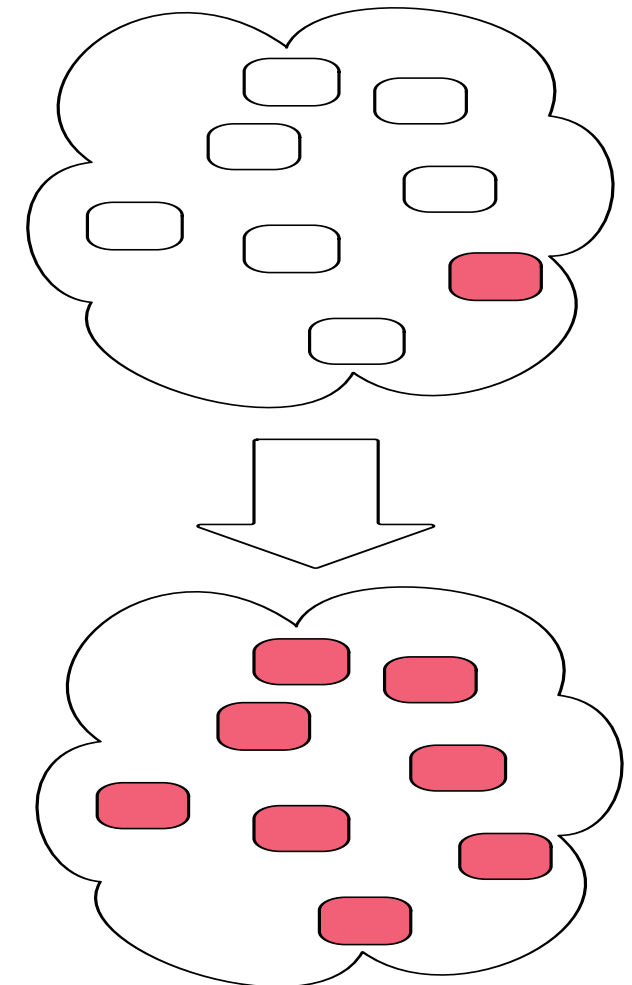


Strategies to communicate

17

Objective: One node will send to All other nodes
Commonly called **One-to-All Broadcast**

Simple Strategy: One node sends to each of the other $(p-1)$ nodes



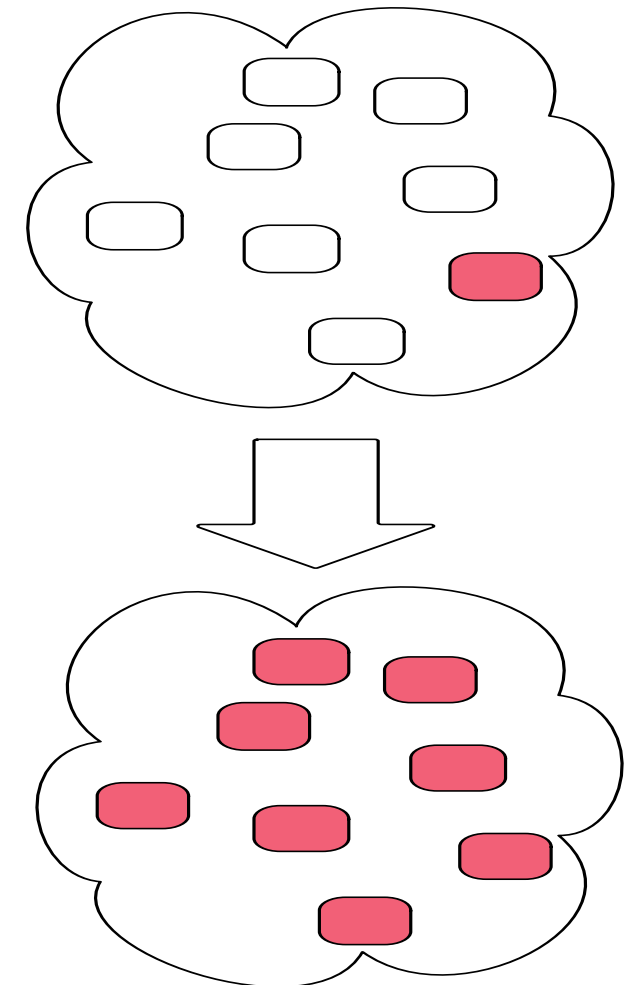
Strategies to communicate

17

Objective: One node will send to All other nodes
Commonly called **One-to-All Broadcast**

Simple Strategy: One node sends to each of the other $(p-1)$ nodes

Why is this a bad idea?



Strategies to communicate

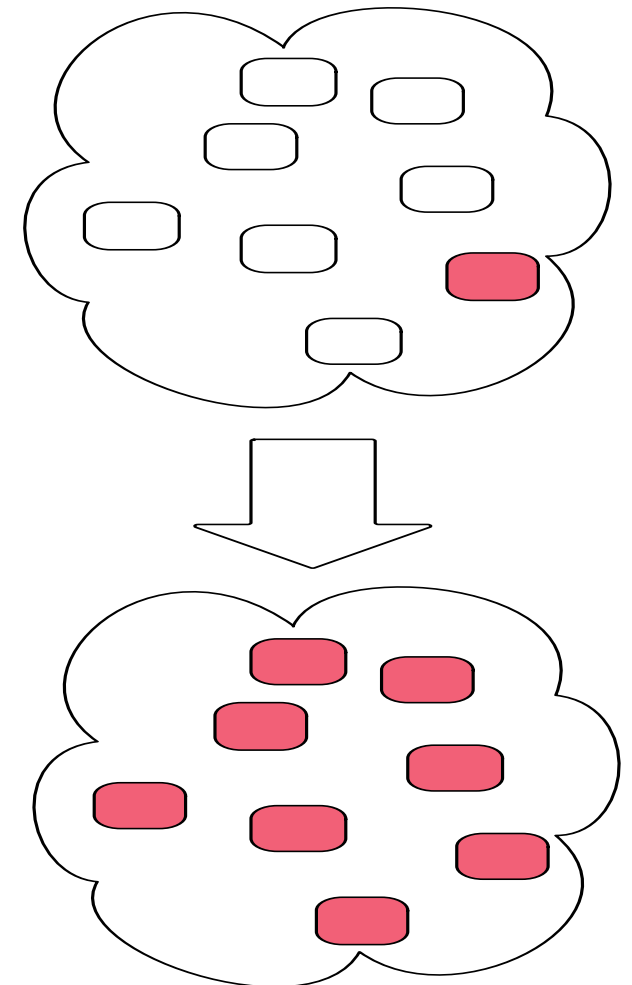
17

Objective: One node will send to All other nodes
Commonly called **One-to-All Broadcast**

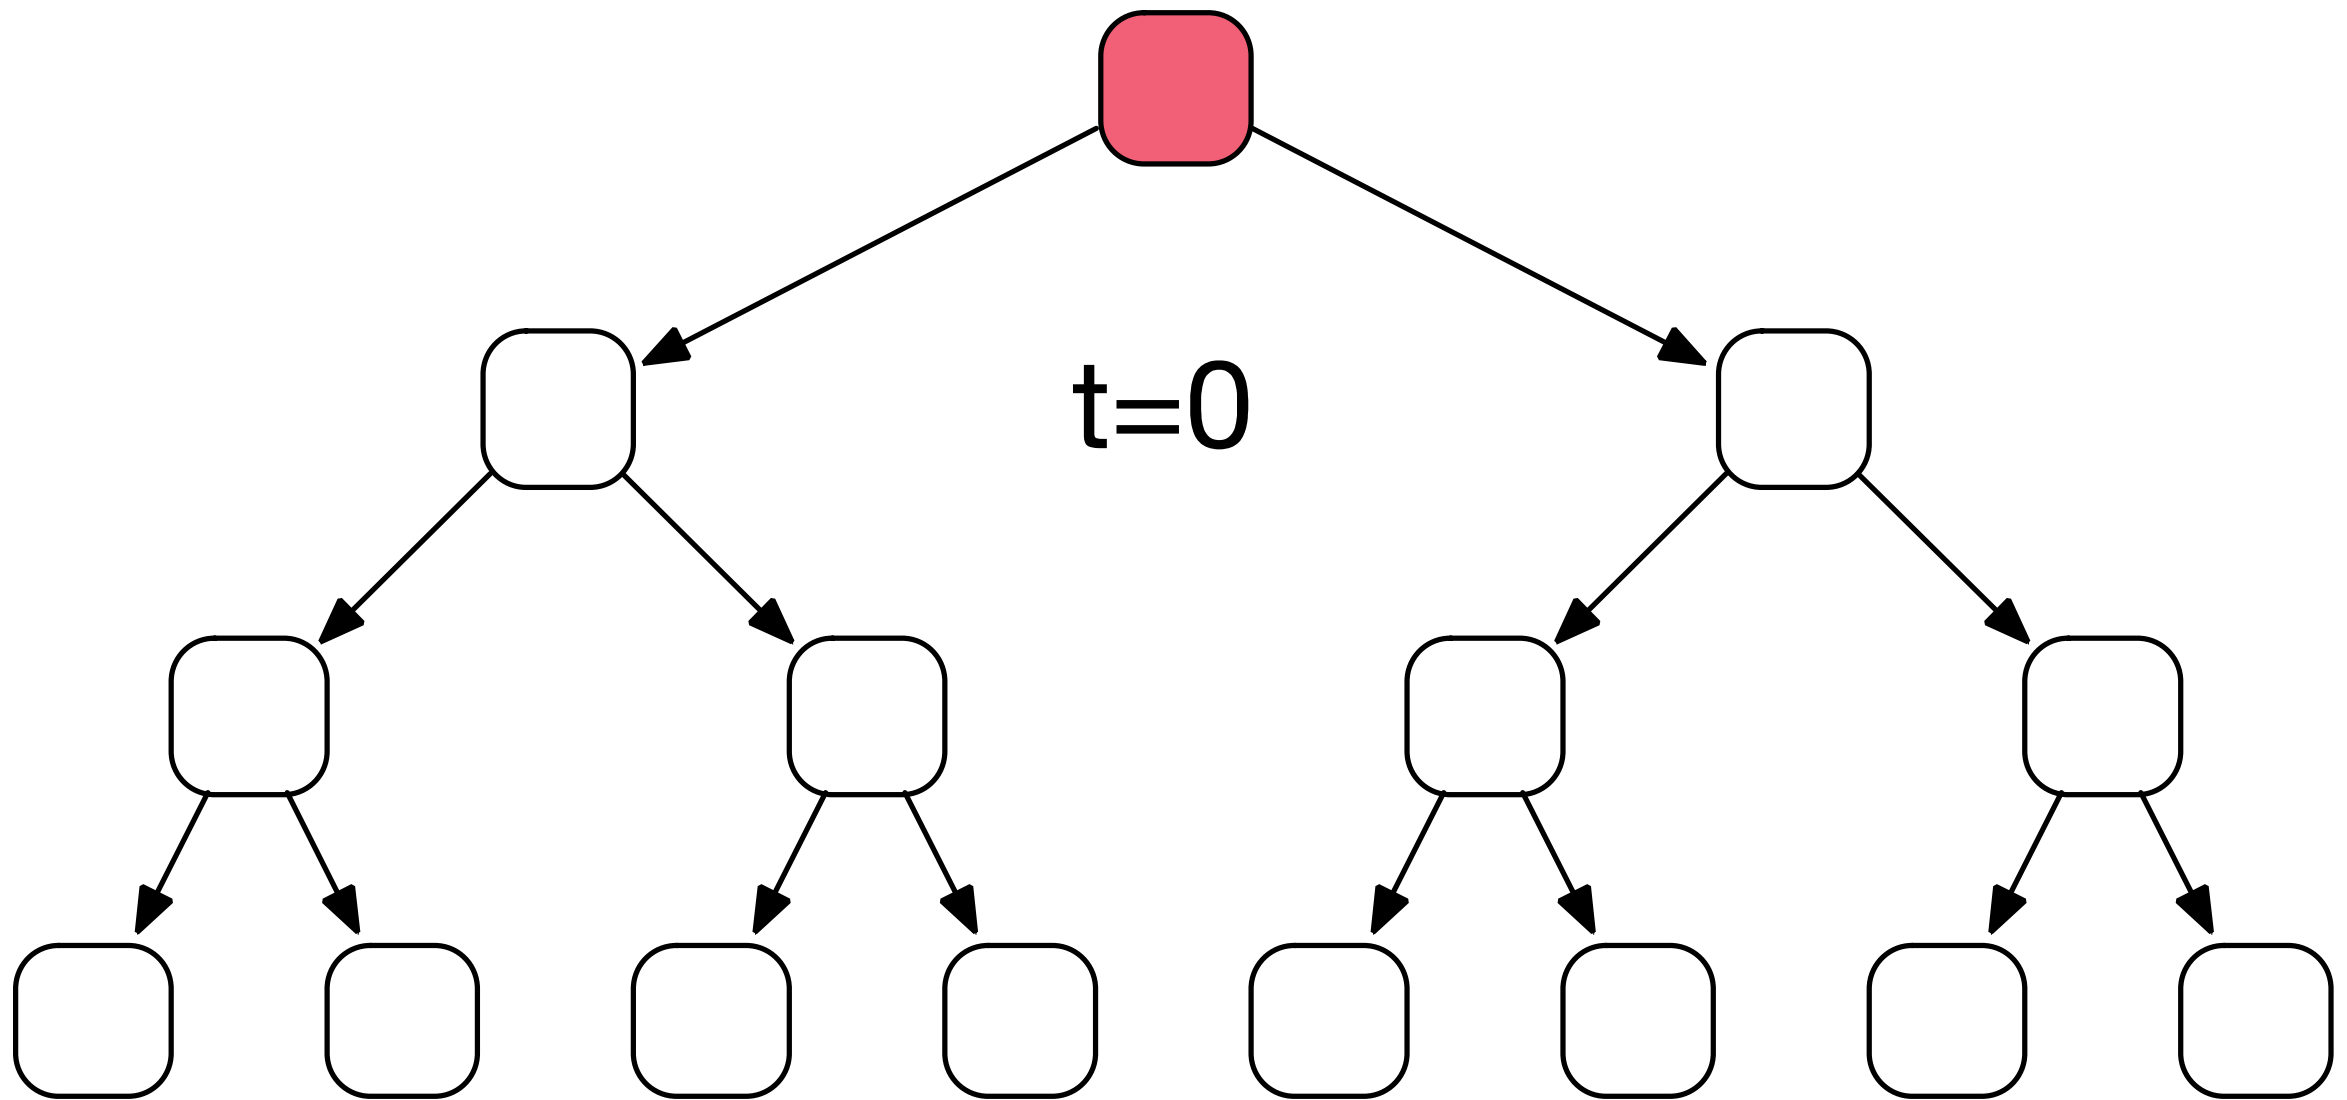
Simple Strategy: One node sends to each of the other $(p-1)$ nodes

Why is this a bad idea?

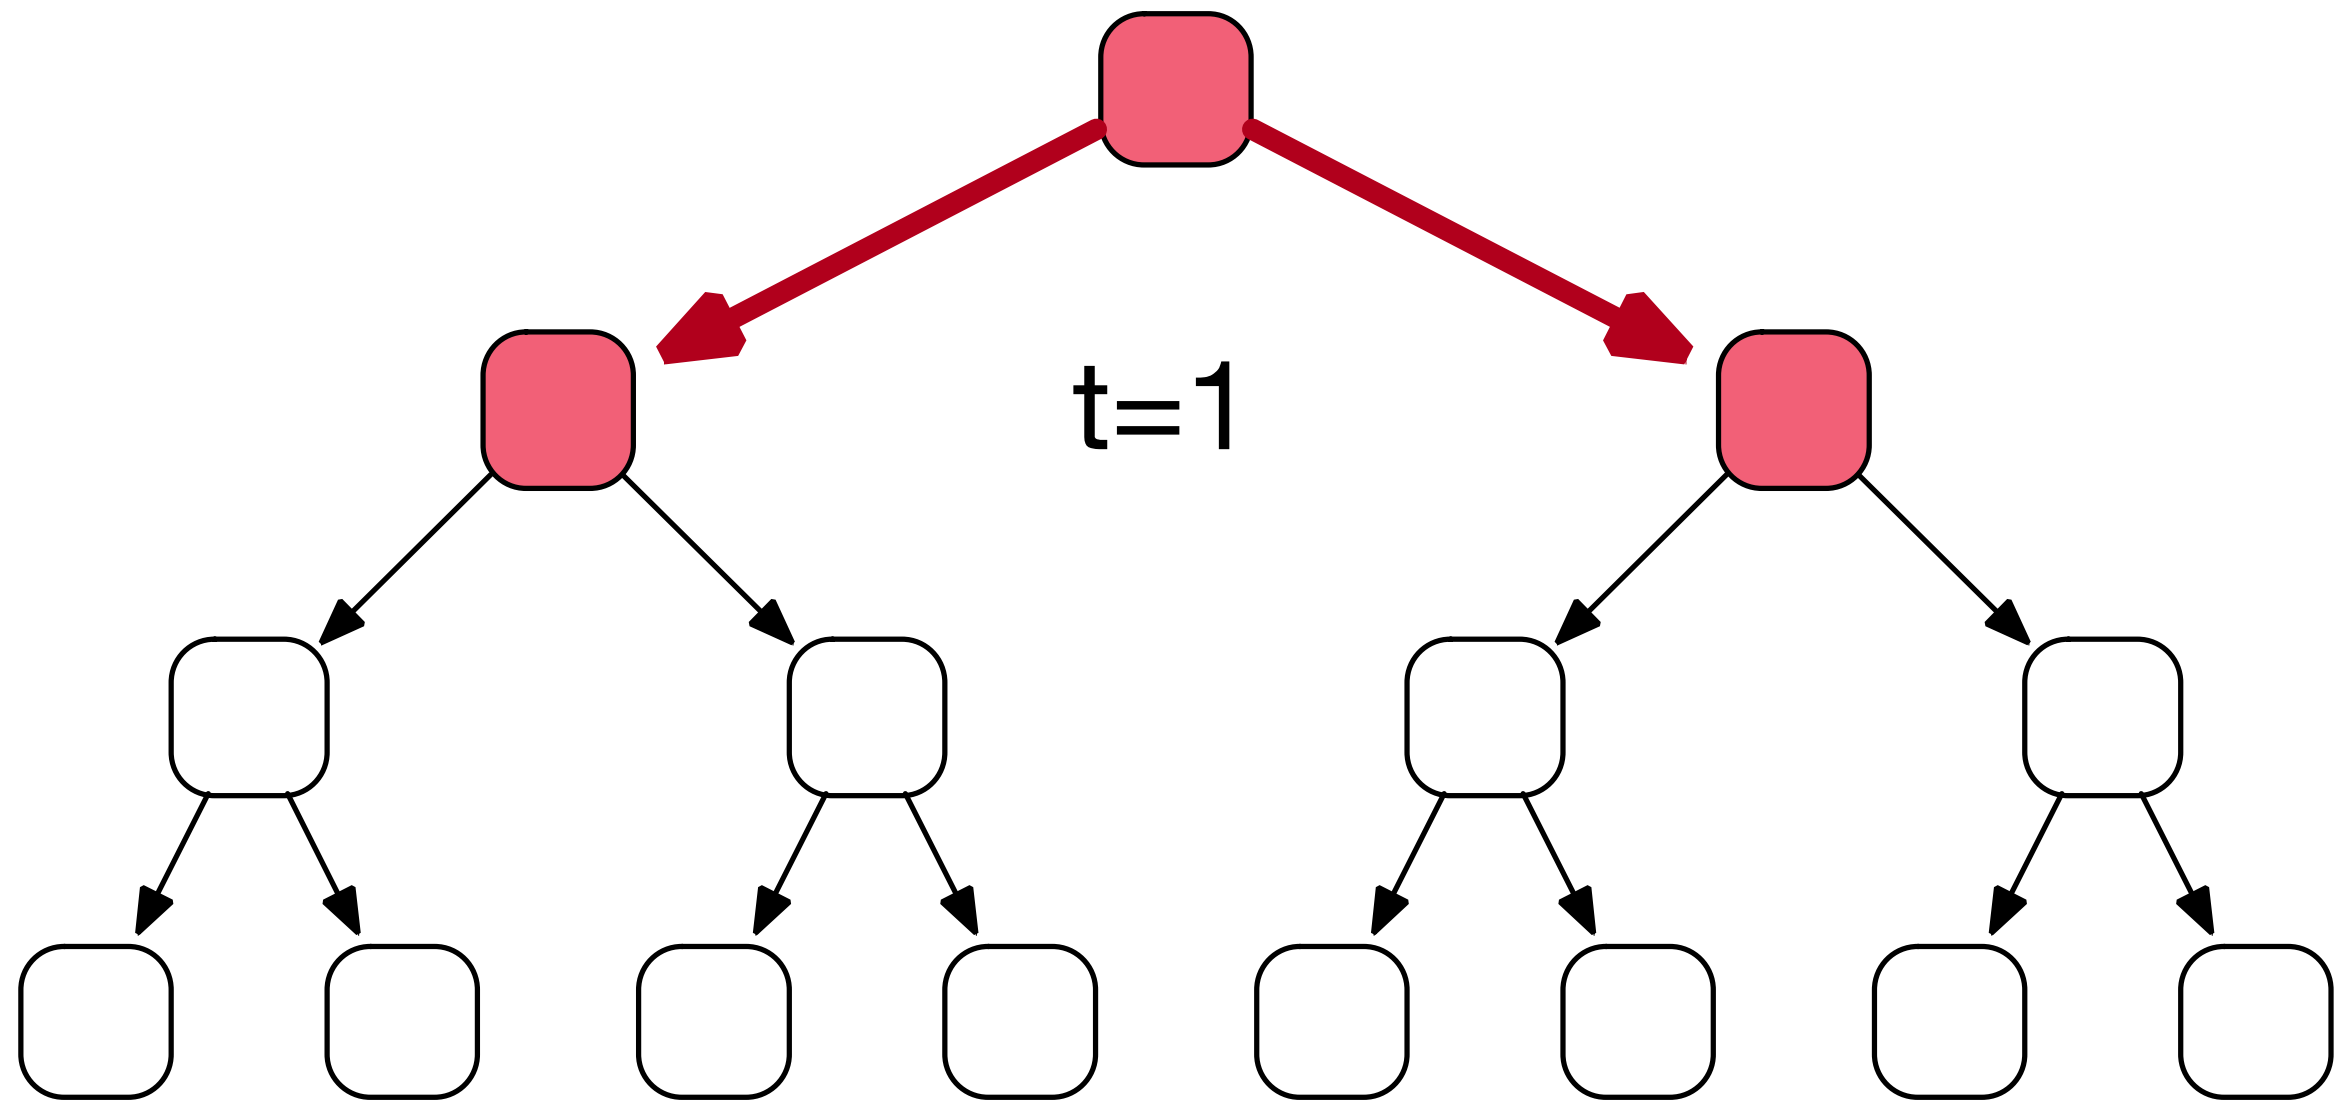
You are using a fraction of the network
i.e., you are only using one node at a time



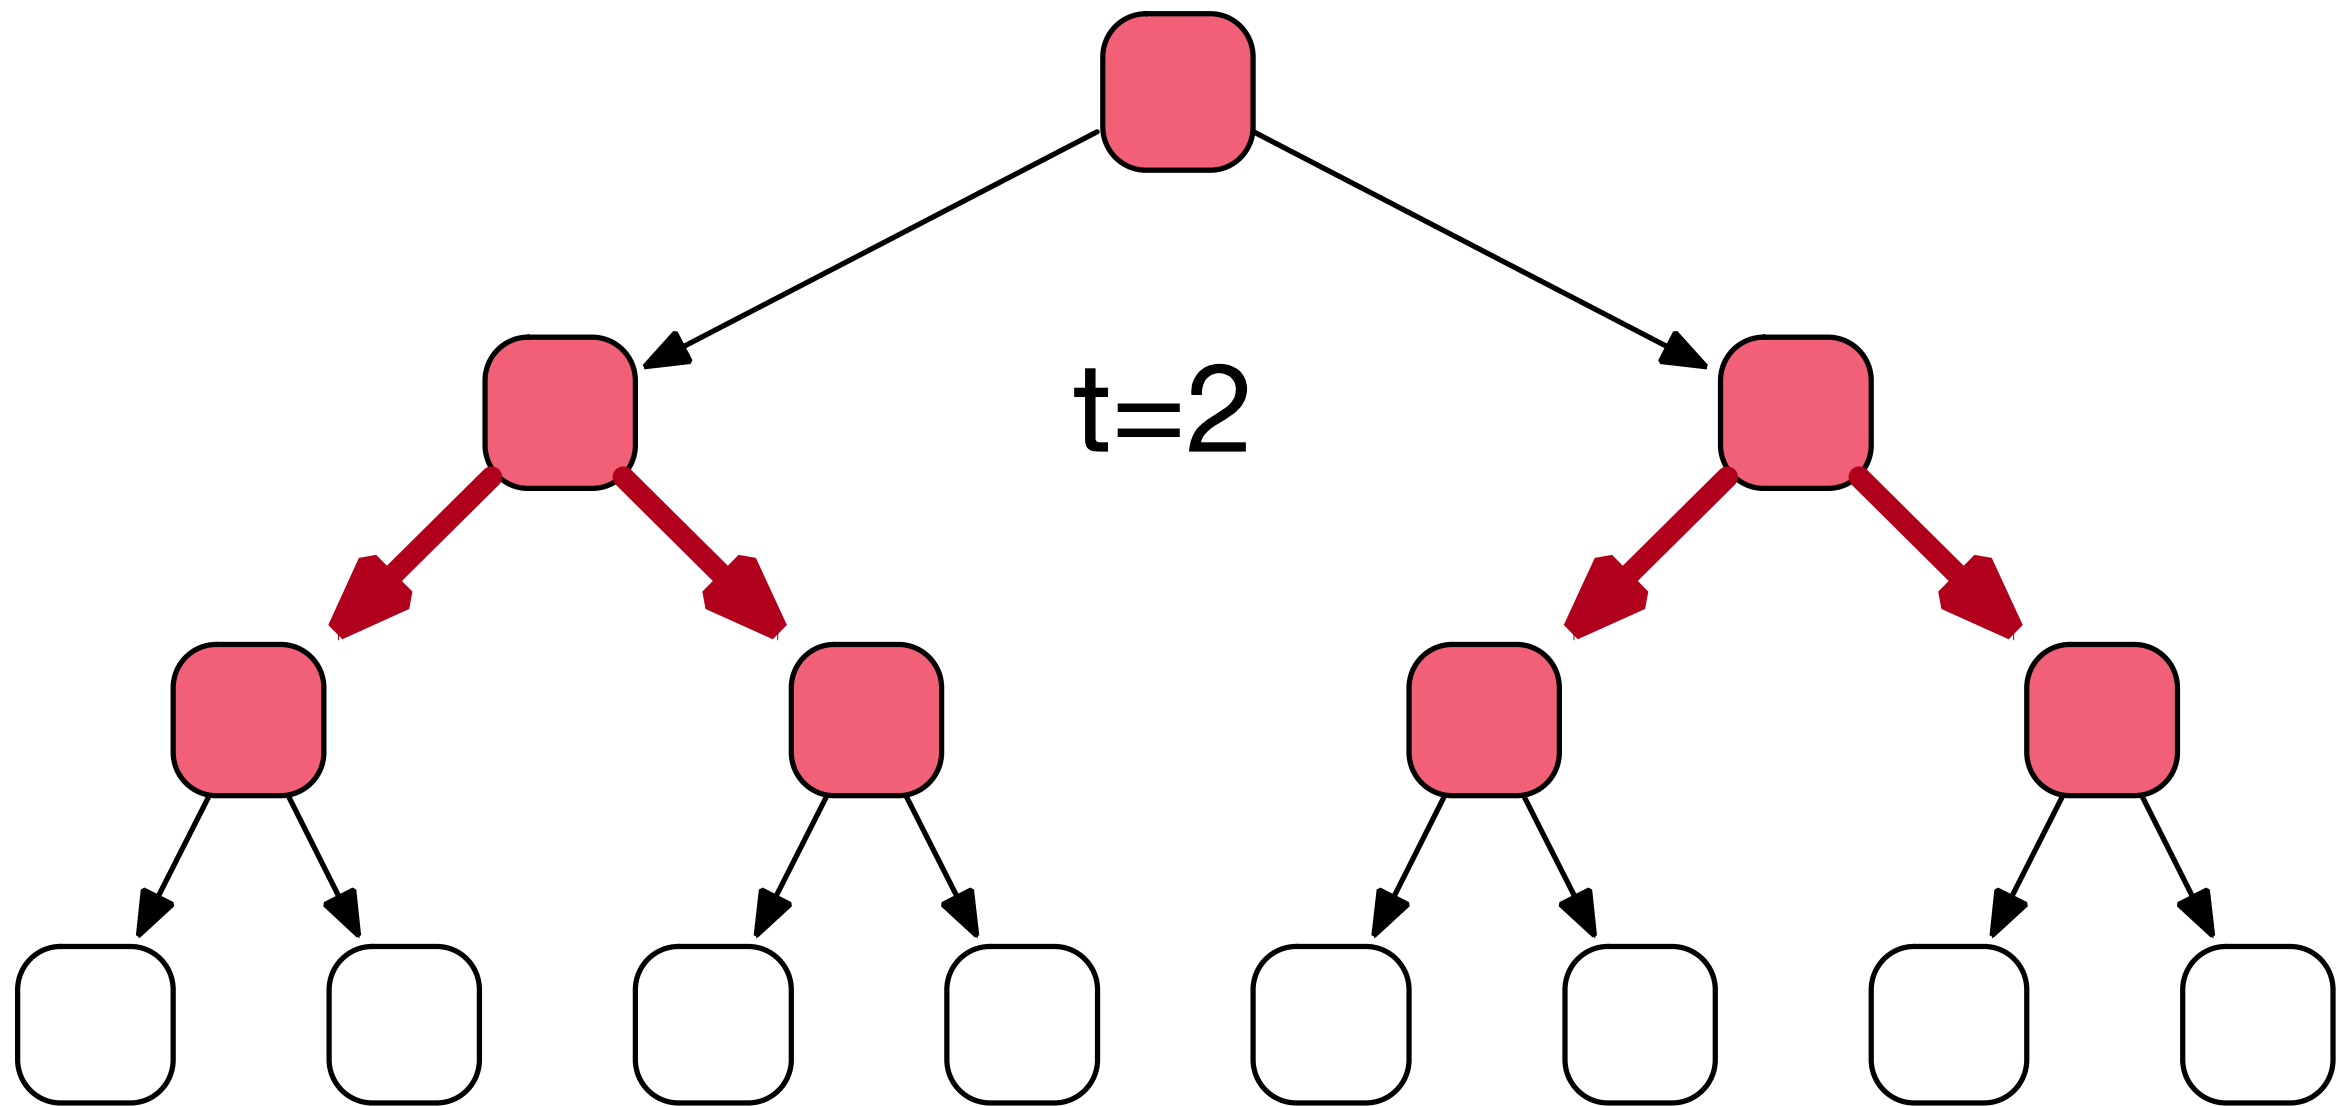
Better Strategy → Recursive Doubling



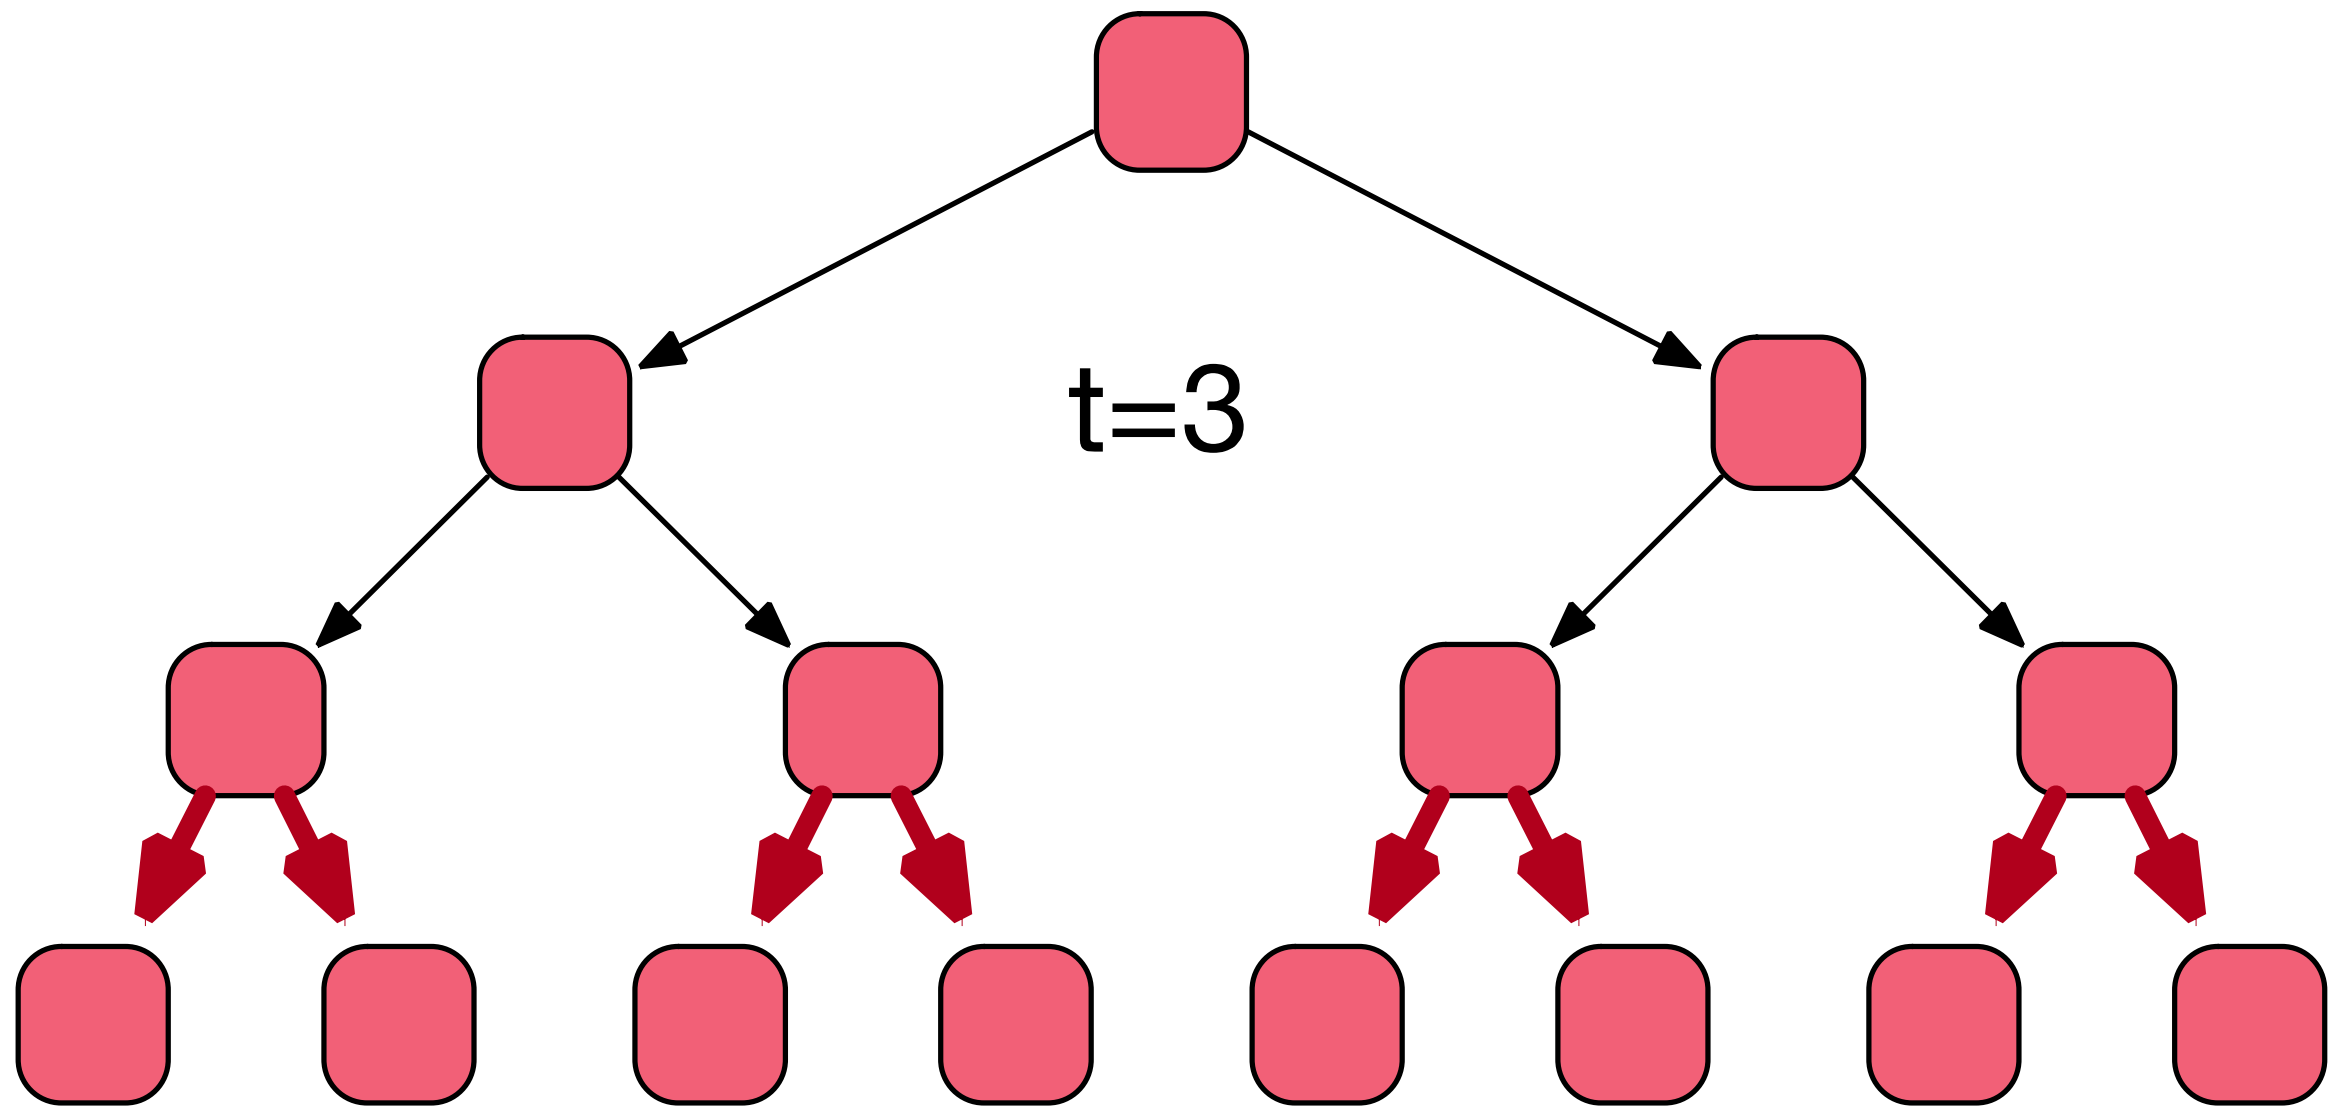
Better Strategy → Recursive Doubling



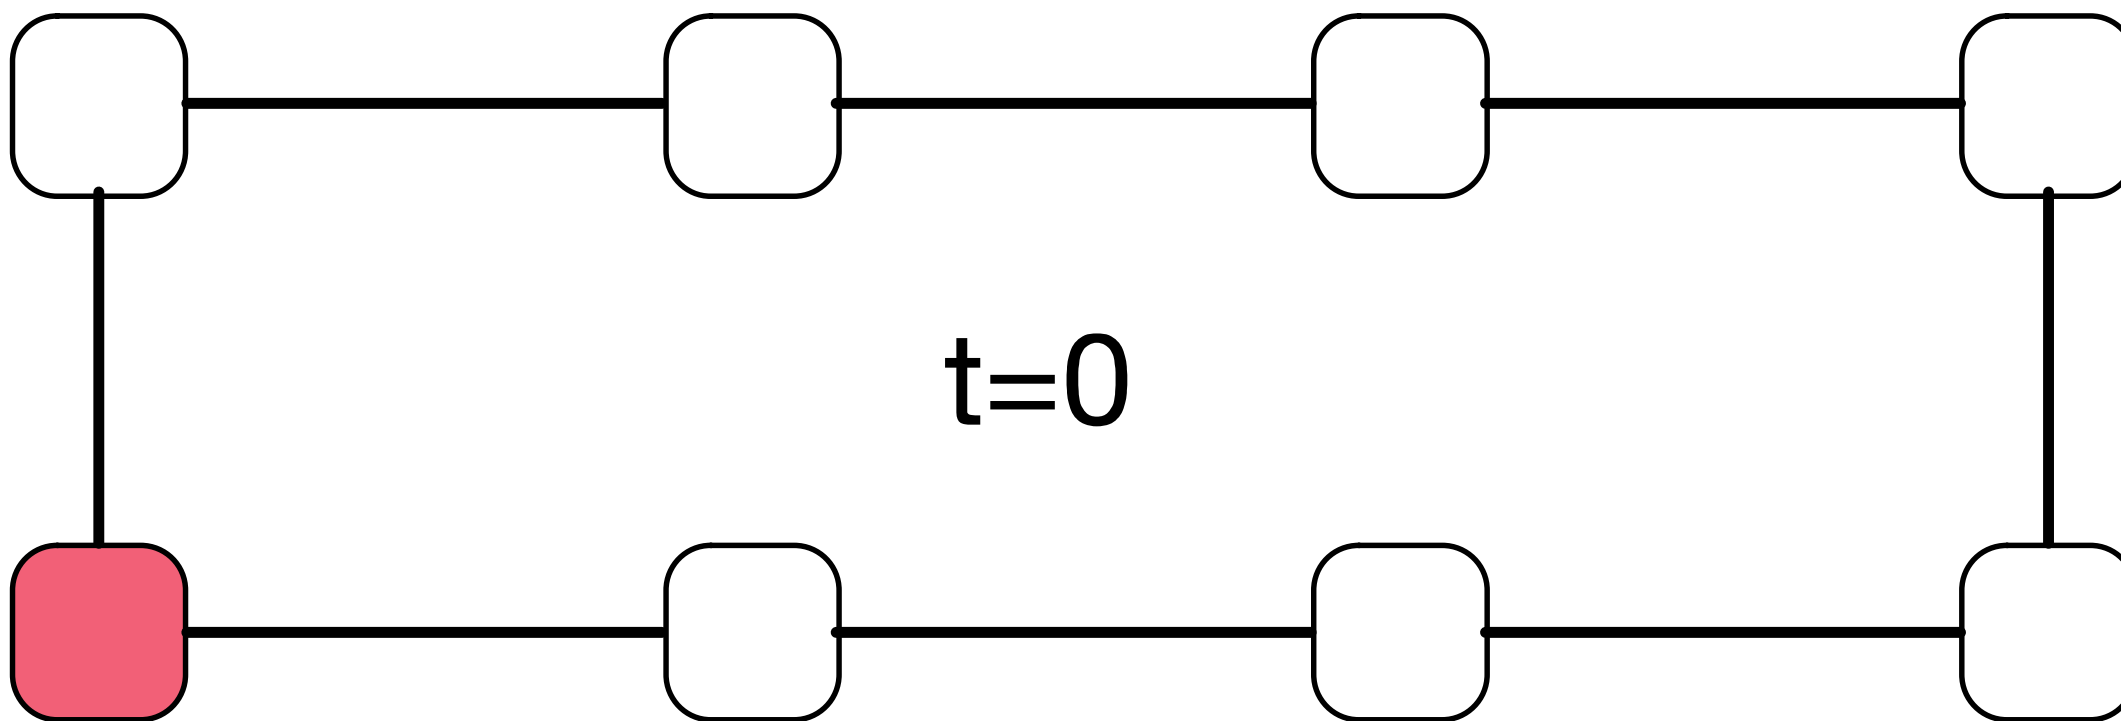
Better Strategy → Recursive Doubling

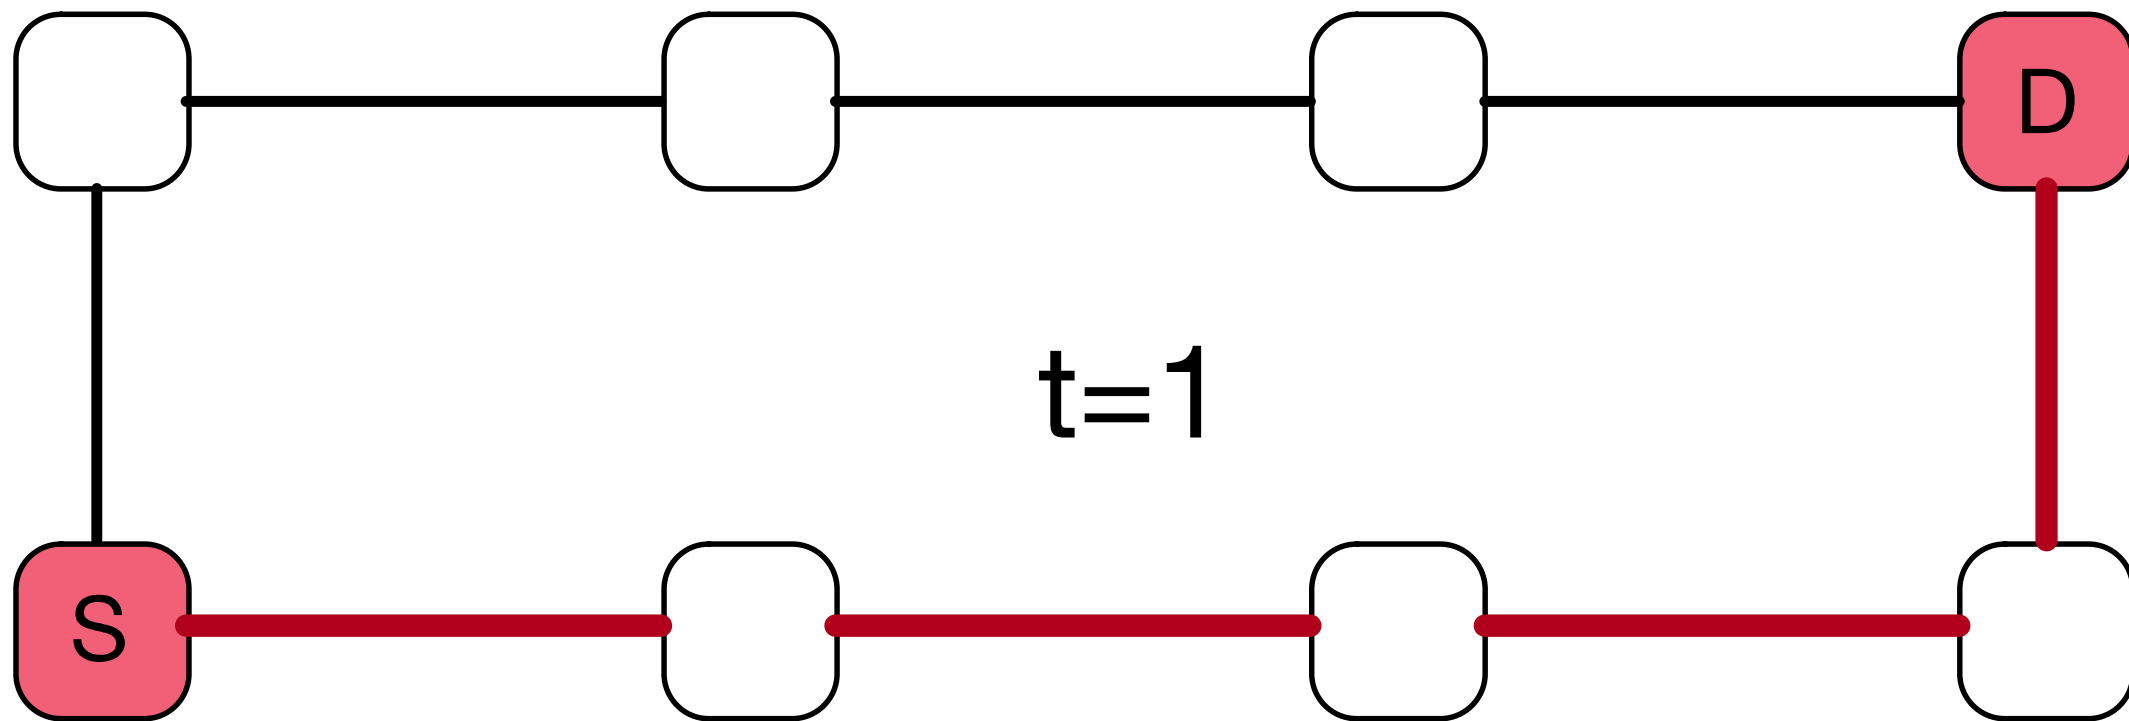


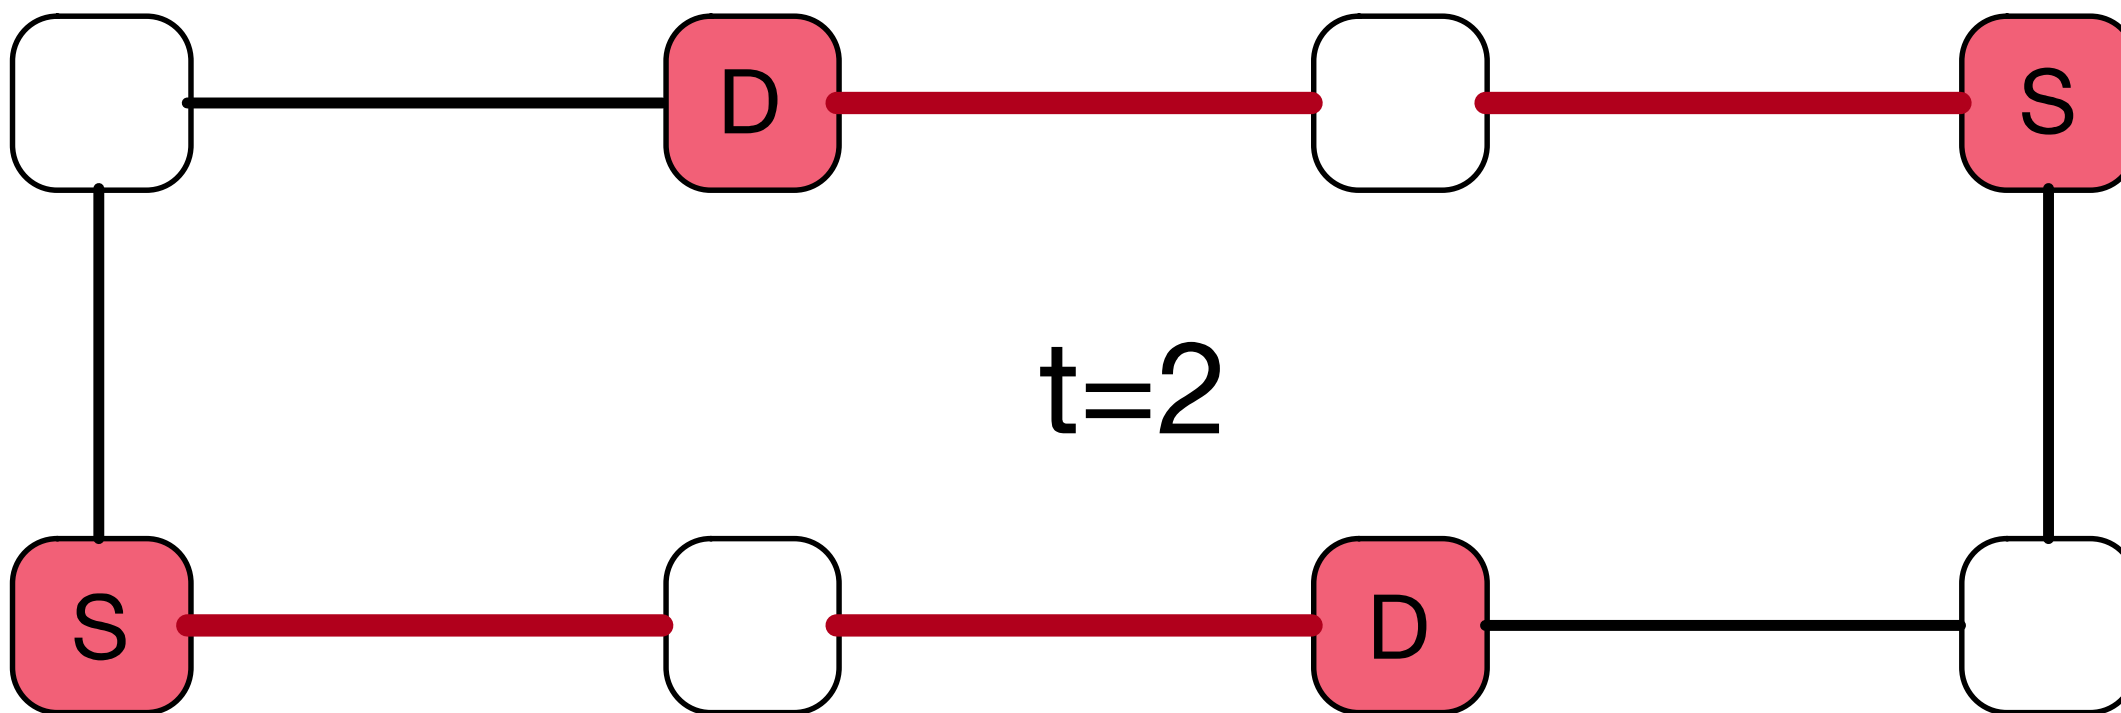
Better Strategy → Recursive Doubling

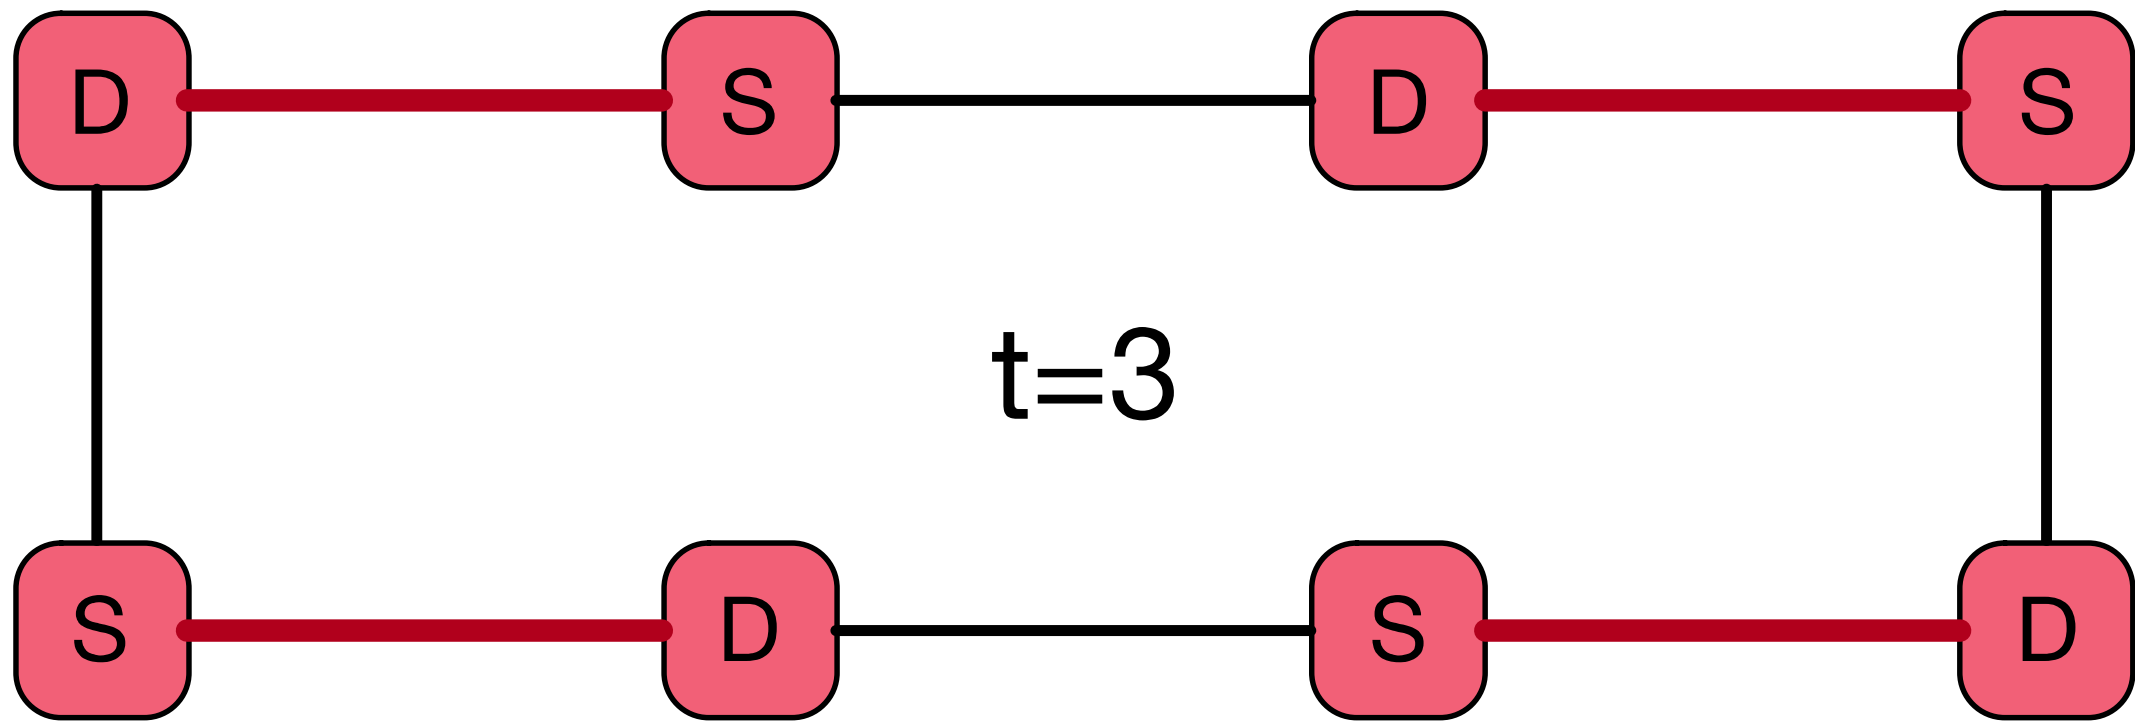


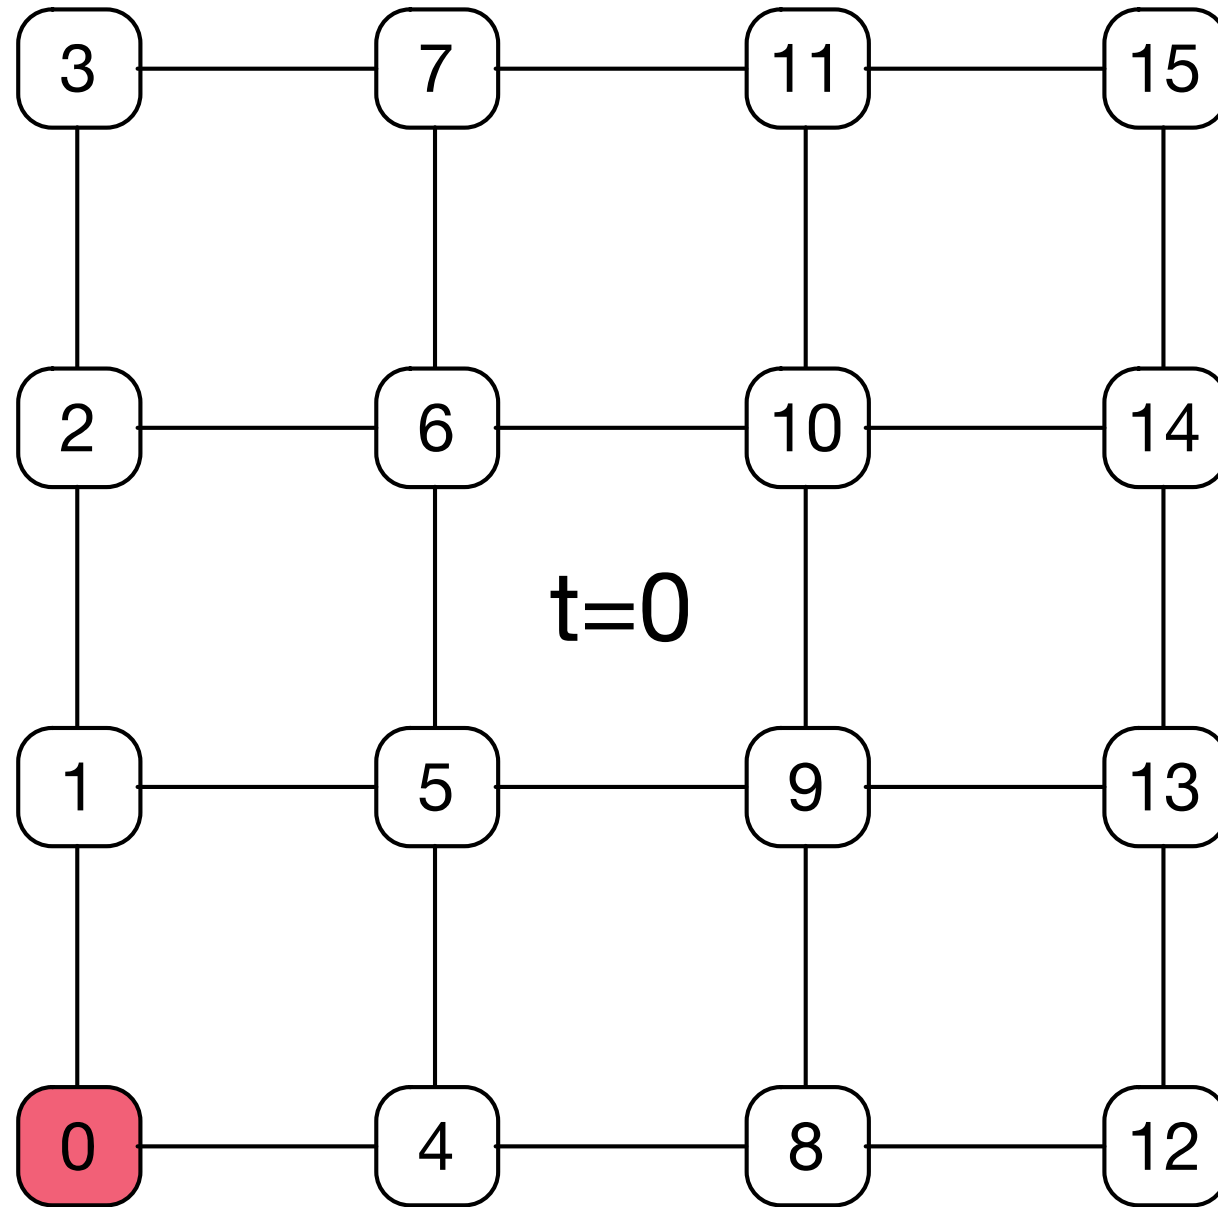
Better Strategy → Recursive Doubling

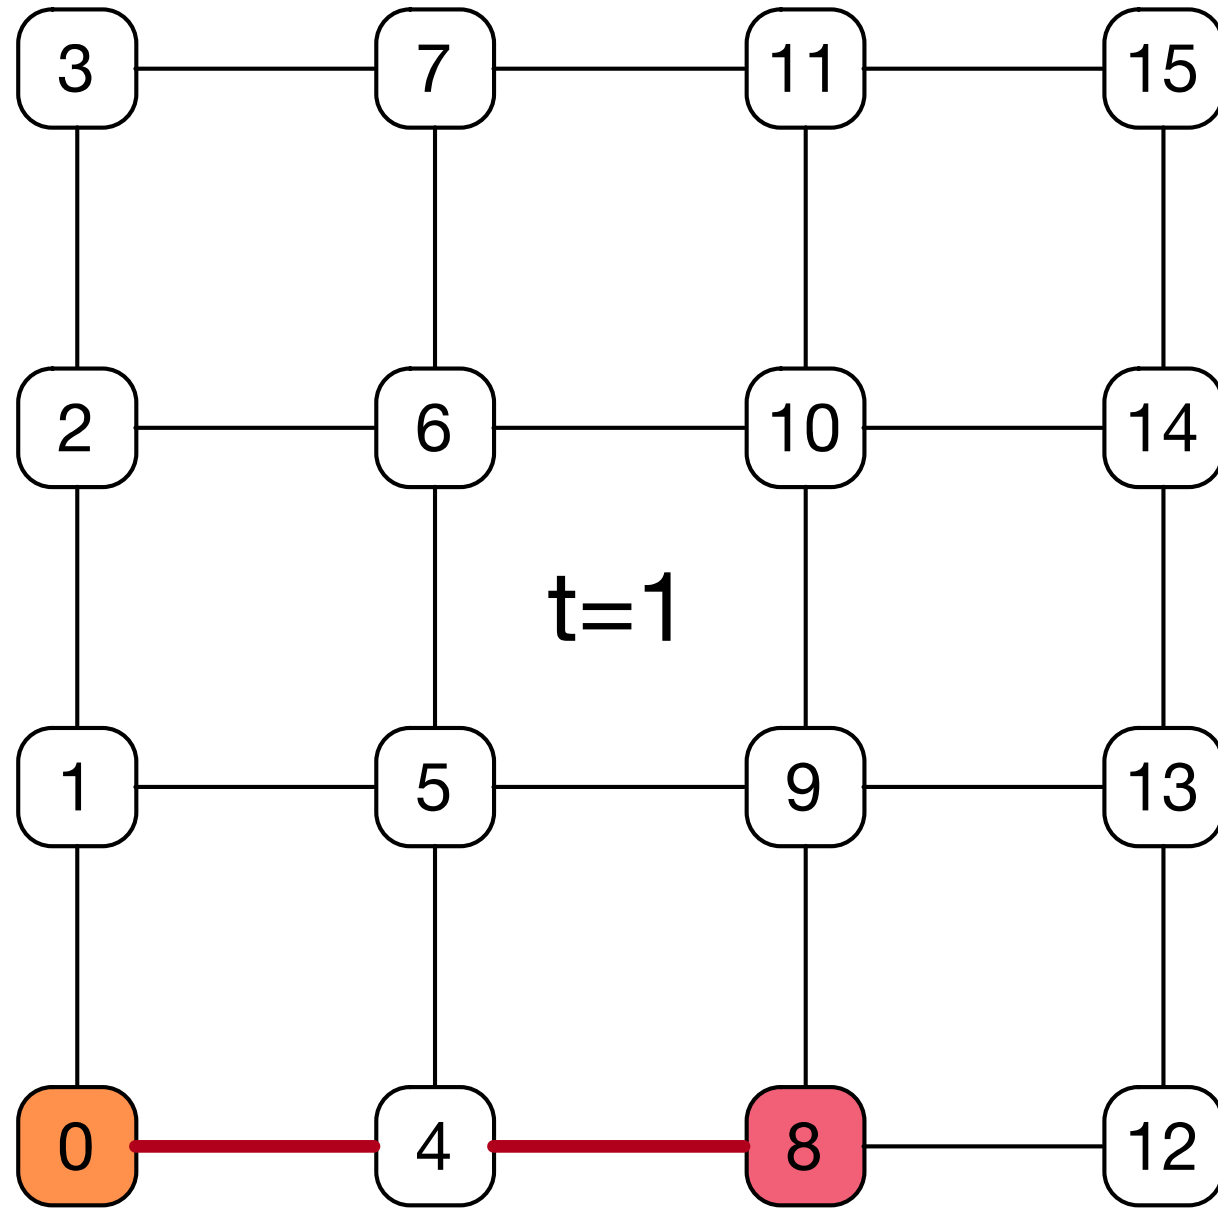


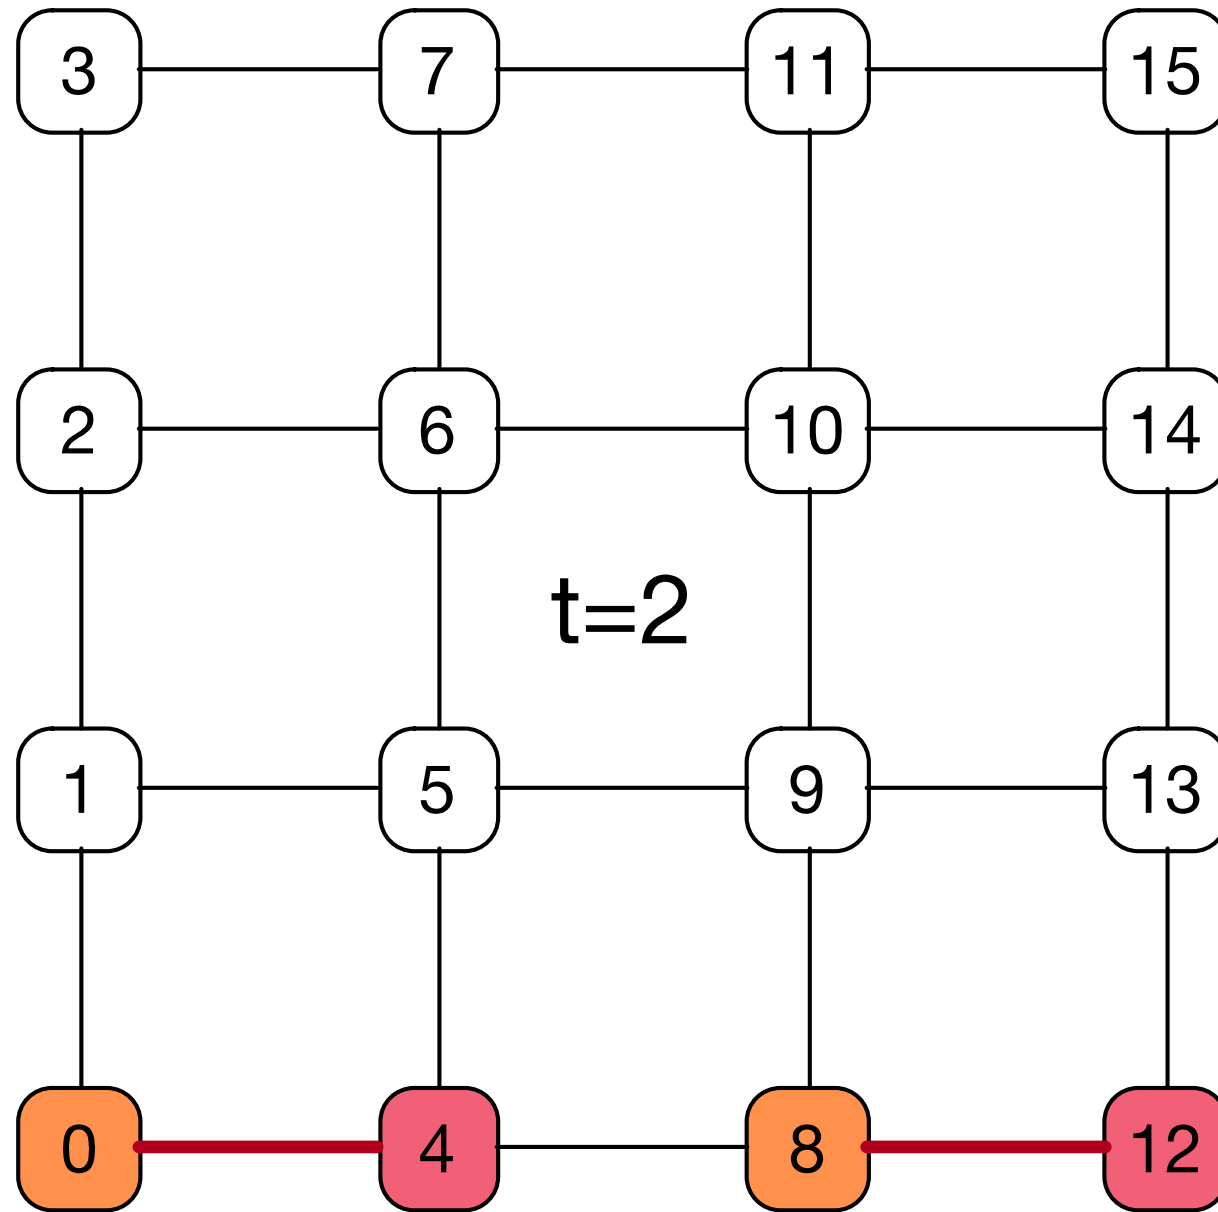


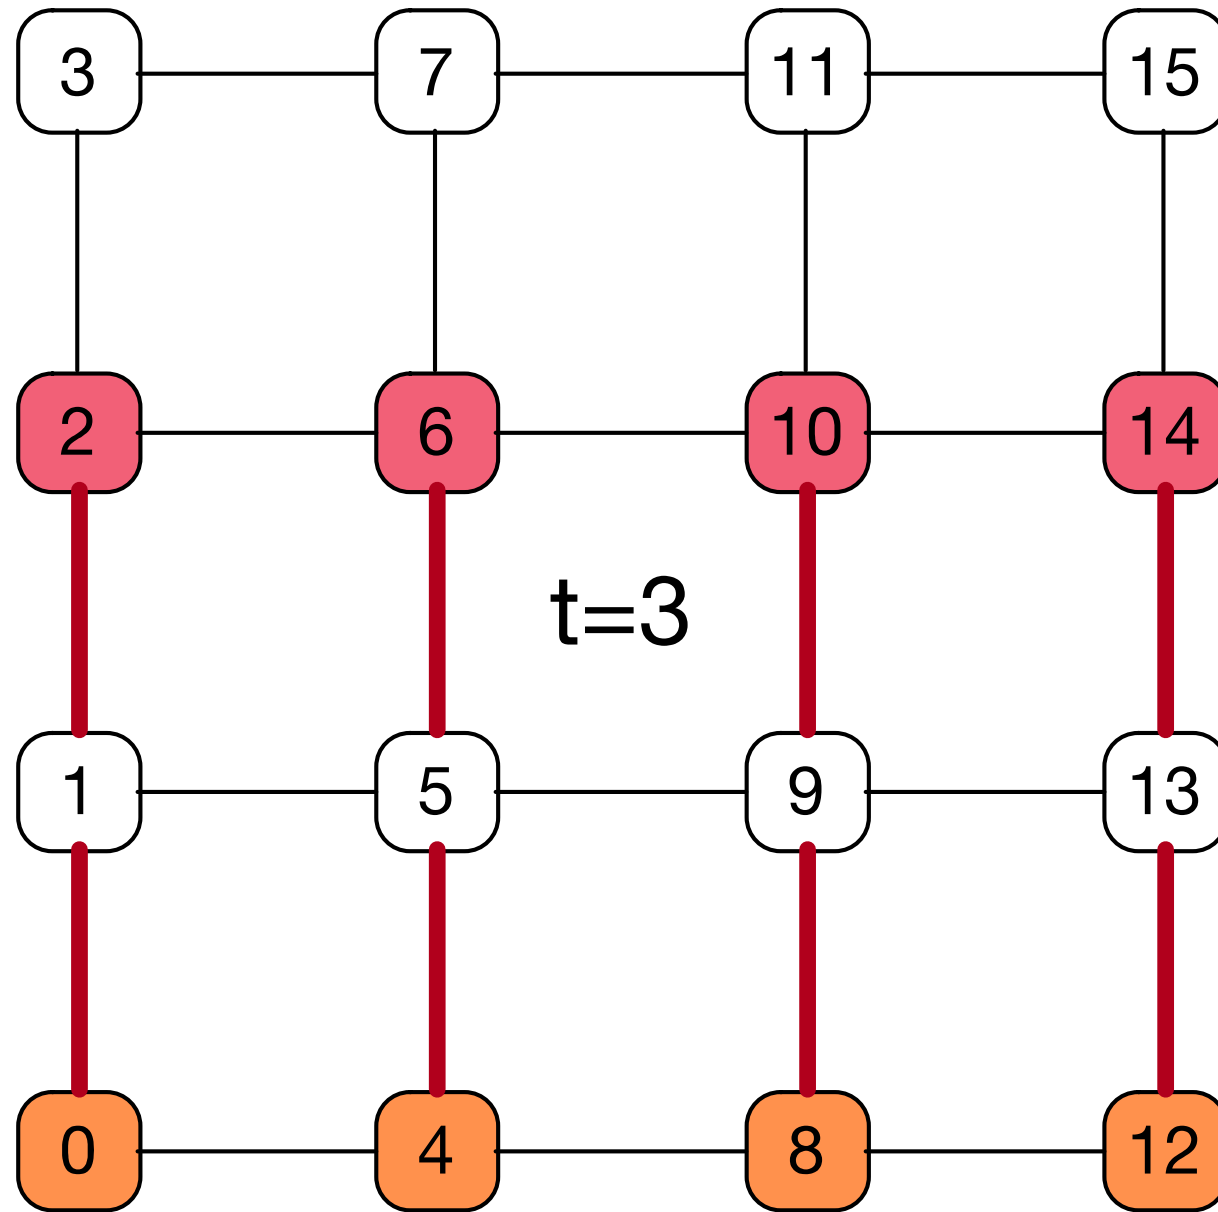


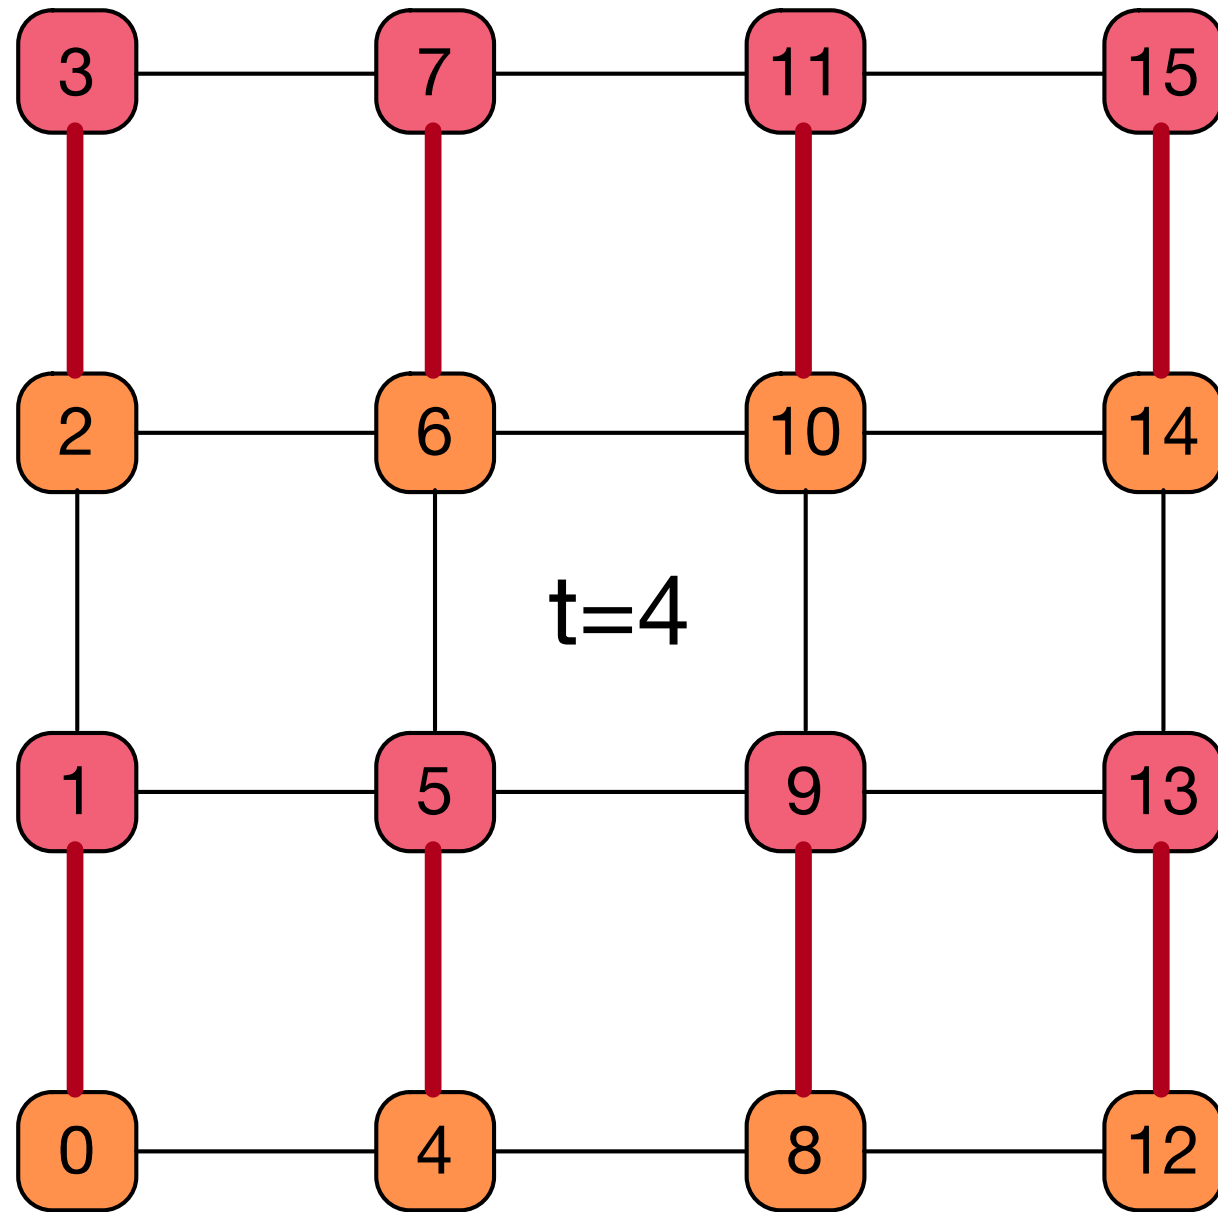


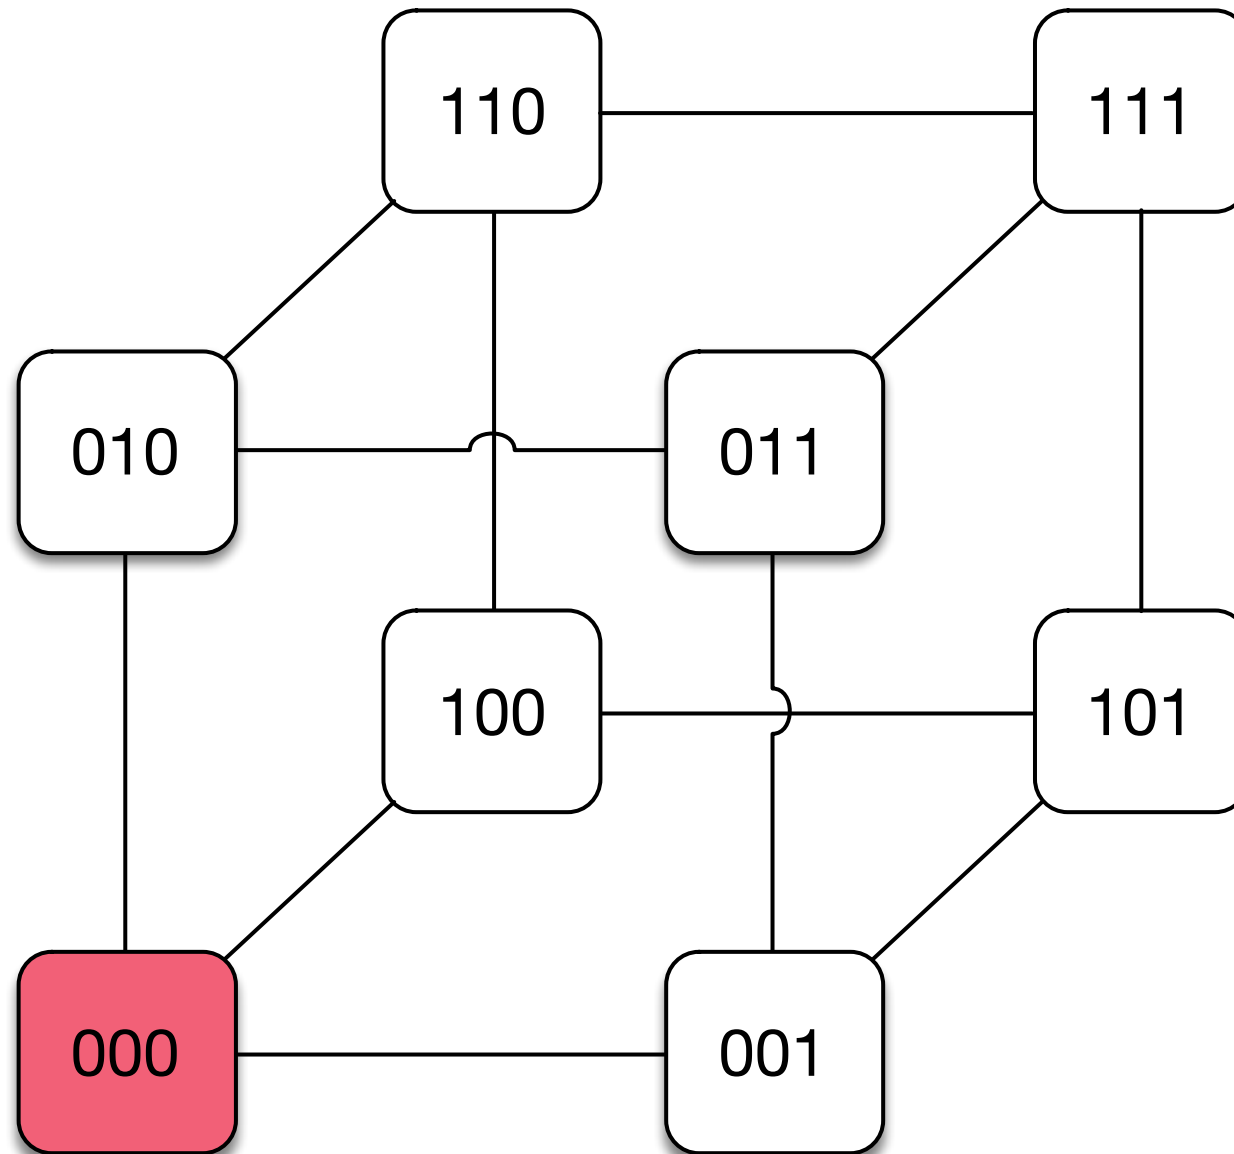


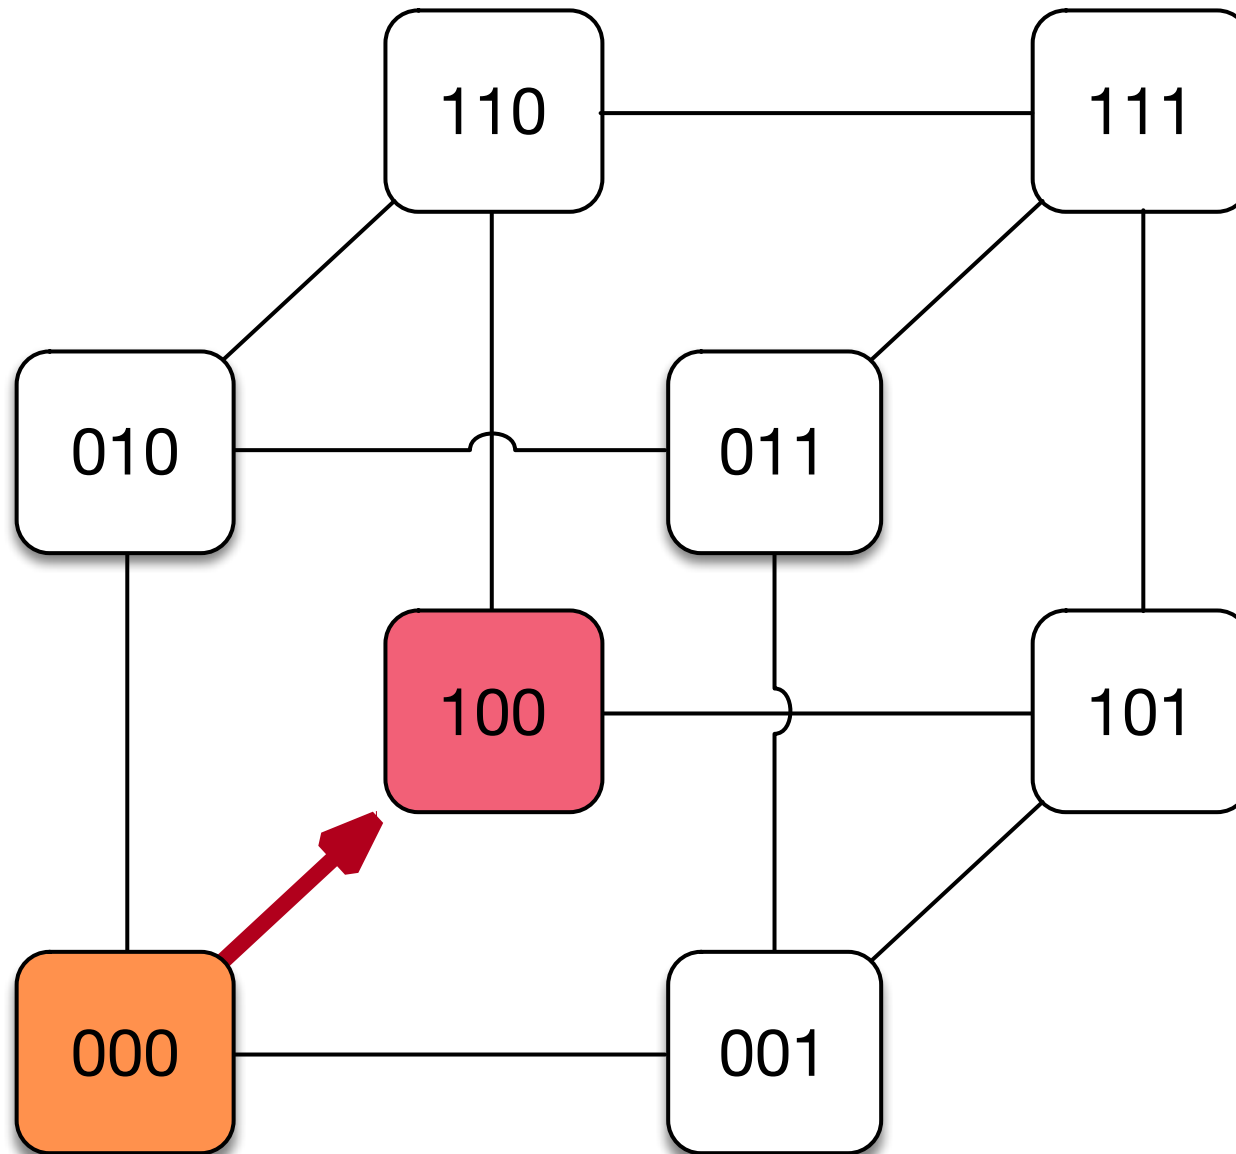


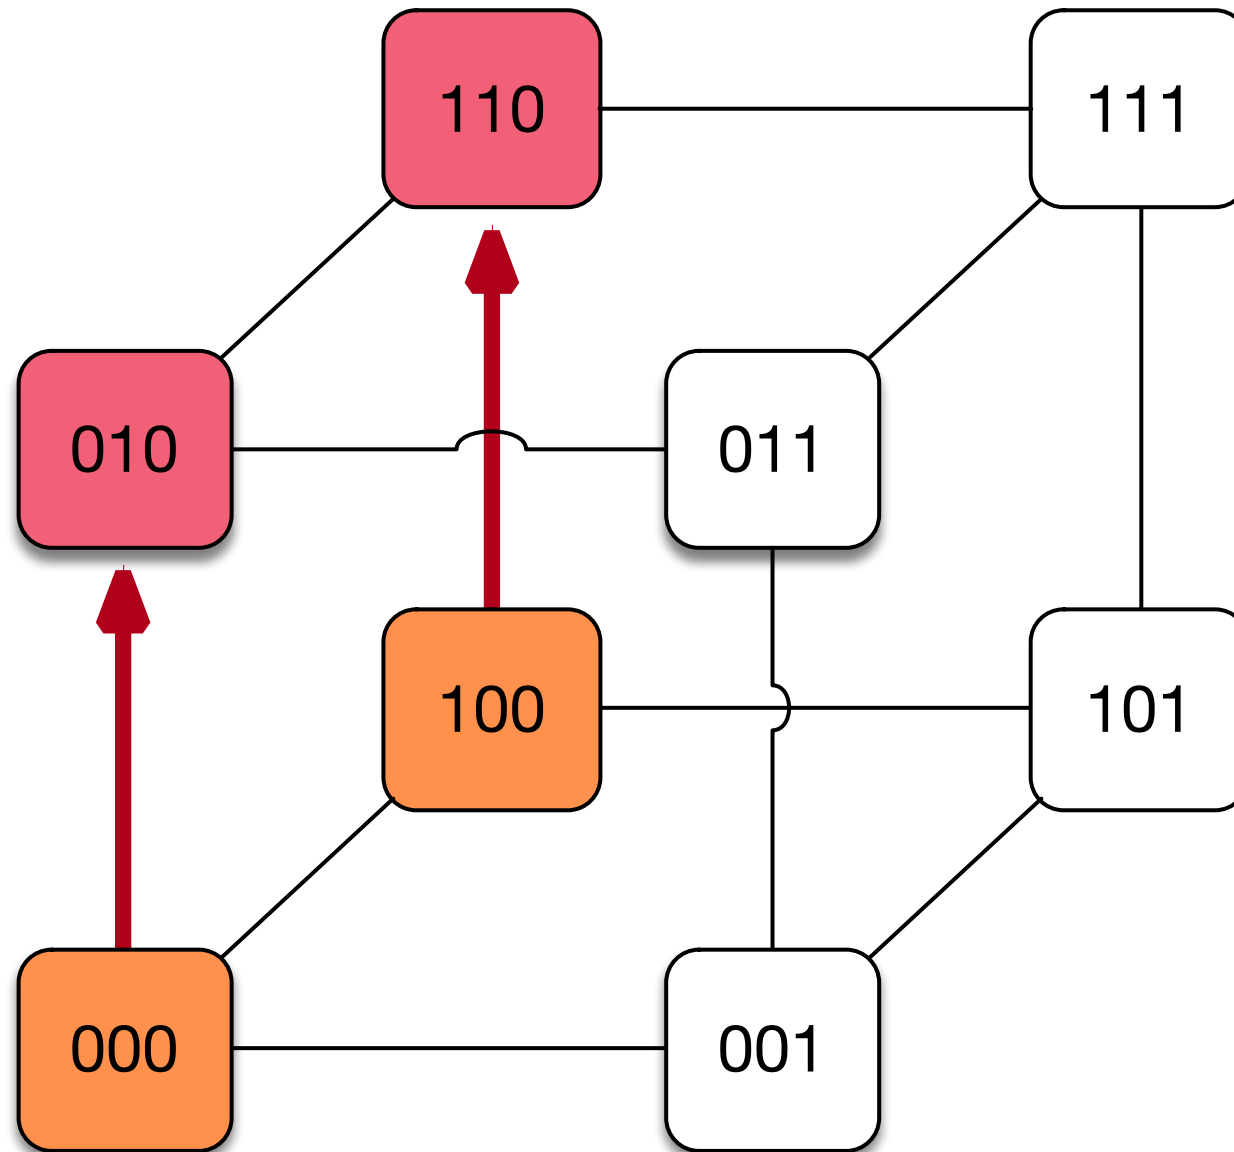


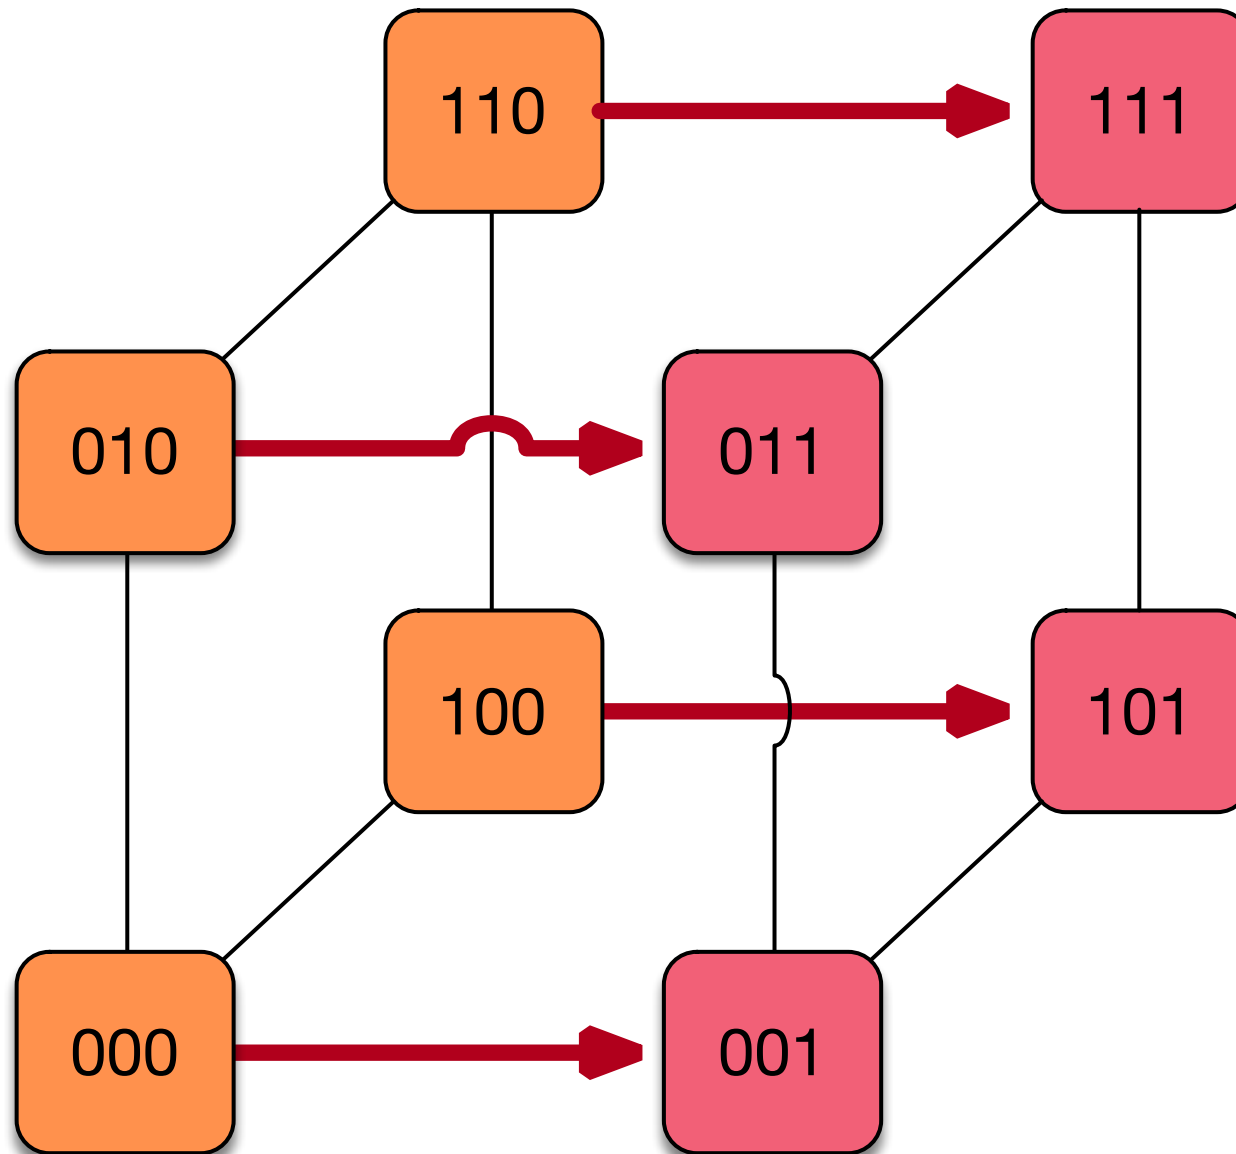












Cost Analysis for (1-n) in a Hypercube

22

- Cost of sending a single message = $T_{\text{setup}} + T_{\text{transmission}}$

$$\rightarrow T_s + T_w \times \text{Size of data}$$

$$\rightarrow T_s + T_w \times m$$

$$\text{Total Cost} = (T_s + T_w \times m) \times \text{number of messages}$$

$\log p$ messages

$$\text{Cost} \rightarrow (T_s + T_w \times m) \times \log p$$

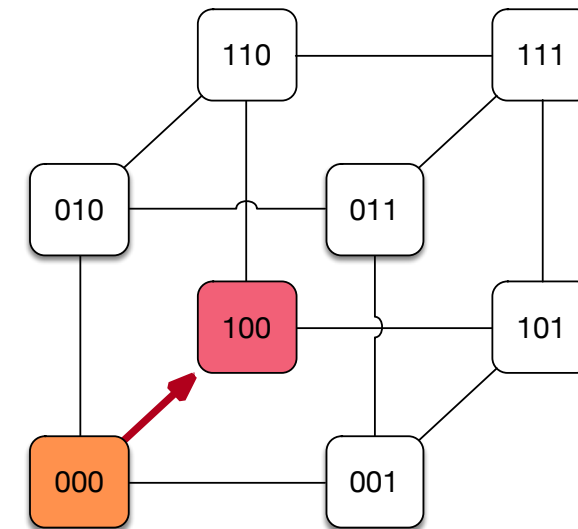
```

1.  procedure ONE_TO_ALL_BC( $d$ ,  $my\_id$ ,  $X$ )
2.  begin
3.       $mask := 2^d - 1$ ;                                /* Set all  $d$  bits of  $mask$  to 1 */ (111)
4.      for  $i := d - 1$  downto 0 do                        /* Outer loop */  $i = d-1 = 2$ 
5.           $mask := mask \text{ XOR } 2^i$ ;                  /* Set bit  $i$  of  $mask$  to 0 */ 111 XOR 100 = (011)
6.          if ( $my\_id \text{ AND } mask$ ) = 0 then              /* If lower  $i$  bits of  $my\_id$  are 0 */ (000 & 011) = 000
7.              if ( $my\_id \text{ AND } 2^i$ ) = 0 then          (100 & 100) = 100
8.                   $msg\_destination := my\_id \text{ XOR } 2^i$ ;  (000 & 100) = 000
9.                  send  $X$  to  $msg\_destination$ ;              (000 XOR 100) = 100
10.             else                                         (100 XOR 100) = 000
11.                  $msg\_source := my\_id \text{ XOR } 2^i$ ;
12.                 receive  $X$  from  $msg\_source$ ;
13.             endelse;
14.         endif;
15.     endfor;
16. end ONE_TO_ALL_BC

```

(100 & 011) = 000
 (100 & 100) = 100

0 → 4



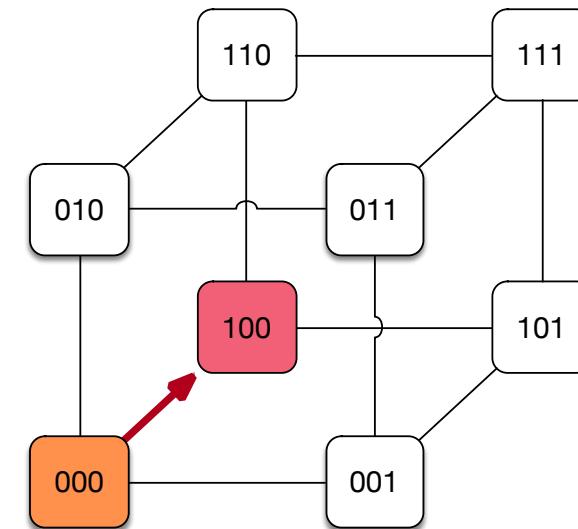
my_id = 0

```
1.  procedure ONE_TO_ALL_BC(d, my_id, X)
2.  begin
3.      mask :=  $2^d - 1$ ;                /* Set all d bits of mask to 1 */ (111)
4.      for i := d - 1 downto 0 do      /* Outer loop */ i = d-1 = 2
5.          mask := mask XOR  $2^i$ ;      /* Set bit i of mask to 0 */ 111 XOR 100 = (011)
6.          if (my_id AND mask) = 0 then /* If lower i bits of my_id are 0 */ (000 & 011) = 000
7.              if (my_id AND  $2^i$ ) = 0 then
8.                  msg_destination := my_id XOR  $2^i$ ;
9.                  send X to msg_destination;
10.             else
11.                 msg_source := my_id XOR  $2^i$ ;
12.                 receive X from msg_source;
13.             endelse;
14.         endif;
15.     endfor;
16. end ONE_TO_ALL_BC
```

(100 & 011) = 000
(100 & 100) = 100

(000 & 100) = 000
(000 XOR 100) = 100
0 → 4

(100 XOR 100) = 000



1 → All (Case 0 to 4)

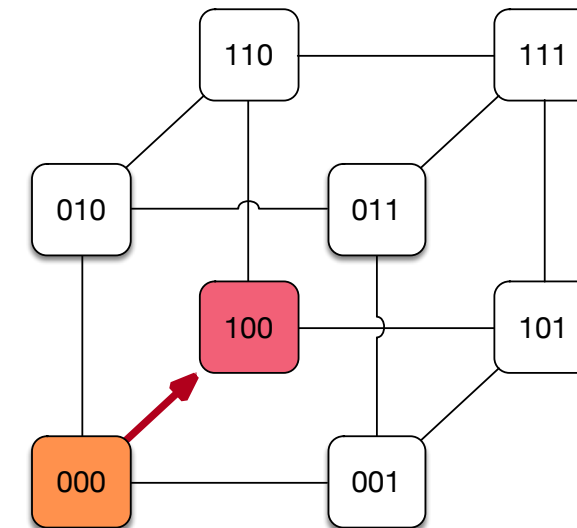
my_id = 0

```
1.  procedure ONE_TO_ALL_BC( $d, my\_id, X$ )
2.  begin
3.       $mask := 2^d - 1;$                                 /* Set all  $d$  bits of  $mask$  to 1 */ (111)
4.      for  $i := d - 1$  downto 0 do                      /* Outer loop */  $i = d-1 = 2$ 
5.           $mask := mask \text{ XOR } 2^i;$                 /* Set bit  $i$  of  $mask$  to 0 */ 111 XOR 100 = (011)
6.          if ( $my\_id \text{ AND } mask$ ) = 0 then            /* If lower  $i$  bits of  $my\_id$  are 0 */ (000 & 011) = 000
7.              if ( $my\_id \text{ AND } 2^i$ ) = 0 then
8.                   $msg\_destination := my\_id \text{ XOR } 2^i;$ 
9.                  send  $X$  to  $msg\_destination$ ;
10.             else
11.                  $msg\_source := my\_id \text{ XOR } 2^i;$     (100 XOR 100) = 000
12.                 receive  $X$  from  $msg\_source$ ;
13.             endelse;
14.         endif;
15.     endfor;
16. end ONE_TO_ALL_BC
```

(100 & 011) = 000
(100 & 100) = 100

(000 & 100) = 000
(000 XOR 100) = 100
0 → 4

NODE 0

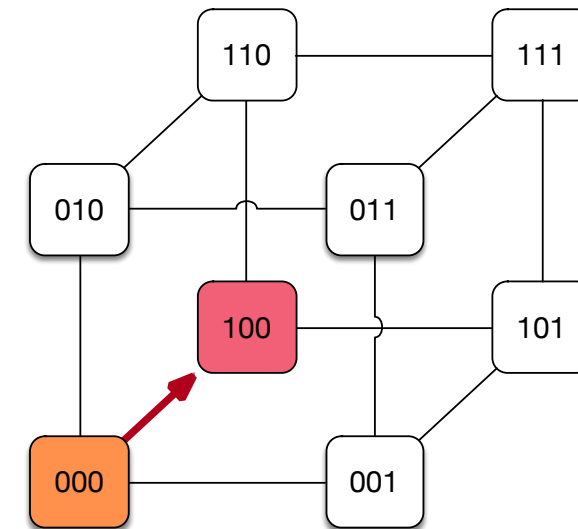


1 → All (Case 0 to 4)

```

1.  procedure ONE_TO_ALL_BC( $d$ ,  $my\_id$ ,  $X$ )
2.  begin
3.       $mask := 2^d - 1$ ;                                /* Set all  $d$  bits of  $mask$  to 1 */ (111)
4.      for  $i := d - 1$  downto 0 do                        /* Outer loop */  $i = d-1 = 2$ 
5.           $mask := mask \text{ XOR } 2^i$ ;                /* Set bit  $i$  of  $mask$  to 0 */ 111 XOR 100 = (011)
6.          if ( $my\_id \text{ AND } mask$ ) = 0 then            /* If lower  $i$  bits of  $my\_id$  are 0 */ (000 & 011) = 000
7.              if ( $my\_id \text{ AND } 2^i$ ) = 0 then          (100 & 100) = 100
8.                   $msg\_destination := my\_id \text{ XOR } 2^i$ ;  (000 & 100) = 000
9.                  send  $X$  to  $msg\_destination$ ;              (000 XOR 100) = 100
10.             else                                         (100 XOR 100) = 000
11.                  $msg\_source := my\_id \text{ XOR } 2^i$ ;
12.                 receive  $X$  from  $msg\_source$ ;
13.             endelse;
14.         endif;
15.     endfor;
16. end ONE_TO_ALL_BC

```



```

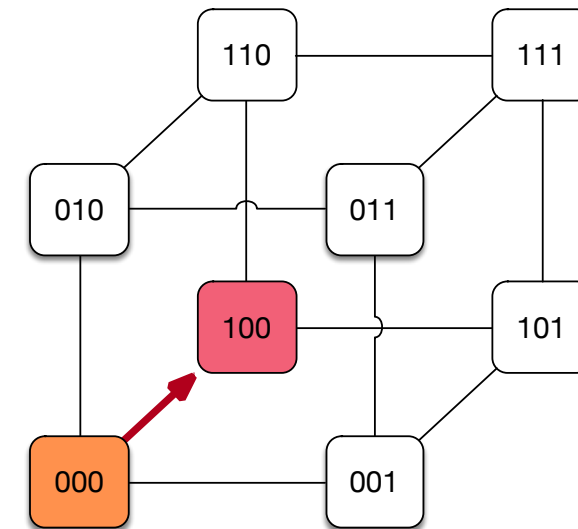
1.  procedure ONE_TO_ALL_BC( $d$ ,  $my\_id$ ,  $X$ )
2.  begin
3.       $mask := 2^d - 1$ ;                                /* Set all  $d$  bits of  $mask$  to 1 */ (111)
4.      for  $i := d - 1$  downto 0 do                        /* Outer loop */  $i = d-1 = 2$ 
5.           $mask := mask \text{ XOR } 2^i$ ;                /* Set bit  $i$  of  $mask$  to 0 */ 111 XOR 100 = (011)
6.          if ( $my\_id \text{ AND } mask$ ) = 0 then            /* If lower  $i$  bits of  $my\_id$  are 0 */ (000 & 011) = 000
7.              if ( $my\_id \text{ AND } 2^i$ ) = 0 then
8.                   $msg\_destination := my\_id \text{ XOR } 2^i$ ;
9.                  send  $X$  to  $msg\_destination$ ;
10.             else
11.                  $msg\_source := my\_id \text{ XOR } 2^i$ ;    (100 XOR 100) = 000
12.                 receive  $X$  from  $msg\_source$ ;
13.             endelse;
14.         endif;
15.     endfor;
16. end ONE_TO_ALL_BC

```

(100 & 011) = 000
(100 & 100) = 100

NODE 4

(000 & 100) = 000
(000 XOR 100) = 100
0 → 4



1 → All (Case 0 to 4)

my_id = 4

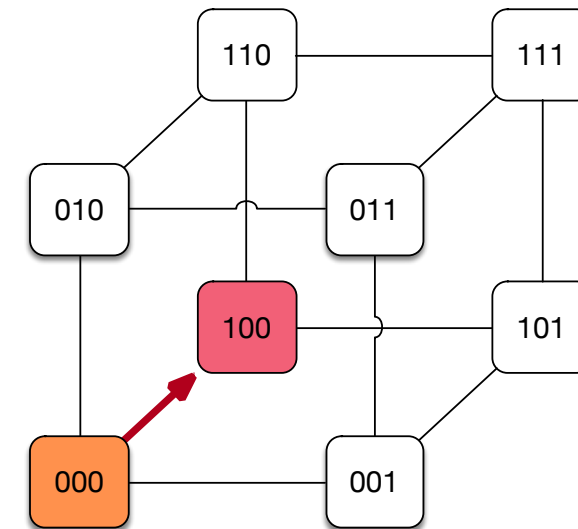
```
1.  procedure ONE_TO_ALL_BC(d, my_id, X)
2.  begin
3.      mask :=  $2^d - 1$ ;                /* Set all d bits of mask to 1 */ (111)
4.      for i := d - 1 downto 0 do      /* Outer loop */ i = d-1 = 2
5.          mask := mask XOR  $2^i$ ;      /* Set bit i of mask to 0 */ 111 XOR 100 = (011)
6.          if (my_id AND mask) = 0 then /* If lower i bits of my_id are 0 */ (000 & 011) = 000
7.              if (my_id AND  $2^i$ ) = 0 then
8.                  msg_destination := my_id XOR  $2^i$ ;  (000 & 100) = 000
9.                  send X to msg_destination;          (000 XOR 100) = 100
10.             else
11.                 msg_source := my_id XOR  $2^i$ ;        (100 XOR 100) = 000
12.                 receive X from msg_source;
13.             endelse;
14.         endif;
15.     endfor;
16. end ONE_TO_ALL_BC
```

(100 & 011) = 000

(100 & 100) = 100

NODE 4

0 → 4

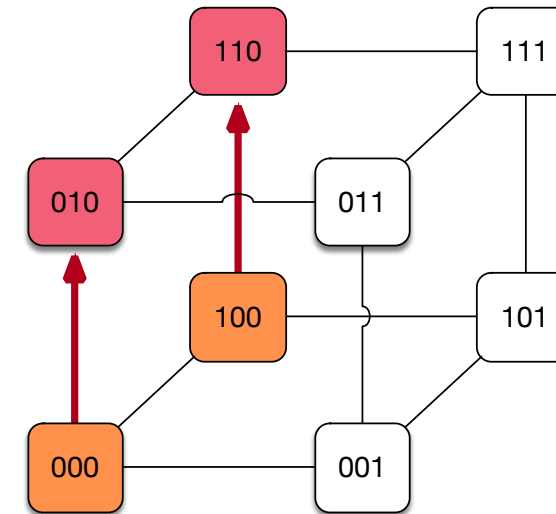


1 → All (Case 0 to 4)

```

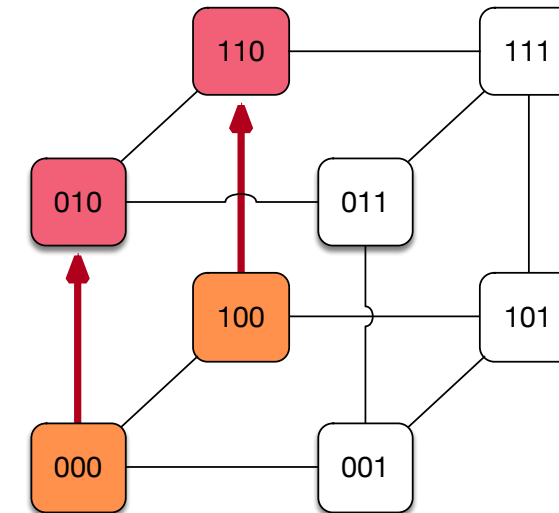
1.  procedure ONE_TO_ALL_BC( $d, my\_id, X$ )
2.  begin
3.       $mask := 2^d - 1;$                                 /* Set all  $d$  bits of  $mask$  to 1 */ (111)
4.      for  $i := d - 1$  downto 0 do                        /* Outer loop */       $i=1$ 
5.           $mask := mask \text{ XOR } 2^i;$                 /* Set bit  $i$  of  $mask$  to 0 */ 111 XOR 010 = (101)
6.          if ( $my\_id \text{ AND } mask$ ) = 0 then            /* If lower  $i$  bits of  $my\_id$  are 0 */ (000 & 101) = 000
7.              if ( $my\_id \text{ AND } 2^i$ ) = 0 then          (010 & 010) = 010
8.                   $msg\_destination := my\_id \text{ XOR } 2^i;$  (000 & 010) = 000
9.                  send  $X$  to  $msg\_destination$ ;          (000 XOR 010) = 010
10.                 else                                     0 → 2
11.                      $msg\_source := my\_id \text{ XOR } 2^i;$  (010 XOR 010) = 000
12.                     receive  $X$  from  $msg\_source$ ;
13.                 endelse;
14.             endif;
15.         endfor;
16.     end ONE_TO_ALL_BC

```



my_id = 0

```
1.  procedure ONE_TO_ALL_BC(d, my_id, X)
2.  begin
3.      mask :=  $2^d - 1$ ;                /* Set all d bits of mask to 1 */ (111)
4.      for i := d - 1 downto 0 do      /* Outer loop */      i=1
5.          mask := mask XOR  $2^i$ ;      /* Set bit i of mask to 0 */ 111 XOR 010 = (101)
(010 & 101) = 000 if (my_id AND mask) = 0 then /* If lower i bits of my_id are 0 */ (000 & 101) = 000
(010 & 010) = 010 if (my_id AND  $2^i$ ) = 0 then (000 & 010) = 000
8.          msg_destination := my_id XOR  $2^i$ ; (000 XOR 010) = 010
9.          send X to msg_destination;      0 → 2
10.         else (010 XOR 010) = 000
11.             msg_source := my_id XOR  $2^i$ ;
12.             receive X from msg_source;
13.         endelse;
14.     endif;
15. endfor;
16. end ONE_TO_ALL_BC
```



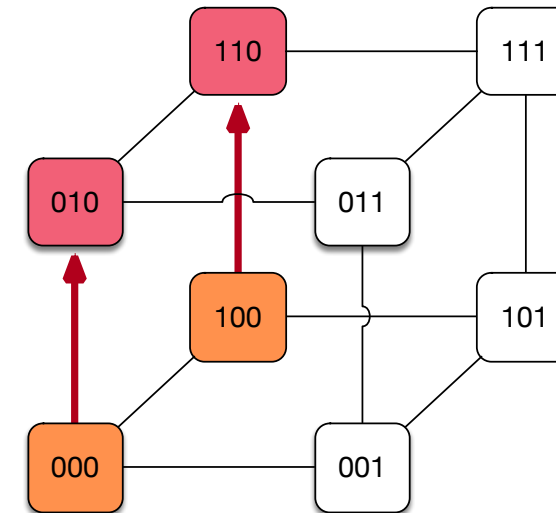
1 → All (Case 0 to 2)

my_id = 0

```
1.  procedure ONE_TO_ALL_BC(d, my_id, X)
2.  begin
3.      mask :=  $2^d - 1$ ;                /* Set all d bits of mask to 1 */ (111)
4.      for i := d - 1 downto 0 do      /* Outer loop */      i=1
5.          mask := mask XOR  $2^i$ ;      /* Set bit i of mask to 0 */ 111 XOR 010 = (101)
6.          if (my_id AND mask) = 0 then /* If lower i bits of my_id are 0 */ (000 & 101) = 000
7.              if (my_id AND  $2^i$ ) = 0 then
8.                  msg_destination := my_id XOR  $2^i$ ; (010 & 010) = 010
9.                  send X to msg_destination;
10.             else
11.                 msg_source := my_id XOR  $2^i$ ; (010 XOR 010) = 000
12.                 receive X from msg_source;
13.             endelse;
14.         endif;
15.     endfor;
16. end ONE_TO_ALL_BC
```

(000 & 010) = 000
(000 XOR 010) = 010
0 → 2

NODE 0

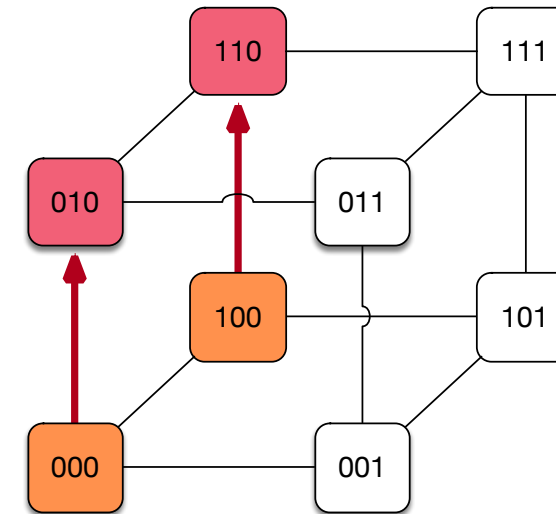


1 → All (Case 0 to 2)

```

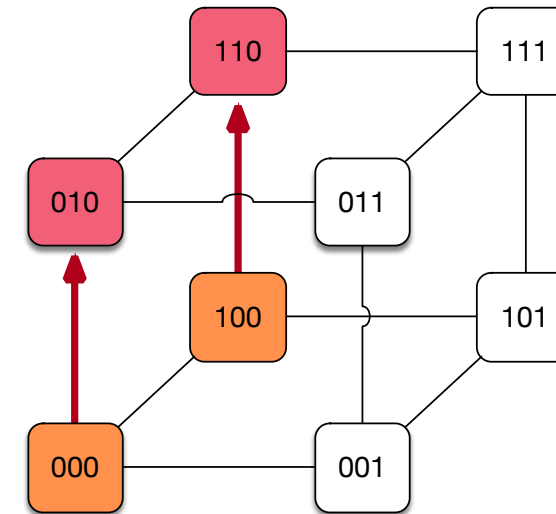
1.  procedure ONE_TO_ALL_BC( $d$ ,  $my\_id$ ,  $X$ )
2.  begin
3.       $mask := 2^d - 1$ ;                                /* Set all  $d$  bits of  $mask$  to 1 */ (111)
4.      for  $i := d - 1$  downto 0 do                        /* Outer loop */       $i=1$ 
5.           $mask := mask \text{ XOR } 2^i$ ;                /* Set bit  $i$  of  $mask$  to 0 */ 111 XOR 010 = (101)
6.          if ( $my\_id \text{ AND } mask$ ) = 0 then            /* If lower  $i$  bits of  $my\_id$  are 0 */ (000 & 101) = 000
7.              if ( $my\_id \text{ AND } 2^i$ ) = 0 then          (010 & 010) = 010
8.                   $msg\_destination := my\_id \text{ XOR } 2^i$ ; (000 & 010) = 000
9.                  send  $X$  to  $msg\_destination$ ;          (000 XOR 010) = 010
10.             else                                       (010 XOR 010) = 000
11.                  $msg\_source := my\_id \text{ XOR } 2^i$ ;
12.                 receive  $X$  from  $msg\_source$ ;
13.             endelse;
14.         endif;
15.     endfor;
16. end ONE_TO_ALL_BC

```



my_id = 2

```
1.  procedure ONE_TO_ALL_BC(d, my_id, X)
2.  begin
3.      mask :=  $2^d - 1$ ;                /* Set all d bits of mask to 1 */ (111)
4.      for i := d - 1 downto 0 do      /* Outer loop */      i=1
5.          mask := mask XOR  $2^i$ ;      /* Set bit i of mask to 0 */ 111 XOR 010 = (101)
(010 & 101) = 000 if (my_id AND mask) = 0 then /* If lower i bits of my_id are 0 */ (000 & 101) = 000
(010 & 010) = 010 if (my_id AND  $2^i$ ) = 0 then      (000 & 010) = 000
8.          msg_destination := my_id XOR  $2^i$ ; (000 XOR 010) = 010
9.          send X to msg_destination;      0 → 2
10.         else
11.             msg_source := my_id XOR  $2^i$ ; (010 XOR 010) = 000
12.             receive X from msg_source;
13.         endelse;
14.     endif;
15. endfor;
16. end ONE_TO_ALL_BC
```



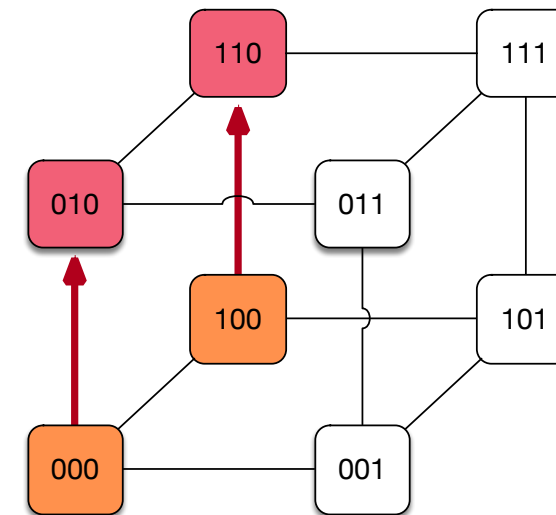
1 → All (Case 0 to 2)

my_id = 2

```
1.  procedure ONE_TO_ALL_BC(d, my_id, X)
2.  begin
3.      mask :=  $2^d - 1$ ;                /* Set all d bits of mask to 1 */ (111)
4.      for i := d - 1 downto 0 do      /* Outer loop */      i=1
5.          mask := mask XOR  $2^i$ ;      /* Set bit i of mask to 0 */ 111 XOR 010 = (101)
6.          if (my_id AND mask) = 0 then /* If lower i bits of my_id are 0 */ (000 & 101) = 000
7.              if (my_id AND  $2^i$ ) = 0 then
8.                  msg_destination := my_id XOR  $2^i$ ; (000 & 010) = 000
9.                  send X to msg_destination; (000 XOR 010) = 010
10.             else
11.                 msg_source := my_id XOR  $2^i$ ; (010 XOR 010) = 000
12.                 receive X from msg_source;
13.             endelse;
14.         endif;
15.     endfor;
16. end ONE_TO_ALL_BC
```

(010 & 101) = 000
(010 & 010) = 010

NODE 2



1 → All (Case 0 to 2)

```
1.  procedure ONE_TO_ALL_BC( $d$ ,  $my\_id$ ,  $X$ )
```

```
2.  begin
```

```
3.     $mask := 2^d - 1$ ;                                /* Set all  $d$  bits of  $mask$  to 1 */ (111)
```

```
4.    for  $i := d - 1$  downto 0 do                        /* Outer loop */
```

```
5.       $mask := mask \text{ XOR } 2^i$ ;                    /* Set bit  $i$  of  $mask$  to 0 */ 111 XOR 001 = (110)
```

```
(001 & 110) = 000 if ( $my\_id \text{ AND } mask$ ) = 0 then /* If lower  $i$  bits of  $my\_id$  are 0 */ (000 & 110) = 000
```

```
(001 & 001) = 001 if ( $my\_id \text{ AND } 2^i$ ) = 0 then
```

```
8.       $msg\_destination := my\_id \text{ XOR } 2^i$ ;
```

```
9.      send  $X$  to  $msg\_destination$ ;
```

```
10.     else                                           (001 XOR 001) = 000
```

```
11.       $msg\_source := my\_id \text{ XOR } 2^i$ ;
```

```
12.      receive  $X$  from  $msg\_source$ ;
```

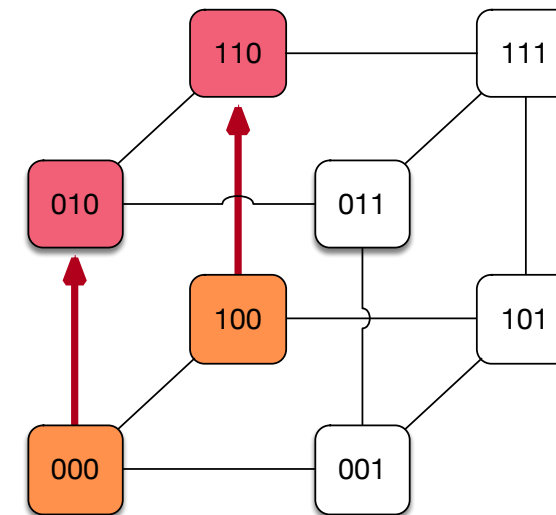
```
13.    endelse;
```

```
14.  endif;
```

```
15.  endfor;
```

```
16. end ONE_TO_ALL_BC
```

(000 & 001) = 000
(000 XOR 001) = 001
0 → 1



1 → All (Case 0 to 1)

my_id = 0

```
1.  procedure ONE_TO_ALL_BC( $d$ , my_id,  $X$ )
2.  begin
3.      mask :=  $2^d - 1$ ;          /* Set all  $d$  bits of mask to 1 */
4.      for  $i := d - 1$  downto 0 do /* Outer loop */
5.          mask := mask XOR  $2^i$ ; /* Set bit  $i$  of mask to 0 */
6.          if (my_id AND mask) = 0 then /* If lower  $i$  bits of my_id are 0 */
7.              if (my_id AND  $2^i$ ) = 0 then
8.                  msg_destination := my_id XOR  $2^i$ ;
9.                  send  $X$  to msg_destination;
10.             else
11.                 msg_source := my_id XOR  $2^i$ ;
12.                 receive  $X$  from msg_source;
13.             endelse;
14.         endif;
15.     endfor;
16. end ONE_TO_ALL_BC
```

/* Set all d bits of mask to 1 */ (111)

/* Outer loop */

/* Set bit i of mask to 0 */

/* If lower i bits of my_id are 0 */

if (my_id AND 2^i) = 0 then

msg_destination := my_id XOR 2^i ;

send X to msg_destination;

else

msg_source := my_id XOR 2^i ;

receive X from msg_source;

endelse;

endif;

endfor;

end ONE_TO_ALL_BC

$i=0$

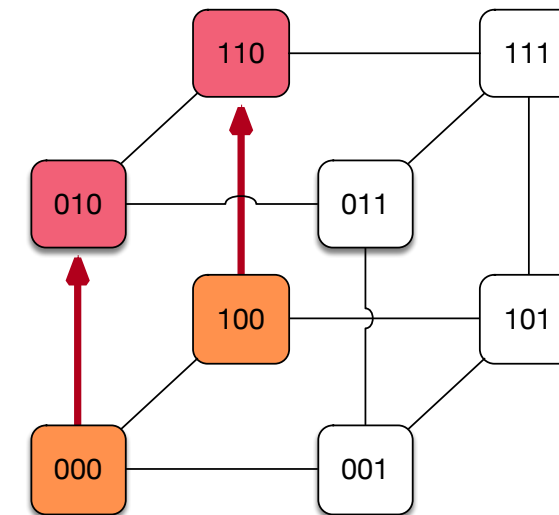
111 XOR 001 = (110)

(000 & 110) = 000

(000 & 001) = 000

(000 XOR 001) = 001

0 → 1



1 → All (Case 0 to 1)

my_id = 0

```
1.  procedure ONE_TO_ALL_BC( $d, my\_id, X$ )
2.  begin
3.       $mask := 2^d - 1;$                 /* Set all  $d$  bits of  $mask$  to 1 */ (111)
4.      for  $i := d - 1$  downto 0 do      /* Outer loop */                 $i=0$ 
5.           $mask := mask \text{ XOR } 2^i;$  /* Set bit  $i$  of  $mask$  to 0 */    111 XOR 001 = (110)
6.          if ( $my\_id \text{ AND } mask$ ) = 0 then /* If lower  $i$  bits of  $my\_id$  are 0 */
7.              if ( $my\_id \text{ AND } 2^i$ ) = 0 then
8.                   $msg\_destination := my\_id \text{ XOR } 2^i;$ 
9.                  send  $X$  to  $msg\_destination$ ;
10.             else
11.                  $msg\_source := my\_id \text{ XOR } 2^i;$ 
12.                 receive  $X$  from  $msg\_source$ ;
13.             endelse;
14.         endif;
15.     endfor;
16. end ONE_TO_ALL_BC
```

(001 & 110) = 000

(001 & 001) = 001

if ($my_id \text{ AND } mask$) = 0 then

if ($my_id \text{ AND } 2^i$) = 0 then

$msg_destination := my_id \text{ XOR } 2^i;$

send X to $msg_destination$;

else

(001 XOR 001) = 000

$msg_source := my_id \text{ XOR } 2^i;$

receive X from msg_source ;

endelse;

endif;

endfor;

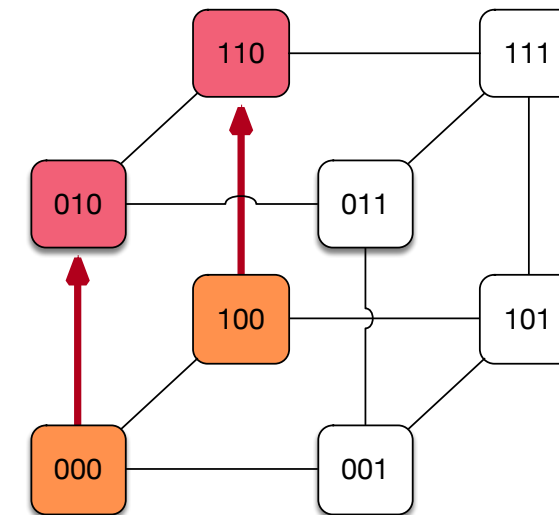
end ONE_TO_ALL_BC

(000 & 001) = 000

(000 XOR 001) = 001

0 → 1

NODE 0



1 → All (Case 0 to 1)

```
1.  procedure ONE_TO_ALL_BC( $d$ ,  $my\_id$ ,  $X$ )
```

```
2.  begin
```

```
3.     $mask := 2^d - 1$ ;                                /* Set all  $d$  bits of  $mask$  to 1 */ (111)
```

```
4.    for  $i := d - 1$  downto 0 do                        /* Outer loop */
```

```
5.       $mask := mask \text{ XOR } 2^i$ ;                    /* Set bit  $i$  of  $mask$  to 0 */ 111 XOR 001 = (110)
```

```
(001 & 110) = 000 if ( $my\_id \text{ AND } mask$ ) = 0 then /* If lower  $i$  bits of  $my\_id$  are 0 */ (000 & 110) = 000
```

```
(001 & 001) = 001 if ( $my\_id \text{ AND } 2^i$ ) = 0 then
```

```
8.       $msg\_destination := my\_id \text{ XOR } 2^i$ ;
```

```
9.      send  $X$  to  $msg\_destination$ ;
```

```
10.     else                                           (001 XOR 001) = 000
```

```
11.       $msg\_source := my\_id \text{ XOR } 2^i$ ;
```

```
12.      receive  $X$  from  $msg\_source$ ;
```

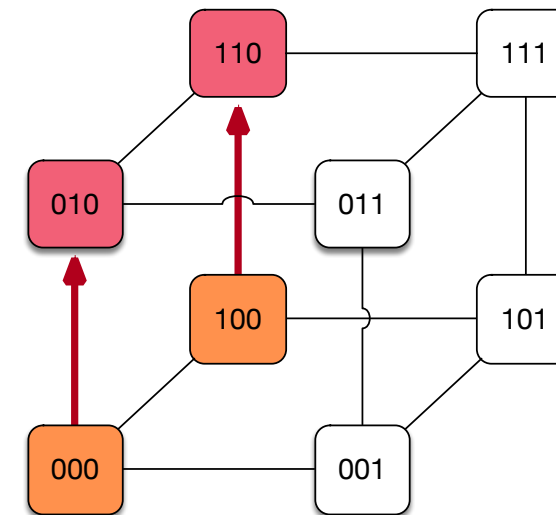
```
13.    endelse;
```

```
14.  endif;
```

```
15.  endfor;
```

```
16. end ONE_TO_ALL_BC
```

(000 & 001) = 000
(000 XOR 001) = 001
0 → 1

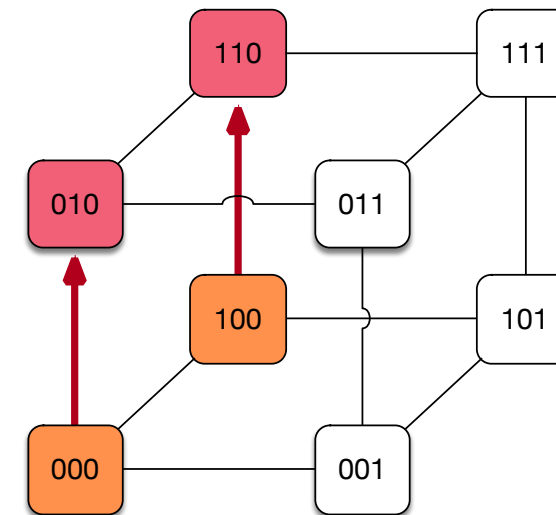


1 → All (Case 0 to 1)

my_id = 1

```
1.  procedure ONE_TO_ALL_BC(d, my_id, X)
2.  begin
3.      mask :=  $2^d - 1$ ;          /* Set all d bits of mask to 1 */ (111)
4.      for i := d - 1 downto 0 do /* Outer loop */ i=0
5.          mask := mask XOR  $2^i$ ; /* Set bit i of mask to 0 */ 111 XOR 001 = (110)
6.          if (my_id AND mask) = 0 then /* If lower i bits of my_id are 0 */ (000 & 110) = 000
7.              if (my_id AND  $2^i$ ) = 0 then (001 & 001) = 001
8.                  msg_destination := my_id XOR  $2^i$ ;
9.                  send X to msg_destination;
10.             else
11.                 msg_source := my_id XOR  $2^i$ ; (001 XOR 001) = 000
12.                 receive X from msg_source;
13.             endelse;
14.         endif;
15.     endfor;
16. end ONE_TO_ALL_BC
```

(000 & 110) = 000
(000 & 001) = 000
(000 XOR 001) = 001
0 → 1



1 → All (Case 0 to 1)

my_id = 1

```
1.  procedure ONE_TO_ALL_BC(d, my_id, X)
2.  begin
3.      mask :=  $2^d - 1$ ;
4.      for i := d - 1 downto 0 do
5.          mask := mask XOR  $2^i$ ;
6.          if (my_id AND mask) = 0 then
7.              if (my_id AND  $2^i$ ) = 0 then
8.                  msg_destination := my_id XOR  $2^i$ ;
9.                  send X to msg_destination;
10.             else
11.                 msg_source := my_id XOR  $2^i$ ;
12.                 receive X from msg_source;
13.             endelse;
14.         endif;
15.     endfor;
16. end ONE_TO_ALL_BC
```

(001 & 110) = 000
(001 & 001) = 001

NODE 1

(001 XOR 001) = 000

/* Set all d bits of *mask* to 1 */ (111)

/* Outer loop */

/* Set bit i of *mask* to 0 */

/* If lower i bits of *my_id* are 0 */

(000 & 001) = 000

(000 XOR 001) = 001

0 → 1

$i=0$

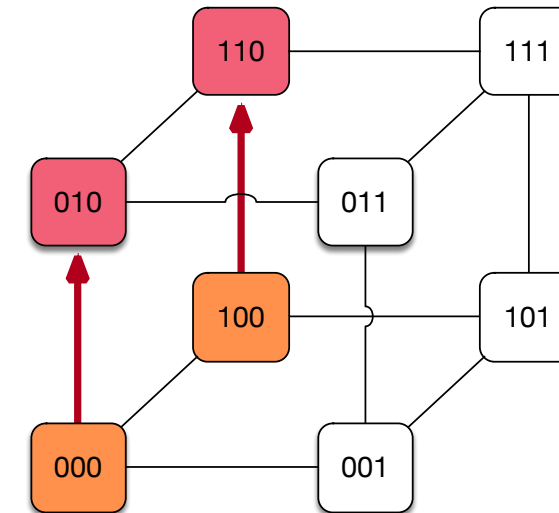
111 XOR 001 = (110)

(000 & 110) = 000

(000 & 001) = 000

(000 XOR 001) = 001

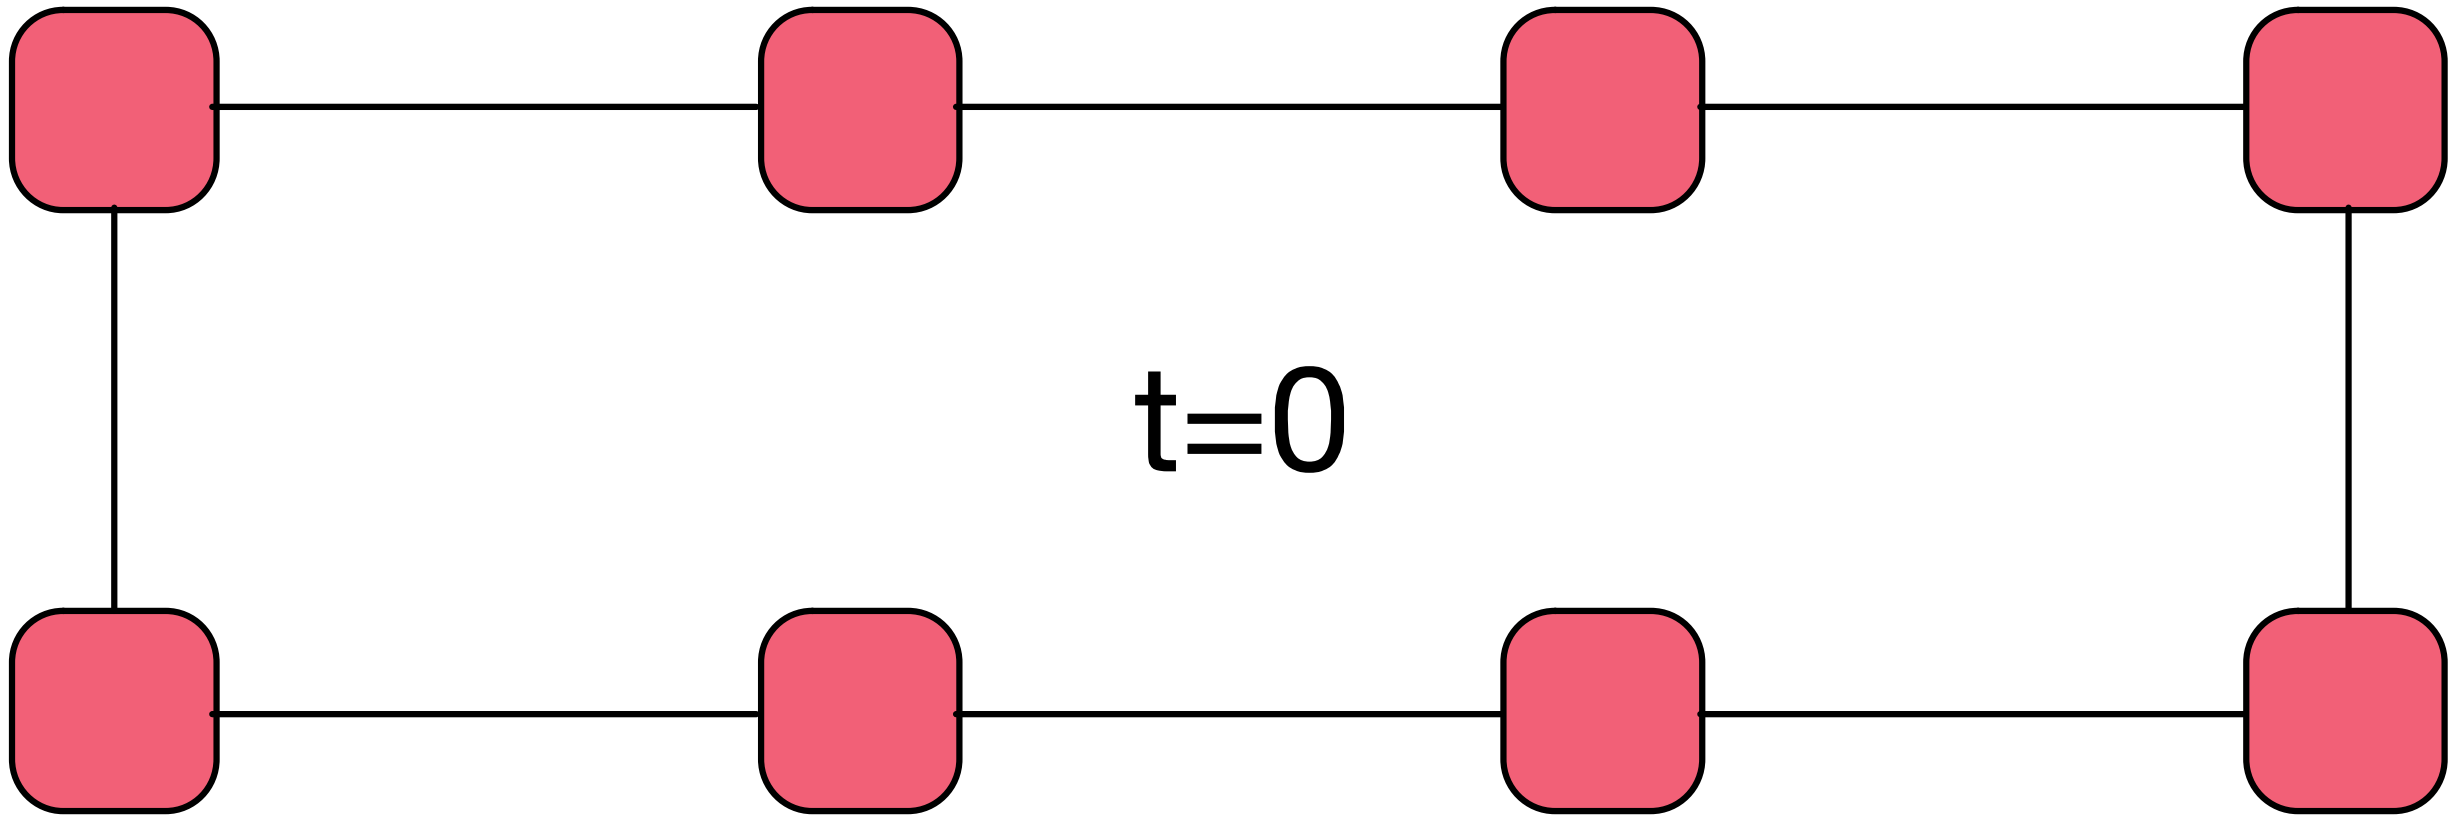
0 → 1



1 → All (Case 0 to 1)

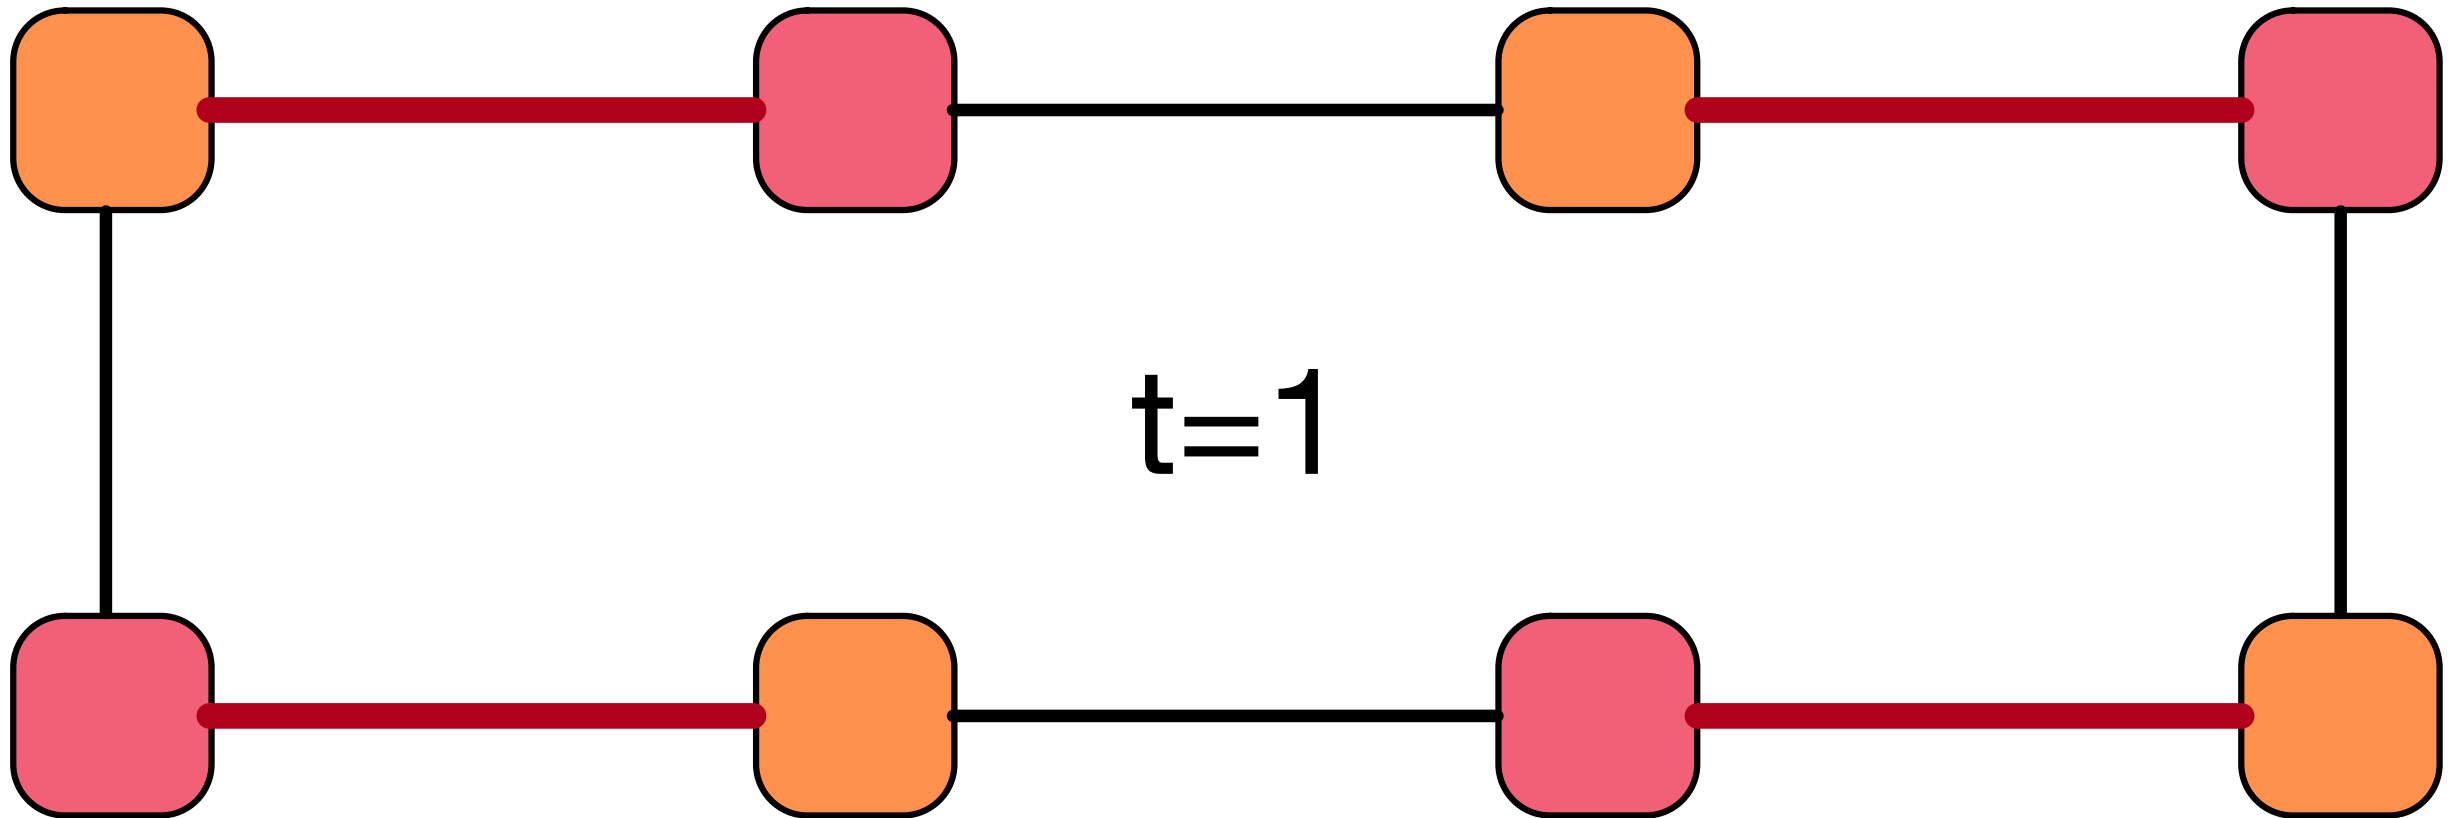
Reduction

26



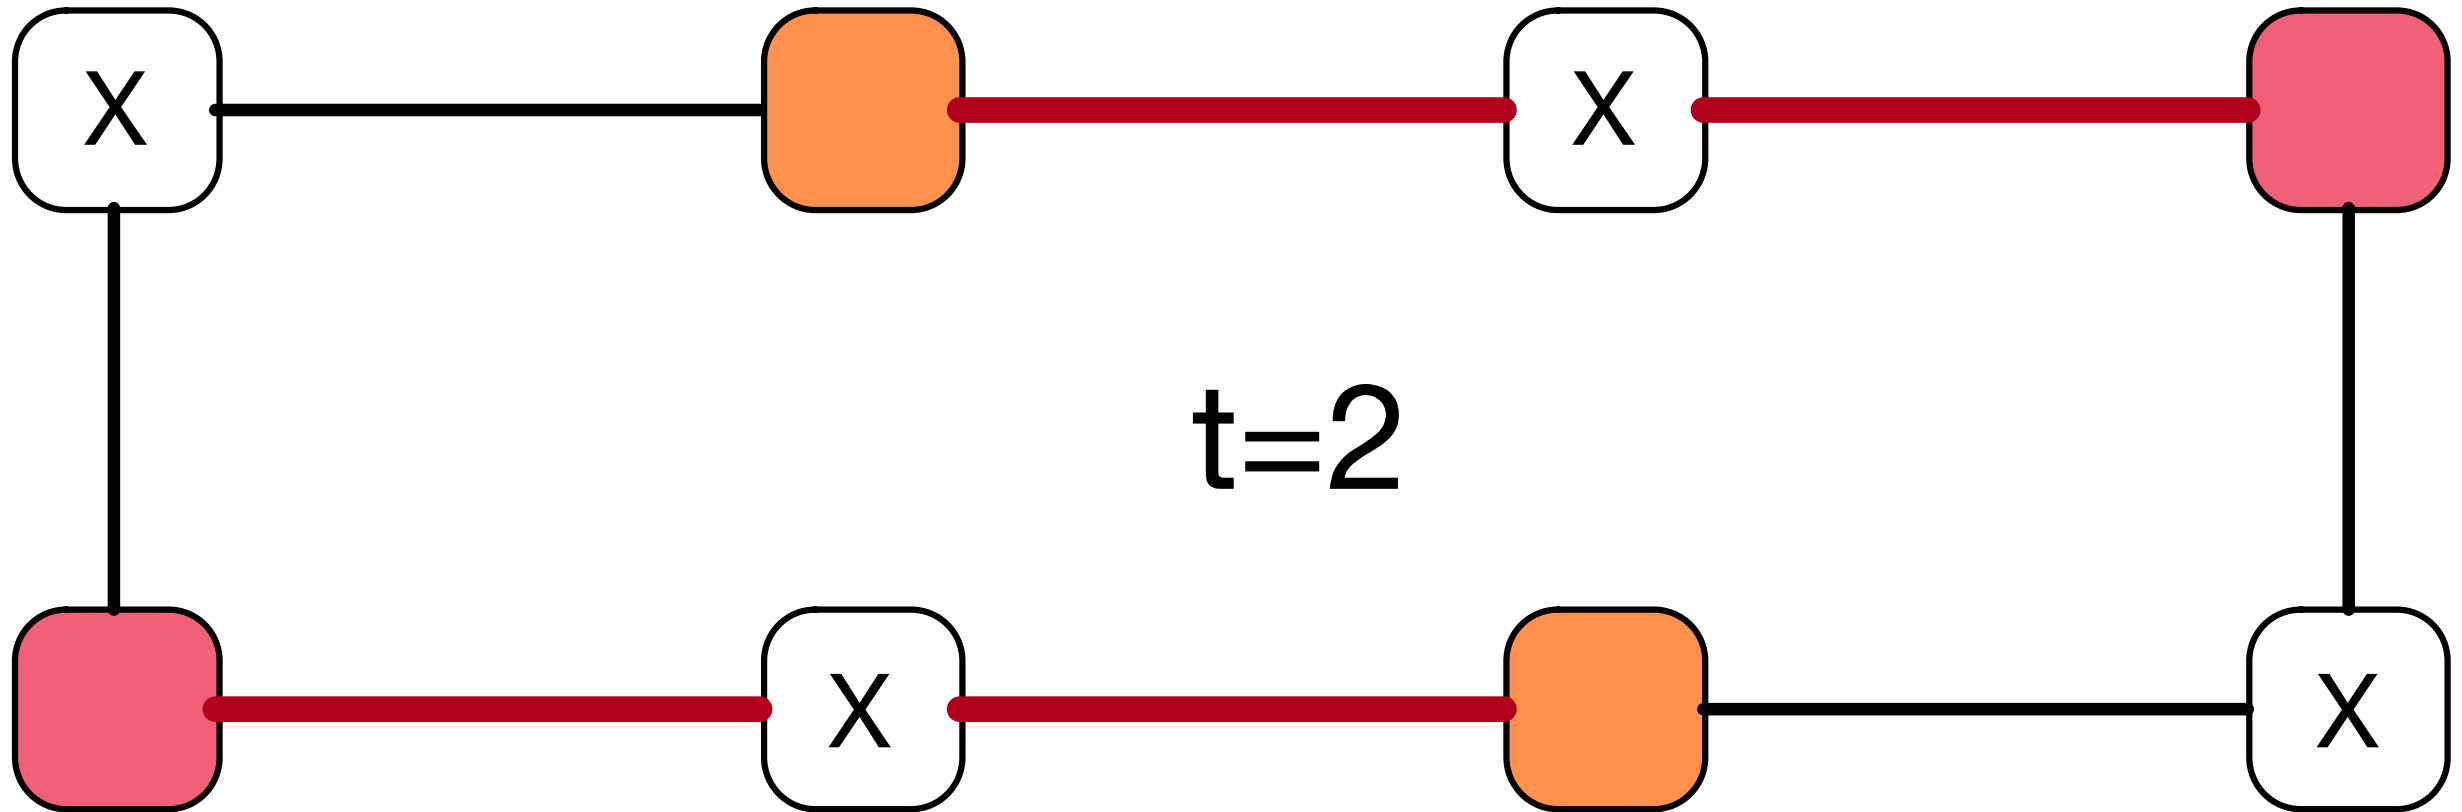
Reduction

26



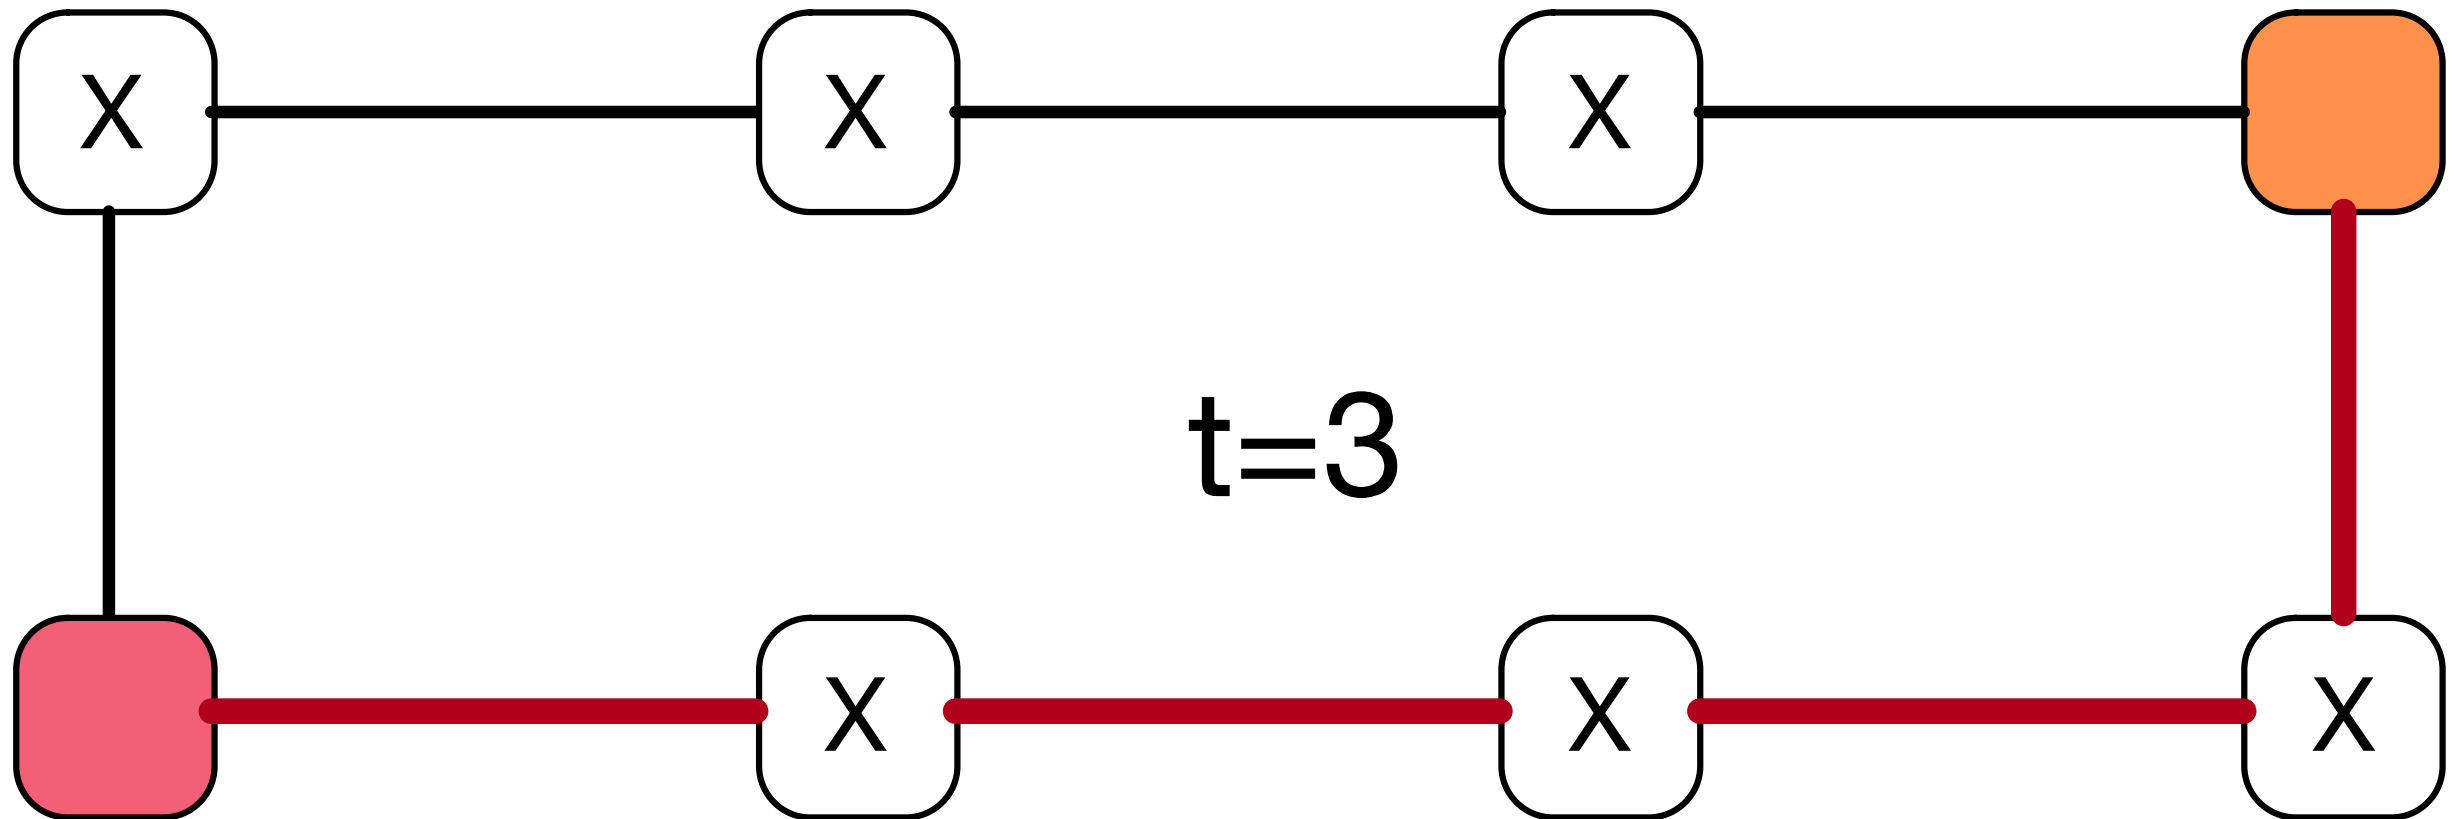
Reduction

26



Reduction

26



Reduction

26

