

OpenMP (cont...)

Acknowledgements: OpenMP Tutorial (LLNL, XSEDE) and OpenMP by Example (Texas A&M)

OpenMP → [Directives]

Sentinel + Directive + [clause[[,]
clause]...]

#pragma omp parallel private(a,b)

Directive: Parallel

#pragma omp parallel

Spawns new threads...

How do you set the max number of threads?

omp_set_num_threads()
OMP_NUM_THREADS

```
#include <omp.h>
#include <stdio.h>
int main()
{
    #pragma omp parallel
    printf("Hello from thread %d,\n",
           nthreads %d\n",
           omp_get_thread_num(),
           omp_get_num_threads());
}
```

Parallel Clauses

Variable Scoping

1. `shared (list)`
2. `private (list)`
3. `firstprivate (list)`

Parallel Clauses

- 4. copyin (list)
- 5. if (*scalar expression*)
- 6. reduction (operator:list)

You want to dynamically switch
between serial and parallel.

How would you do it?

```
#pragma omp parallel if (n>100000)
```

Other Directives

Loops

a.k.a. **Parallelize a loop**

```
#pragma omp parallel  
#pragma omp for ... →  
#pragma omp parallel for ...
```

1. Limited to the loop that immediately follows it
2. DO NOT change the **iteration variable**

Clauses:

- ✓ private(list)
- ✓ firstprivate(list)
- 1. lastprivate(list)
- 2. reduction
- 3. schedule(list)
- 4. ordered
- 5. nowait

lastprivate (list)

Similar to private

Value is available outside the loop
Value → What would be at the end of the loop, if code was serial

Could be 'ready', but not accessible as other threads maybe busy

```

13 #include <omp.h>
14 #include <stdio.h>
15 #include <stdlib.h>
16 int main (int argc, char *argv[])
17 {
18     int i, n;
19     float a[100], b[100], sum;
20     /* Some initializations */
21     n = 100;
22     for (i=0; i < n; i++)
23         a[i] = b[i] = i * 1.0;
24     sum = 0.0;
25     #pragma omp parallel
26     {
27         #pragma omp for reduction(+:sum)
28         for (i=0; i < n; i++)
29             sum = sum + (a[i] * b[i]);
30         printf("Completed work in thread: %d\n", omp_get_thread_num());
31     }
32     printf("Sum = %f\n", sum);
33 }
34

```

```

dhcp061174:OpenMP ashish$ gcc -fopenmp reduction.c
dhcp061174:OpenMP ashish$ ./reduction.c
Completed work in thread: 4
Completed work in thread: 1
Completed work in thread: 2
Completed work in thread: 6
Completed work in thread: 7
Completed work in thread: 0
Completed work in thread: 5
Completed work in thread: 3
Sum = 328350.000000
dhcp061174:OpenMP ashish$

```

Reduction clause: In a for loop

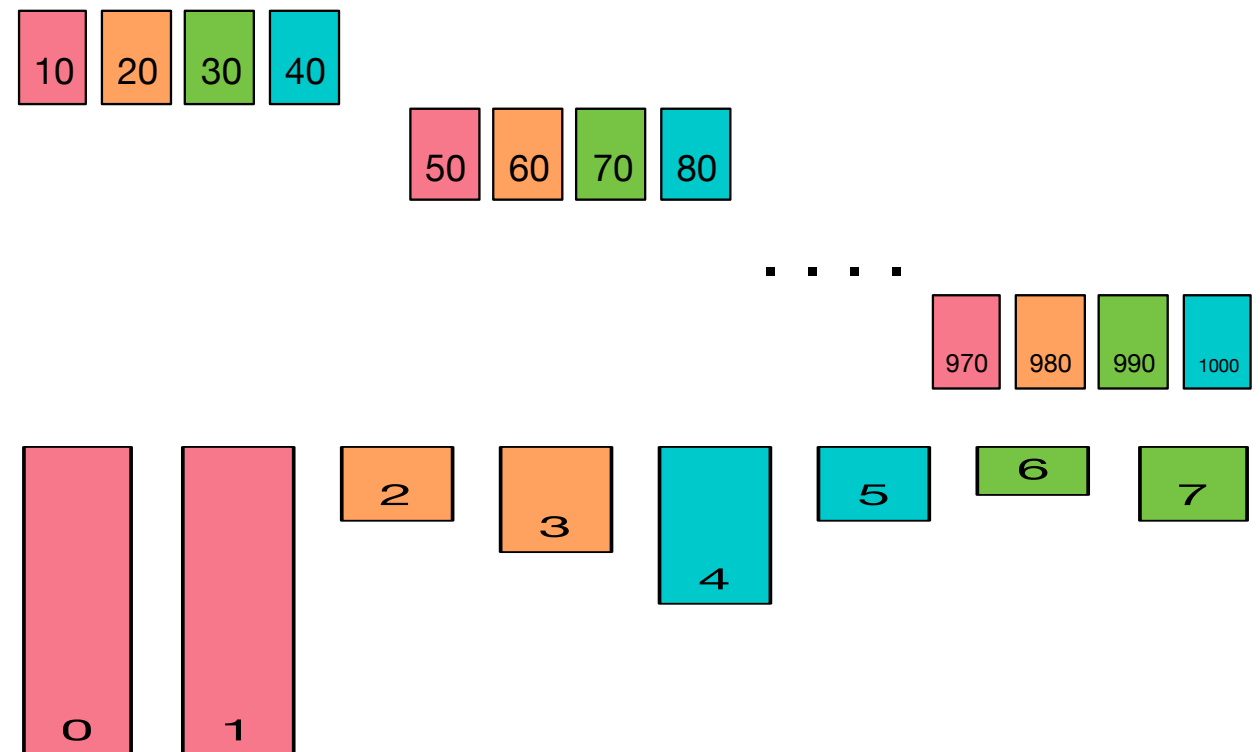
(see reduction.c)

Static Scheduling:

#pragma omp parallel for **schedule**(static, 10)

12

- OpenMP does not specify how loops are partitioned, the typical implementation is one where a loop is split equally across the number of threads
- Chunk size, if specified, indicates the number of iterations
→ Size = 1000;
4 threads; Chunk = 10
- Pros: Easy
- Cons → zz..... (load imbalance)



Runtime, Dynamic and Guided Scheduling

13

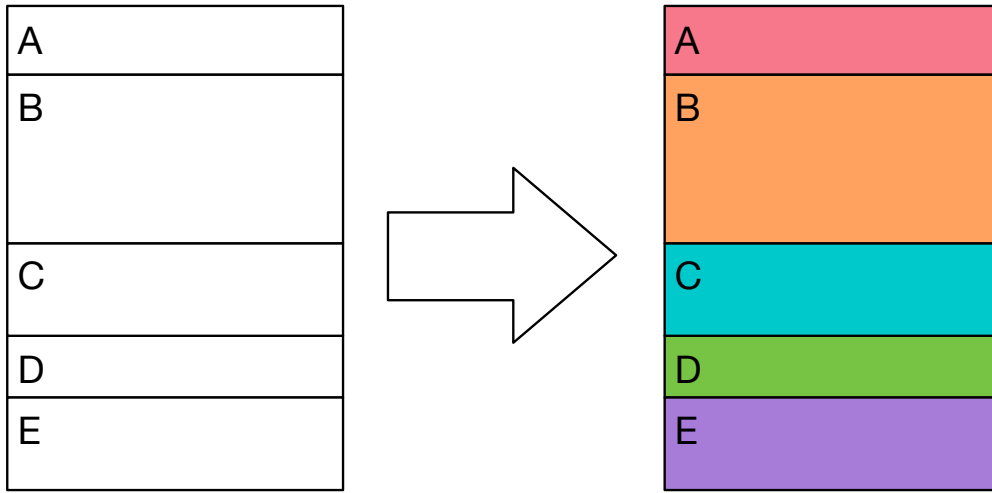
- runtime → specify using ENV variable (OMP_SCHEDULE)
- #pragma omp parallel for `schedule (dynamic, n)`
 - Chunk size of n (default 1)
 - When thread finishes one chunk, it is assigned a new one
- #pragma omp parallel for `schedule (guided, n)`
 - Chunk size is relative to the num of iterations left
 - Iterations are dynamically assigned to threads in blocks as threads request them
 - Similar to dynamic except that the block size decreases each time a parcel of work is given to a thread.
 - The size of the initial block is proportional to: $\text{\#iterations} / \text{\#threads}$
 - Subsequent blocks are proportional to: $\text{\#iterations_remaining} / \text{\#threads}$
 - The chunk parameter defines the minimum block size. The default chunk size is 1.

nowait

Don't wait for all threads to finish,
you can proceed

```
#pragma omp parallel
{
    #pragma omp for nowait
    for (i=1; i<n; i++)
        b[i] = (a[i] + a[i-1]) / 2.0;

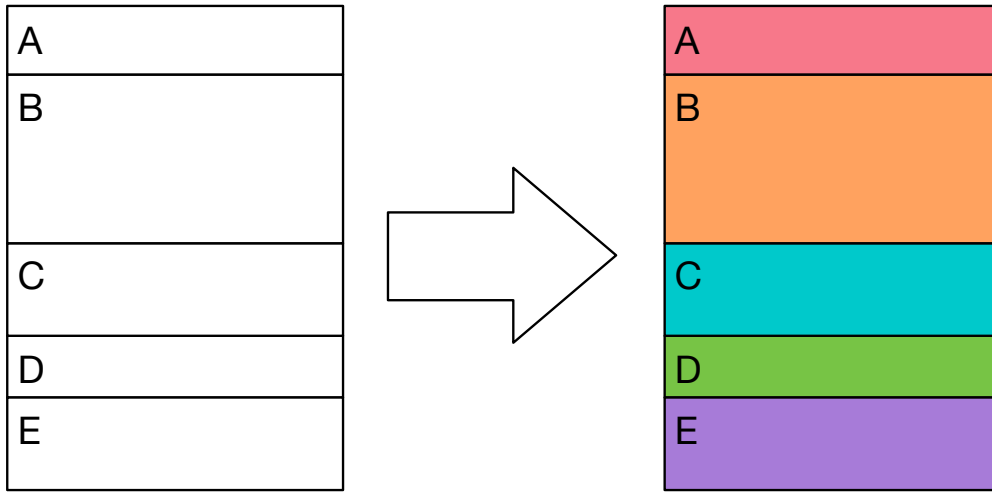
    #pragma omp for nowait
    for (i=0; i<m; i++)
        y[i] = sqrt(z[i]);
}
```



- Independent SECTION directives are nested within a SECTIONS directive.
- Each SECTION is executed once by a thread in the team.
- Different sections may be executed by different threads.
- It is possible for a thread to execute more than one section if it is quick enough and the implementation permits such.

Sections

a non-iterative work-sharing construct



Sections

a non-iterative work-sharing construct

```
1  #pragma omp sections [clause ...] newline
2      private (list)
3      firstprivate (list)
4      lastprivate (list)
5      reduction (operator: list)
6      nowait
7  {
8      #pragma omp section
9          structured_block
10
11     #pragma omp section
12         structured_block
13 }
```



```
1  #pragma omp single [clause ...] newline
2      private (list)
3      firstprivate (list)
4      nowait
5
6      structured_block
```

Other threads will wait
Useful for thread-unsafe code
Useful for I/O operations

Single

Only one thread will execute this

Synchronization

18

- OpenMP programs use shared variables to communicate. We need to make sure these variables are not accessed at the same time by different threads — **Why?**
- Available Directives:
 - Master
 - Critical
 - Atomic
 - Barrier

Master → #pragma omp master

- This Directive ensures that only the master threads executes instructions in the block. There is no implicit barrier so other threads will not wait for master to finish

Critical → #pragma omp critical

20

- Specifies a region of code that must be executed by only one thread at a time.
- If another threads reaches the critical section it will wait until the current thread finishes this critical section.
- Every thread will execute the critical block and they will synchronize at end of critical section

```
#include <omp.h>
main()
{
    int x;
    x = 0;
    #pragma omp parallel shared(x) {
        #pragma omp critical
            x = x + 1;
    } /* end of parallel section */
}
```

Barrier

`#pragma omp barrier`

A barrier will force every thread to wait at the barrier until all threads have reached the barrier.

Atomic

`#pragma omp atomic`

This Directive is very similar to the **CRITICAL** directive on the previous slide. Difference is that **ATOMIC** is only used for the update of a memory location.

Write a simple matrix multiplier

$$C[N][1] = A[N][N] * B[N][1]$$

1. Initialize as follows:

```
srand((int)time(NULL));  
X = rand() % 10
```

2. Surround timers around the loop construct

```
Start = omp_get_wtime();  
    {C = A x B}  
End = omp_get_wtime();  
printf("---- Parallel done in %f seconds.\n", End - Start);
```

3. Run for N = 100000