

Message Passing (cont...)

- Job Name

```
#!/bin/sh
```

```
#PBS -N HelloWorld
```

```
# nodes is the number of computers to use. Adjust as needed. MAX is 8. Leave ppn alone. Adjust walltime to max time to allow your binary to run.
```

- n = num of nodes (Max 8)
- Max time for Job

```
#PBS -l nodes=4:ppn=32,mem=2GB, walltime=00:00:15
```

```
#PBS -q batch
```

```
# If you want email notifications, remove the extra # in front of the #PBS -M and #PBS -m lines
```

```
##PBS -M <YOUR EMAIL ADDRESS HERE>
```

```
##PBS -m abe
```

- All transactions relative to this directory

```
# Set the correct PATH and env for the MPI implementation
```

```
#!/usr/lib64/{openmpi,mpich,mvapich2}/bin
```

```
MPIRUN='/usr/lib64/openmpi/bin/mpirun -mca plm_rsh_agent rsh --map-by node --display-allocation -mca orte_base_help_aggregate 0 -mca btl ^openib'
```

- np → num of proc
- Name of parallel executable

```
WORK_HOME=/home/ashish/MPI/HelloWorld
```

```
cd $WORK_HOME
```

```
$MPIRUN -np 6 --machinefile $PBS_NODEFILE ./hello_world
```

Submitting your jobs (all files in `/opt/tools/mpi`)

0. Simple Setup

(5mins)

- Work individually or 2/group
- Login
- Makefile and PBS scripts (git)



1. Hello World

(15mins)

1. Each node figures out its rank/identity
2. Each node figures out size of community
3. Print "I am rank/size"

- Edit and run on 4 and 8 nodes

```
#include <mpi.h>
#include <stdio.h>
int main(int argc, char** argv){
    ...
}
```

Setup/Cleanup

```
MPI_Init(&argc, &argv);
MPI_Finalize();
```

Gather Information

```
MPI_Comm_rank (MPI_COMM_WORLD,
               &my_rank);
```

```
MPI_Comm_size (MPI_COMM_WORLD,
               &num_procs);
```



2. Comm.. in a ring

(10mins)

1. $0 \rightarrow 1, 1 \rightarrow 2, \dots N \rightarrow 0$
2. Send your rank
3. Print "<rank> got <data> from <sender>"

- Run on 8 nodes

Basics

```
MPI_Init(&argc, &argv);  
MPI_Finalize();  
MPI_Comm_rank(MPI_COMM_WORLD,  
               &my_rank);  
MPI_Comm_size(MPI_COMM_WORLD,  
               &num_procs);
```

Send/Recv.

```
MPI_Send(&message, sizeof(message),  
         MPI_INT, destination, tag,  
         MPI_COMM_WORLD);  
  
MPI_Recv(&message, sizeof(message)+1,  
         MPI_INT, source, tag,  
         MPI_COMM_WORLD, &status)
```

```
MPI_Status status
```



3. Broadcast (10mins)

1. Node 0 will send a number (say 27)
2. Nodes !0 will receive the number
3. Nodes !0 will print rank, size and number received.


- Edit and run on 8 nodes

Basics

```
MPI_Init(&argc, &argv);  
MPI_Finalize();  
MPI_Comm_rank(MPI_COMM_WORLD,  
               &my_rank);  
MPI_Comm_size(MPI_COMM_WORLD,  
               &num_procs);
```


Broadcast

```
MPI_Bcast(&message, sizeof(message),  
          MPI_INT, sender_rank,  
          MPI_COMM_WORLD);
```



```
rbuf  
real a(25), rbuf(MAX)  
...  
call mpi_gather(a, 25, MPI_REAL, rbuf, 25, MPI_REAL, root, comm, ierr)  
...
```

MPI_GATHER



```
sbuf  
real sbuf(MAX), rbuf(25)  
...  
call mpi_scatter(sbuf, 25, MPI_REAL, rbuf, 25, MPI_REAL, root, comm, ierr)  
...
```

MPI_SCATTER

$$\begin{array}{c}
 A \\
 \left[\begin{array}{c} \text{Process 0} \\ \text{-----} \\ \text{Process 1} \\ \text{-----} \\ \text{Process 2} \\ \text{-----} \\ \text{Process 3} \end{array} \right]
 \end{array}
 * b = c$$

$$\begin{array}{c}
 \\
 \left[\begin{array}{c} \\ \\ \\ \end{array} \right]
 \end{array}
 =
 \begin{array}{c}
 \\
 \left[\begin{array}{c} 0 \\ \text{-----} \\ 1 \\ \text{-----} \\ 2 \\ \text{-----} \\ 3 \end{array} \right]
 \end{array}$$

- Matrix is distributed by rows (i.e. row-major order)
- Vector is needed by all
- MPI_Gather will be used to collect the product

4. Matrix Multiply (40mins)

$$C[N] = A[N][N] \times B[N]$$

What is your decomposition strategy?
What is your aggregation strategy?

Solutions

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv)
{
    int rank, size;
    MPI_Init(&argc, &argv);
    MPI_Comm_rank (MPI_COMM_WORLD, &rank);
    MPI_Comm_size (MPI_COMM_WORLD, &size);
    printf("I am %d / %d\n", rank, size);
    MPI_Finalize();
}
```



```

#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv)
{
    int rank, size;
    MPI_Init(&argc, &argv);
    MPI_Comm_rank (MPI_COMM_WORLD, &rank);
    MPI_Comm_size (MPI_COMM_WORLD, &size);
    int send_msg = 27;
    int recv_msg;
    int dest;
    MPI_Status status;
    if(rank == size-1)
        dest = 0;
    else
        dest = rank+1;

    MPI_Send(&send_msg, sizeof(send_msg), MPI_INT, dest, 22, MPI_COMM_WORLD);
    MPI_Recv(&recv_msg, 20, MPI_INT, MPI_ANY_SOURCE, 22, MPI_COMM_WORLD, &status);

    printf("%d got %d from %d\n",rank, recv_msg, status.MPI_SOURCE);
    fflush(stdout);
    MPI_Finalize();
}

```



```
#include <stdio.h>
#include <mpi.h>
int main(int argc, char **argv)
{
    int message;
    int rank, size;
    MPI_Status status;
    int sender= 0;

    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    if (rank == sender)
    {
        message = 20;
    }
    MPI_Bcast(&message, sizeof(message), MPI_INT, sender, MPI_COMM_WORLD);
    if (rank != sender)
        printf( "Message from process %d : %d\n", rank, message);
    MPI_Finalize();
}
```

3. Broadcast (lab5.c)

