

# Housekeeping

Lookup the cluster script pinned to #general. **Run asap**

**No class on 3/16. In it's place we will extend the class on 3/14 from 4 – 6:30**

## **Final Project:**

1. Proposal Presentations due 3/14
2. Proposal writeup due 3/21

# Final Project

2

- Introductions (Name, members and project/idea) - **TODAY**
- **Present your proposal – 3/14**
  - Each team member will present one aspect of the project
  - Describe the project, related work, deliverables, schedule, and intellectual challenges (or what do you expect to gain from it).
  - Will be Q&A style with feedback
  - 15-20mins
- **Submit your proposal – 3/21**
  - 3-4 pages, single spaced
  - Response to any questions/concerns raised during the presentation
  - Details that could not be covered in the proposal, as well as schedule and deliverables.

# Final Project (cont...)

3

- Project Update – 4/6 (tentative)
  - 5-10 min presentation.
  - Are you on track?
- Final Presentation – 4/25 & 4/27 (tentative)
- Final Report – 5/2

# Some final words on Parallelization

Sources for this lecture:

- Jim Demmel/Kathy Yelick (Berkeley - CS 267)
- Mary Hall (Utah – CS4961)
- Software Systems for Numerical Analysis (Emory CS561)

<http://www.mathcs.emory.edu/~cheung/Courses/561/>

- Our focus has been Dense arrays and loop-based data-parallel computation has been the focus of this class so far
- Real world is rarely that dense and organized

**SPARSE MATRICES**

# Dense Linear Algebra vs. Sparse Linear Algebra


5



- Matrix vector multiply:  
    for (i=0; i<n; i++)  
        for (j=0; j<n; j++)  
            a[i] += c[j][i]\*b[j];
- What if n is very large, and some large percentage (say 90%) of c is zeros?
- Should you represent all those zeros?
- If not, how to represent "c"?

# Not a hypothetical situation

6

- Most matrices arising from real applications are sparse.
- A 1M-by-1M submatrix of the web connectivity graph, constructed from an archive at the Stanford WebBase.





### Matrix properties

number of rows	52,329
number of columns	52,329
nonzeros	2,698,463

Examples from Tim Davis's Sparse Matrix Collection,  
<http://www.cise.ufl.edu/research/sparse/matrices/>

Boeing CT20 ENGINE BLOCK -- STIFFNESS MATRIX

# Sparse Linear Algebra

8

- Suppose you are applying matrix-vector multiply and the matrix has lots of zero elements
    - Computation cost? Space requirements?
  - General sparse matrix representation concepts
    - Primarily only represent the nonzero data values
    - Auxiliary data structures describe placement of nonzeros in “dense matrix”
1. COO → Coordinate-wise Storage Method
  2. CSR/CCR → Compressed Sparse Row (or Column) Method





Sparse Matrix Representations and Iterative Solvers – Nathan Bell, nVidia

COO → Coordinate-wise Storage Method

	0	1	2	3	4	5	6	7		
0	[	11	12		14					]
1	[		22	23	25					]
2	[	31		33	34					]
3	[		42		45	46				]
4	[				55					]
5	[				65	66	67			]
6	[				75		77	78		]
7	[						87	88		]

Representation using Coordinate-wise method:											
Index	0	1	2	3	4	5	6	7	8	9	10
Val	11	12	14	22	23	25	31	33	34	42	45
Row	0	0	0	1	1	1	2	2	2	3	3
Col	0	1	3	1	2	4	0	2	3	1	4
Index	11	12	13	14	15	16	17	18	19	20	21
Val	46	55	65	66	67	75	77	78	87	88	-
Row	3	4	5	5	5	6	6	6	7	7	-
Col	5	4	4	5	6	4	6	7	6	7	-

# Detailed Example

Val[N]: value of the non-zero elements

Row[N]: row-index of non-zero elements

Col[N]: column -index of non-zero elements

	0	1	2	3	4	5	6	7	
0	[	11	12	14					]
1	[		22	23	25				]
2	[	31		33	34				]
3	[		42		45	46			]
4	[				55				]
5	[				65	66	67		]
6	[				75		77	78	]
7	[						87	88	]

Representation using Coordinate-wise method:

Index	0	1	2	3	4	5	6	7	8	9	10
Val	11	12	14	22	23	25	31	33	34	42	45
Row	0	0	0	1	1	1	2	2	2	3	3
Col	0	1	3	1	2	4	0	2	3	1	4

Index	11	12	13	14	15	16	17	18	19	20	21
Val	46	55	65	66	67	75	77	78	87	88	-
Row	3	4	5	5	5	6	6	6	7	7	-
Col	5	4	4	5	6	4	6	7	6	7	-

# Detailed Example

# Matrix Multiplication

```
for (k = 0; k < N; k = k + 1)
    result[i] = 0;
```

```
for (k = 0; k < nnz; k = k + 1)
    result[Row[k]] += Val[k]*d[Col[k]];
```

	0	1	2	3	4	5	6	7		
0	[ 11	12		14					]	[ 3 ]
1	[	22	23		25				]	[ 5 ]
2	[ 31		33	34					]	[ 2 ]
3	[	42			45	46			]	[ 1 ]
4	[				55				]	x [ 0 ] = ???
5	[				65	66	67		]	[ 1 ]
6	[				75		77	78	]	[ 2 ]
7	[						87	88	]	[ 4 ]

**A[7][7]** **v[7]**

Suppose  $d[0..N-1] = (1, 0, 0, 0, 0, 0, 0, 0)$

[ 11	12	14									[ 1 ]	[ 11 ]
[	22	23	25								[ 0 ]	[ 0 ]
[ 31		33	34								[ 0 ]	[ 31 ]
[	42		45	46							[ 0 ]	[ 0 ]
[			55								[ 0 ]	[ 0 ]
[			65	66	67						[ 0 ]	[ 0 ]
[			75		77	78					[ 0 ]	[ 0 ]
[				87	88						[ 0 ]	[ 0 ]

After:

```

for (k = 0; k < N; k = k + 1)
    result[i] = 0;
    ----> [0, 0, 0, 0, 0, 0, 0, 0]

k = 0;
result[Row[0]] = result[Row[0]] + Val[0]*d[Col[0]];
result[0] = result[0] + Val[0]*d[0];
result[0] = 0 + 11*1;
    ----> [11, 0, 0, 0, 0, 0, 0, 0]

```

```

k = 1;
result[Row[1]] = result[Row[1]] + Val[1]*d[Col[1]];
result[0] = result[0] + Val[1]*d[1];
result[0] = 0 + 12*0;
    ----> [11, 0, 0, 0, 0, 0, 0, 0]

k = 2;
result[Row[2]] = result[Row[2]] + Val[2]*d[Col[2]];
result[0] = result[0] + Val[2]*d[3];
result[0] = 0 + 14*0;
    ----> [11, 0, 0, 0, 0, 0, 0, 0]

...
k = 6;
result[Row[6]] = result[Row[6]] + Val[6]*d[Col[6]];
result[2] = result[2] + Val[6]*d[0];
result[2] = 0 + 31*1;
    ----> [11, 0, 31, 0, 0, 0, 0, 0]

```

## Sample Execution

Slides from: Emory – CS561 (<http://goo.gl/HUYKzP>)

```

[ 11 12      14      ]
[      22 23      25  ]
[ 31      33 34      ]
[      42      45 46  ]
[      55      ]
[      65 66 67      ]
[      75      77 78  ]
[      87 88  ]

```

Representation using Compressed Sparse Row method:

Index	0	1	2	3	4	5	6	7	8	9	10
Val	11	12	14	22	23	25	31	33	34	42	45
RowPtr	0	3	6	9	12	13	16	19	21	-	-
Col	0	1	3	1	2	4	0	2	3	1	4
Index	11	12	13	14	15	16	17	18	19	20	21
Val	46	55	65	66	67	75	77	78	87	88	-
RowPtr	-	-	-	-	-	-	-	-	-	-	-
Col	5	4	4	5	6	4	6	7	6	7	-

# Compressed Sparse Row

Val[N]: contains the value of the non-zero elements

Col[N]: contains the column-index of the non-zero elements

RowPtr[N]: contains the row-index range of the non-zero elements

```

[ 11 12      14      ]
[      22 23      25  ]
[ 31      33 34      ]
[      42      45 46  ]
[      55      ]
[      65 66 67      ]
[      75      77 78  ]
[      87 88  ]

```

Representation using Compressed Sparse Row method:

Index	0	1	2	3	4	5	6	7	8	9	10
Val	11	12	14	22	23	25	31	33	34	42	45
RowPtr	0	3	6	9	12	13	16	19	21	-	-
Col	0	1	3	1	2	4	0	2	3	1	4

Index	11	12	13	14	15	16	17	18	19	20	21
Val	46	55	65	66	67	75	77	78	87	88	-
RowPtr	-	-	-	-	-	-	-	-	-	-	-
Col	5	4	4	5	6	4	6	7	6	7	-

# Compressed Sparse Row

Val[N]: contains the value of the non-zero elements

Col[N]: contains the column-index of the non-zero elements

RowPtr[N]: contains the row-index range of the non-zero elements

```

[ 11 12      14      ]
[      22 23      25  ]
[ 31      33 34      ]
[      42      45 46  ]
[      55      ]
[      65 66 67      ]
[      75      77 78  ]
[      87 88  ]

```

Representation using Compressed Sparse Row method:

Index	0	1	2	3	4	5	6	7	8	9	10
Val	11	12	14	22	23	25	31	33	34	42	45
Col	0	1	3	1	2	4	0	2	3	1	4

Index	11	12	13	14	15	16	17	18	19	20	21
Val	46	55	65	66	67	75	77	78	87	88	-
Col	5	4	4	5	6	4	6	7	6	7	-

```

0 1 3 | 1 2 4 | 0 2 3 | 1 4 5 | 4 | 4 5 6 | 4 6 7 | 6 7 |
0    | 3    | 6    | 9    | 12 | 13    | 16    | 19 | 21

```

# Compressed Sparse Row

Val[N]: contains the value of the non-zero elements

Col[N]: contains the column-index of the non-zero elements

RowPtr[N]: contains the row-index range of the non-zero elements



# Matrix Multiplication

```
for (k = 0; k < N; k = k + 1)
    result[i] = 0;

for (i = 0; i < N; i = i + 1)
{
    for (k = RowPtr[i]; k < RowPtr[i+1]; k++)
    {
        result[i] += Val[k]*d[Col[k]];
    }
}
```

Suppose  $d[0..N-1] = (1, 0, 0, 0, 0, 0, 0, 0)$

[ 11	12		14					]	[ 1 ]	[ 11 ]
[		22	23		25			]	[ 0 ]	[ 0 ]
[ 31			33	34				]	[ 0 ]	[ 31 ]
[		42			45	46		]	x [ 0 ]	= [ 0 ]
[					55			]	[ 0 ]	[ 0 ]
[					65	66	67	]	[ 0 ]	[ 0 ]
[					75		77	78 ]	[ 0 ]	[ 0 ]
[							87	88 ]	[ 0 ]	[ 0 ]

After:

```
for (k = 0; k < N; k = k + 1)
    result[i] = 0;
    ----> [0, 0, 0, 0, 0, 0, 0, 0]

i = 0;
k = 0;
result[0] = result[0] + Val[0]*d[Col[0]];
result[0] = result[0] + Val[0]*d[0];
result[0] = 0 + 11*1;
    ----> [11, 0, 0, 0, 0, 0, 0, 0]

i = 0;
k = 1;
result[0] = result[0] + Val[1]*d[Col[1]];
result[0] = result[0] + Val[1]*d[1];
result[0] = 0 + 12*0;
    ----> [11, 0, 0, 0, 0, 0, 0, 0]
```

```
i = 0;
k = 2;
result[0] = result[0] + Val[2]*d[Col[2]];
result[0] = result[0] + Val[2]*d[3];
result[0] = 0 + 14*0;
    ----> [11, 0, 0, 0, 0, 0, 0, 0]

...
i = 2;
k = RowPtr[2] = 6;
result[2] = result[2] + Val[6]*d[Col[6]];
result[2] = result[2] + Val[6]*d[0];
result[2] = 0 + 31*1;
    ----> [11, 0, 31, 0, 0, 0, 0, 0]
```

# HW 3

Individual or Team of 2  
Due 3/20

Large scale problems so develop a testing and evaluation strategy.

1. Implement the quicksort on the cluster. Implement for a  $nD$  hypercube
2. Implement Google's PageRank algorithm using MPI.
  1. Use data from the Stanford Network Repository
  2. Assume a damping of 0.85
  3. Due 3/20
  4. Submit a brief report w/ your code.